
A Survey of Multimodal Language Models and Related Technologies

www.surveyx.cn

Abstract

This survey paper provides a comprehensive examination of multimodal language models (MMLMs), highlighting their significance in integrating diverse data modalities to enhance human-machine interaction. By leveraging advanced algorithms and neural networks, MMLMs address the limitations of unimodal approaches, thereby improving applications across healthcare, education, creative industries, and more. The paper explores the interconnected fields of artificial intelligence, computational linguistics, and machine learning, which collectively contribute to the development of robust models capable of processing and generating human-like communication. Key methodologies for multimodal data integration, such as transformer-based architectures and encoder-decoder frameworks, are discussed alongside the challenges of data alignment, computational complexity, and interpretability. The survey also delves into the capabilities and optimization strategies of large language models (LLMs), emphasizing their transformative impact on natural language processing and multilingual applications. Furthermore, the paper addresses ethical considerations, including bias mitigation and transparency, essential for the responsible deployment of AI-driven language models. Future research directions are identified, focusing on architectural optimization, interdisciplinary collaboration, and the development of explainability and evaluation metrics. By advancing these areas, researchers can unlock new applications and address complex societal challenges, ensuring that MMLMs and LLMs remain at the forefront of artificial intelligence innovation.

1 Introduction

The evolution of language models has significantly transformed human-machine interaction, paving the way for more intuitive and effective communication. Among the most noteworthy advancements in this domain are multimodal language models, which integrate various data modalities such as text, images, and audio. This integration facilitates complex reasoning processes and enhances the ability of machines to generate human-like responses. The significance of these models extends beyond mere technological innovation; they hold the potential to address critical challenges across diverse fields, including healthcare, education, and misinformation detection. This review aims to explore the multifaceted nature of multimodal language models, their interconnectedness with various disciplines, and the implications of their applications in real-world scenarios.

1.1 Significance of Multimodal Language Models

Multimodal language models represent a pivotal advancement in human-machine interaction by integrating diverse data modalities, including text, images, and audio, to enable machines to perform complex reasoning and generate human-like communication. These models address inherent limitations in unimodal approaches by leveraging the synergistic potential of multimodal data, thereby enhancing their applicability across various domains [1]. In healthcare, multimodal models significantly improve diagnostic accuracy and clinical reporting by integrating data from medical imaging

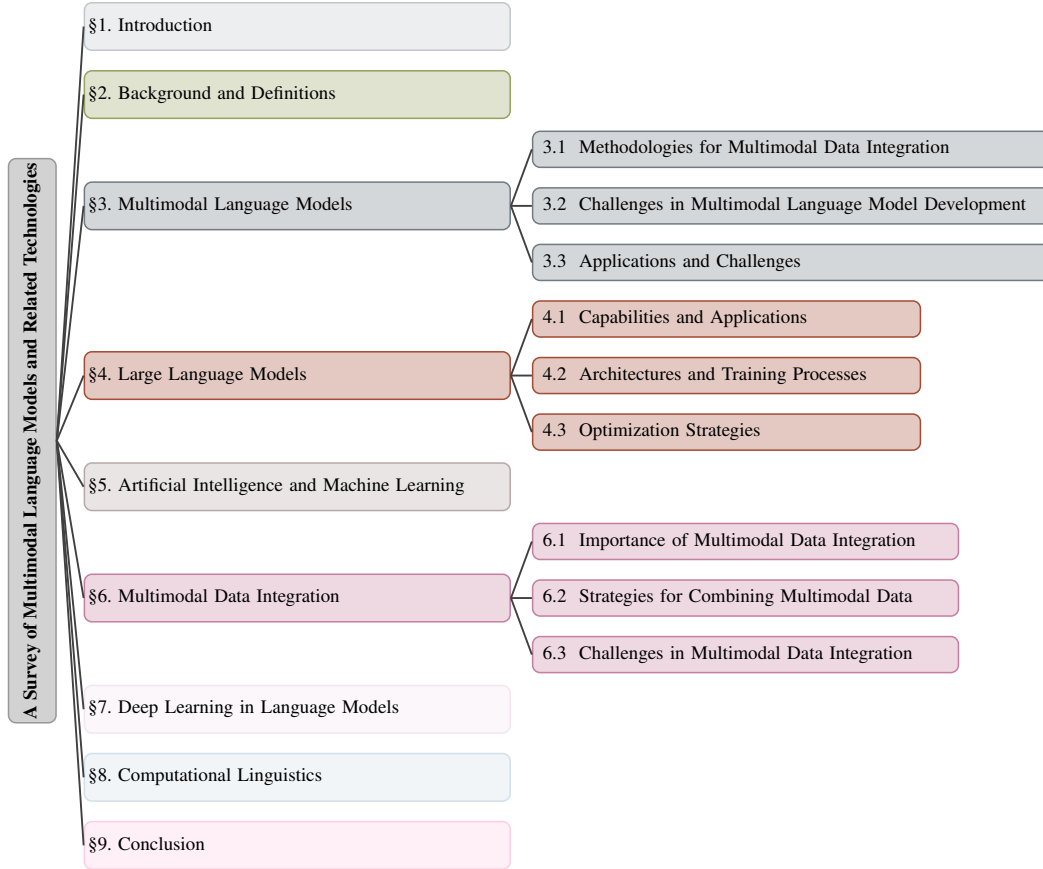


Figure 1: chapter structure

and genomic sources. For instance, they have been effectively utilized in generating clinical reports from chest radiographs and predicting phenotypes from genomic data, showcasing their ability to synthesize complex information. Moreover, the use of automated tools for identifying cognitive impairment underscores the transformative impact of multimodal models in early detection and healthcare improvement. Additionally, the classification of causes of death from verbal autopsy reports highlights their importance in providing accurate information in lower-and-middle-income countries, where traditional data collection methods may be limited.

In education, multimodal language models facilitate adaptive learning experiences and enhance accessibility for diverse learner populations. Benchmarks like UGIF (UI Grounded Instruction Following) offer comprehensive tutorials for smartphone navigation, addressing common user challenges such as blocking calls from unknown numbers. By integrating natural language processing techniques for retrieval, parsing, and grounding, UGIF enhances user experience through multimodal systems that link step-by-step instructions directly to the smartphone’s user interface. This approach not only provides multilingual support with a dataset containing 4,184 tasks across eight languages but also highlights the importance of cross-modal and cross-lingual retrieval. Such advancements significantly improve task completion rates and overall usability for diverse user populations [2, 3, 4, 5]. Furthermore, these models contribute to creative industries by generating innovative outputs, such as Urdu poetry, through deep learning methodologies, showcasing their transformative impact on natural language generation. In the realm of conversational AI, multimodal models utilize prompt engineering techniques to imitate the language style of real individuals using limited text data, further enhancing user engagement.

The integration of multimodal data plays a crucial role in addressing societal challenges, such as combating misinformation, by leveraging diverse data sources—like text, images, and audio—to enhance deception detection across multiple languages and contexts. This approach not only aims to identify false information more effectively but also addresses the complexities of linguistic variations

and the multimodal nature of digital content. Such advancements foster a more comprehensive understanding of how misinformation spreads in today’s interconnected world [6, 5, 3, 7, 8]. Robust detection systems, particularly during crises like the COVID-19 pandemic, have leveraged multimodal architectures to improve public communication and trust. In business operations, multimodal models optimize workflows by enhancing recommendation systems and automating complex tasks, such as network configuration and business process management. Their ability to address challenges like the cold-start problem and improve personalization underscores their value in operational efficiency and customer engagement.

From a foundational research perspective, multimodal models unify generative and discriminative methodologies in visual foundation models, addressing the fragmented nature of existing literature and advancing human-machine interaction [9]. Applications such as TrackGPT further illustrate the significance of multimodal systems in forecasting entity positions using historical track data, emphasizing their versatility across diverse contexts. Transfer learning techniques have also gained prominence in multimodal systems, fostering advancements in computer vision, medical imaging, and natural language processing. The transformative impact of multimodal language models lies in addressing contemporary challenges, such as data alignment, computational complexity, and interpretability. Their applications in sentiment analysis and Visual Question Answering (VQA) exemplify their ability to combine textual and visual data for sophisticated reasoning tasks, enhancing multimodal image interpretation in fields like ophthalmology. As these models continue to evolve, their potential to enhance human-machine collaboration across domains remains indispensable [10]. Furthermore, the integration of multimodal language models in robotic systems, such as embodied robotic waiters, highlights their transformative impact on human-robot interaction, allowing robots to understand language instructions and react accordingly to visual perception [11]. The survey on cross-field engagement quantifies the influence of NLP on other fields, emphasizing the significance of multimodal language models in bridging knowledge gaps and fostering interdisciplinary collaboration [12].

The exploration of multimodal language models reveals their profound significance across various sectors, illustrating their capacity to address complex challenges and enhance human-machine interactions. This multifaceted impact underscores the necessity for continued research and development in this area, as the potential applications of these models are vast and varied, with implications for future technological advancements.

1.2 Interconnected Fields and Technologies

The development of multimodal language models is inherently interdisciplinary, necessitating the integration of diverse domains such as artificial intelligence (AI), computational linguistics, machine learning, and multimodal data integration. These interconnected fields collectively contribute to the advancement of systems capable of understanding and generating human-like communication through the synergistic utilization of diverse data modalities. The synthesis of AI techniques with computational linguistics is crucial for enabling machines to process and interpret natural language, as well as integrate multimodal inputs, including text, images, and audio [13]. This integration is essential for creating robust models that can perform complex reasoning tasks and enhance human-machine interactions.

Artificial intelligence serves as the foundational framework driving innovations across these domains. The integration of AI and machine learning is pivotal in developing effective methodologies, such as video summarization, which require the processing of both visual and textual inputs [14]. Furthermore, the integration of computer vision and natural language processing is essential for the development of multimodal language models, as it allows for the representation and understanding of visual and linguistic data simultaneously [15]. This interconnected nature is crucial in advancing multimodal language models, particularly in tasks such as emotion detection, where the ability to interpret non-verbal cues alongside textual information can significantly enhance the model’s performance.

The role of multimodal data integration further exemplifies the interconnectedness of these disciplines. By combining data from disparate sources, multimodal models achieve greater robustness and accuracy in applications ranging from healthcare to social sciences. The integration of visual and textual inputs, as explored in robotic manipulation tasks, highlights the collaborative nature of AI, robotics, and multimodal data integration [13]. This collaboration is essential for developing systems that can interpret and respond to complex stimuli, enhancing the capabilities of human-

machine interaction systems. The implications of such advancements extend beyond mere technical achievements; they present opportunities for addressing significant societal challenges, including accessibility and the digital divide.

The interplay between these fields not only fosters innovation but also catalyzes the development of novel applications that can transform industries. The convergence of AI, machine learning, and multimodal data integration is poised to redefine how we interact with technology, offering new avenues for enhancing user experiences and improving outcomes across various domains. As research continues to progress, the potential for interdisciplinary collaboration will only grow, driving further advancements in multimodal language models.

1.3 Structure of the Survey

This survey is systematically organized to provide a comprehensive overview of multimodal language models and related technologies. The paper begins with an **Introduction** that sets the stage by highlighting the significance of multimodal language models and their transformative impact on human-machine interaction. This section also delineates the interconnected fields that contribute to the development of these models, including artificial intelligence, computational linguistics, and multimodal data integration.

The **Background and Definitions** section follows, offering foundational insights into key concepts and definitions pertinent to the survey topic. It clarifies the scope and relevance of terms such as large language models, artificial intelligence, and machine learning, establishing a clear framework for understanding multimodal language models. This foundational knowledge is critical for readers to grasp the complexities involved in the development and application of these advanced systems.

In the subsequent section, **Multimodal Language Models**, the survey delves into methodologies and architectures essential for developing these models, alongside the challenges encountered in their creation. This section emphasizes strategies for integrating diverse data types and highlights innovative applications and real-world challenges. Through a thorough examination of various methodologies, the paper aims to illuminate the ongoing advancements and the potential future directions for research in this area.

The examination of **Large Language Models** provides insights into their capabilities, architectures, and optimization strategies. This section explores the transformative impact of large language models on natural language processing and their diverse applications, underscoring the importance of scalability and efficiency in model design.

The role of **Artificial Intelligence and Machine Learning** is discussed next, focusing on key algorithms and techniques that advance language model development. The section also addresses bias and ethical considerations crucial to AI-driven language models, emphasizing the need for responsible development practices that mitigate potential harms associated with these technologies.

The analysis of **Multimodal Data Integration** highlights its critical role in enhancing language models, focusing on effective strategies for combining diverse data types, such as text, numbers, and tables. This section addresses the significant challenges encountered during this integration process, including issues of data quality and model performance in tasks like long-form summarization and data augmentation [16, 17, 8].

The application of **Deep Learning in Language Models** is detailed, focusing on neural network architectures and their efficiency. The exploration of attention mechanisms and model interpretability reveals their significant influence on the performance of language models. Attention mechanisms are categorized into various architectures that enhance understanding of input representations and their effects on output predictions, particularly in tasks such as multiple-choice reading comprehension and sentiment classification, where the semantic meaning of inputs often outweighs linguistic factors [18, 19].

The survey then explores **Computational Linguistics**, highlighting its intersection with language models and its contributions to applications such as hate speech detection and semantic coherence. This exploration underscores the importance of linguistic principles in the development of effective language models that can comprehend and generate nuanced human communication.

Finally, the **Conclusion** summarizes key findings, advancements, challenges, and future research directions. The text explores the latest trends and applications in language model development, highlighting the critical need for explainability and robust evaluation metrics. It identifies nine key challenges across various text generation tasks, including bias, reasoning, and interpretability, while advocating for improved evaluation methodologies to address these issues. The discussion also underscores the importance of developing effective detection mechanisms to combat the misuse of text generation models, such as the creation of fake news and misleading product reviews, thereby ensuring the responsible deployment of these technologies [20, 4]. The following sections are organized as shown in Figure 1.

2 Background and Definitions

2.1 Defining Multimodal Language Models

Multimodal language models signify a pivotal advancement in artificial intelligence, designed to analyze and integrate diverse data modalities such as text, images, audio, and video, enabling human-like communication and complex reasoning. These models utilize attention mechanisms, deep neural networks, and contrastive learning to enhance performance and adaptability. By combining visual analysis with natural language processing, they refine sentiment analysis, offering nuanced insights into user emotions from textual and visual cues. Large-scale datasets encompassing multiple languages and modalities facilitate advancements in tasks like named entity recognition (NER), where language transfer learning and feature sharing are crucial. Multimodal applications like meme generation demonstrate these models' ability to creatively synthesize text and images, capturing contemporary cultural trends in digital interactions [21, 5, 6, 22]. Addressing challenges in multimodal data integration, such as semantic alignment, these models enhance interpretability and robustness across applications.

A critical aspect of multimodal language models is their ability to merge structured and unstructured data, significantly enhancing classification and reasoning tasks. Integrating textual and visual inputs facilitates the generation of visually grounded paraphrases (VGPs), where different phrasal expressions convey the same visual concept [9]. Benchmarks like nvBench simulate translating natural language queries into visualizations, underscoring the importance of synthesizing linguistic and visual information [23]. In meme caption generation, models analyze single-image and multi-image formats to produce contextually relevant outputs [24]. This capability broadens the applicability of multimodal language models across domains.

Architecturally, multimodal language models employ single-stream and dual-stream designs to learn joint representations across modalities. Vision-language pre-training (VLP) integrates visual and linguistic modalities through extensive pre-training on large datasets, equipping models to effectively process multimodal inputs [25]. These architectures are instrumental in tasks like recognizing textual entailment, where models classify relationships between premises and hypotheses as entailment, contradiction, or neutral [26]. Their effectiveness in understanding complex relationships across modalities illustrates potential for real-world applications requiring accurate interpretation and synthesis of information.

Societal applications of multimodal language models are extensive and transformative across sectors. In healthcare, they enhance clinical diagnostics by integrating ophthalmic images with textual inquiries within Visual Question Answering (VQA) systems [27]. They tackle local challenges by analyzing news articles affecting users at neighborhood, city, or county levels [28]. In creative industries, these models facilitate meme caption generation, requiring sophisticated understanding of humor and context [24]. By synthesizing insights from various modalities and disciplines, these models pave the way for context-aware systems addressing complex challenges, such as enhancing long-form summarization of financial reports and improving cross-modal reasoning capabilities.

The integration of multimodal language models into diverse applications illustrates their transformative potential across domains. By addressing challenges related to data integration, semantic alignment, and contextual understanding, these models enhance human-machine interaction and foster interdisciplinary engagement, contributing to advances in technology and society. Ongoing research promises to refine multimodal language models' capabilities, ensuring relevance and effectiveness in future challenges.

2.2 Key Terms and Concepts

Exploring multimodal language models requires understanding foundational terms and concepts underpinning their development and application. Technologies driving advancements in human-machine interaction systems include large language models (LLMs), artificial intelligence (AI), natural language processing (NLP), machine learning (ML), deep learning (DL), and computational linguistics. Large pre-trained transformer-based models like BERT revolutionized NLP by enabling effective pre-training and fine-tuning for various tasks. LLMs excel in complex NLP tasks but face challenges in domain-specific analytical functions, as highlighted by the ResearchArena benchmark evaluating academic surveys through information discovery, selection, and organization [29, 30].

LLMs are sophisticated AI systems designed to process and generate human-like text using vast datasets and complex neural architectures. They excel in knowledge retrieval, summarization, and conversational AI but face challenges related to data privacy, interpretability, and computational efficiency. The imbalance of annotated data in supervised word sense disambiguation tasks exemplifies LLMs' limitations in handling underrepresented linguistic phenomena. Existing benchmarks focus on single language models or inadequately evaluate bilingual capabilities, highlighting a key obstacle in LLM development [31]. Addressing these limitations is essential for advancing LLMs, necessitating innovative approaches to model training and evaluation.

AI encompasses developing systems capable of tasks requiring human intelligence. In multimodal language models, AI integrates diverse data modalities, enabling machines to emulate complex reasoning and communication processes. A core issue in AI is the lack of standardized terminology and relationships, complicating communication and collaboration among researchers. AI's societal implications are profound, with transformative applications across sectors such as healthcare, finance, and education. However, advancements raise ethical concerns, including bias, privacy, and potential social inequalities, necessitating ongoing dialogue among stakeholders to establish responsible usage guidelines and policies [32, 33, 34, 35, 36].

NLP equips machines to understand, interpret, and generate human language using advanced computational techniques, particularly deep learning methods, which create hierarchical representations of language data. NLP encompasses tasks like parsing, part-of-speech tagging, machine translation, and dialogue systems. Neural network architectures have significantly improved language processing efficiency and accuracy, enabling machines to respond to textual and spoken inputs at speeds surpassing human capabilities [37, 38, 39, 40]. However, existing NLP datasets are often task-specific, limiting models' ability to learn task-invariant knowledge. Reinforcement learning (RL) in NLP tasks like syntactic parsing, language understanding, text generation, machine translation, and conversational systems highlights NLP methodologies' versatility and adaptability.

ML automates analytical model building, allowing systems to learn from data and improve over time. ML techniques are integral to multimodal language model development, enhancing training data and optimizing model performance. Successful AI systems rely on integrating machine learning algorithms, effective data generation strategies, and domain expertise, collectively enhancing processing and interpretation of complex information, ensuring ethical standards, and addressing challenges like bias and transparency across fields, including scientific research and education [20, 4, 39].

DL employs deep neural networks—multiple interconnected layers of neurons—to autonomously learn intricate patterns from raw data. This approach has demonstrated success across domains like speech recognition, computer vision, and NLP, modeling complex data relationships. DL's application to information retrieval explores neural embeddings to improve query and document representation. Despite challenges in achieving comparable performance in information retrieval, ongoing research indicates a promising trajectory for DL methodologies [41, 42, 38, 43]. DL techniques are central to multimodal language models' architecture, enabling feature extraction across modalities. Significant DL models in NLP include convolutional neural networks (CNNs), recurrent neural networks (RNNs), and recursive neural networks. Optimizer instability, such as Adam(W), can lead to training challenges, underscoring the need for robust optimization techniques.

Computational Linguistics studies language from a computational perspective, focusing on algorithms and models to process linguistic data. This field provides the theoretical foundation for multimodal language models, enabling tasks like hypothesis assessment and semantic coherence. Benchmarks assessing question-answering models by measuring incorrect predictions' proximity to correct answers refine error analysis and evaluate task difficulty. Metrics like Golden Rank (GR) and

Golden Rank Interpolated Median (GRIM) enhance understanding of model performance beyond traditional metrics like F1 and exact match scores. These benchmarks classify question difficulty and identify model error patterns, guiding improvements in model training and deployment in applications like customer service and academic research [29, 44, 45].

Integrating key terms and concepts in multimodal language model research advances NLP applications like keyword extraction, text generation, and sentiment analysis, improving human-machine interaction systems. This research leverages LLMs like GPT-3.5, Llama2-7B, and Falcon-7B to enhance tasks like information retrieval and document summarization, addressing challenges like model complexity and evaluation metrics. Innovative prompting strategies, such as Role-Playing and Chain-of-Thought, optimize LLM performance in nuanced tasks like sentiment analysis, driving innovation and effectiveness [46, 4, 47]. By addressing inefficiency, bias, and interpretability, researchers advance multimodal language models, fostering interdisciplinary collaboration and real-world impact.

3 Multimodal Language Models

Multimodal language models (MMLMs) represent a significant advancement in AI, integrating diverse data modalities to enhance machine comprehension and interaction capabilities. These models are crucial in practical applications across healthcare, education, and creative industries. As illustrated in Figure 2, the hierarchical structure of key concepts in the development and application of MMLMs encompasses methodologies for data integration, challenges in development, and diverse applications. The methodologies section highlights various strategies, including transformer architectures and neural network methodologies, while the challenges section addresses critical issues such as data alignment and interpretability. Furthermore, the applications section showcases the impact of MMLMs across different domains, including healthcare, education, and robotics, alongside the ongoing challenges faced. Table 1 presents a comparison of key methodologies used for integrating multimodal data in MMLMs, outlining their integration strategies, challenges, and applications to provide insight into their contributions to model development and performance. The following subsection explores methodologies for integrating multimodal data, addressing these challenges to enhance MMLMs' capabilities in solving complex real-world problems.

3.1 Methodologies for Multimodal Data Integration

Integrating diverse data modalities is vital for developing MMLMs, enabling machines to perform complex reasoning and emulate human communication. Researchers have devised strategies to address challenges like semantic alignment and cross-modal representation learning. Transformer architectures have revolutionized modality integration, enhancing model robustness and interpretability for tasks such as text classification and clinical concept extraction [48, 16, 18, 49, 4].

Transformer-based architectures, using attention mechanisms, effectively process heterogeneous inputs, as seen in the Multi-lingual Local News Classifier (MLNC) [28]. Neural network methodologies enhance language model performance by integrating multimodal data, improving feature extraction [1]. Encoder-decoder frameworks, like the XMeCap method, align image and caption features, demonstrating the potential of combining NLP and CV methodologies [24]. Task-specific methodologies, such as nvBench, highlight semantic connections for tasks like text-to-image generation [23].

These methodologies advance MMLMs by addressing semantic alignment, feature extraction, and contextual coherence, enhancing natural language processing tasks like text generation and document information extraction [50, 51, 4].

3.2 Challenges in Multimodal Language Model Development

Developing MMLMs involves overcoming challenges such as data alignment, computational complexity, and interpretability. Data alignment requires harmonizing diverse modalities into cohesive representations, complicated by semantic disparities and the scarcity of annotated datasets [52]. Computational complexity arises from resource-intensive models, necessitating more efficient frameworks [52]. Interpretability challenges stem from the complexity of transformer-based models, demanding transparency in decision-making processes [52].

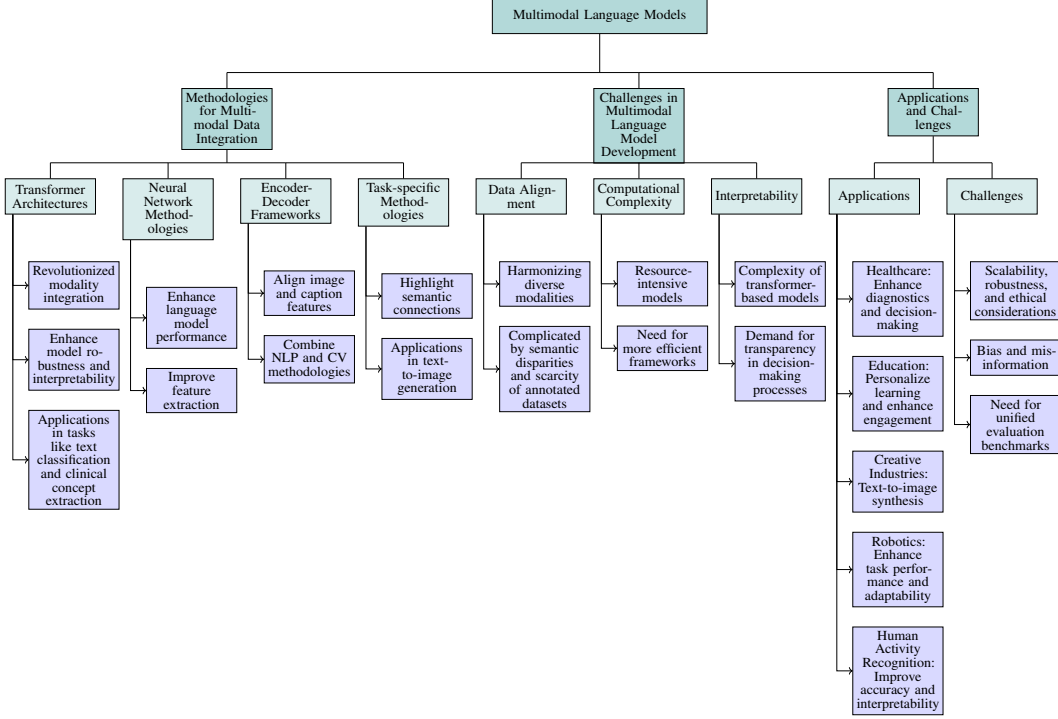


Figure 2: This figure illustrates the hierarchical structure of key concepts in the development and application of Multimodal Language Models (MMLMs), including methodologies for data integration, challenges in development, and diverse applications. The methodologies section highlights various strategies like transformer architectures and neural network methodologies, while the challenges section addresses issues like data alignment and interpretability. The applications section showcases the impact of MMLMs across different domains such as healthcare, education, and robotics, along with the ongoing challenges faced.

Addressing these challenges requires collaboration among academia, industry, and civil society to develop guidelines promoting transparency and accountability [36, 4]. Advanced methodologies for data alignment, scalable architectures for computational efficiency, and principled frameworks for interpretability are essential for robust and applicable MMLMs.

3.3 Applications and Challenges

MMLMs transform domains by integrating data modalities, enabling applications in healthcare, education, creative industries, robotics, and human activity recognition. In healthcare, MMLMs enhance diagnostics and decision-making by integrating multimodal data [53, 54, 55, 56]. In education, they personalize learning and enhance engagement [36, 57]. Creative industries leverage MMLMs for text-to-image synthesis, though biases in outputs remain a challenge [24]. In robotics, MMLMs enhance task performance and adaptability in dynamic environments [58, 14, 6, 8]. Human activity recognition benefits from integrating visual and textual data, improving accuracy and interpretability [59, 8].

As illustrated in Figure 3, the applications of MMLMs span diverse domains such as healthcare, education, robotics, and creative industries, while the challenges include scalability, ethical concerns, and robustness. Despite advancements, challenges in scalability, robustness, and ethical considerations persist [60, 36, 61, 4]. Techniques like RPC-Attention improve robustness, but unified evaluation benchmarks are needed [29, 62, 63, 4, 56]. Innovations like WaveletGPT enhance efficiency, making MMLMs more accessible [47]. Addressing ethical concerns, such as bias and misinformation, is crucial for responsible AI deployment [26].

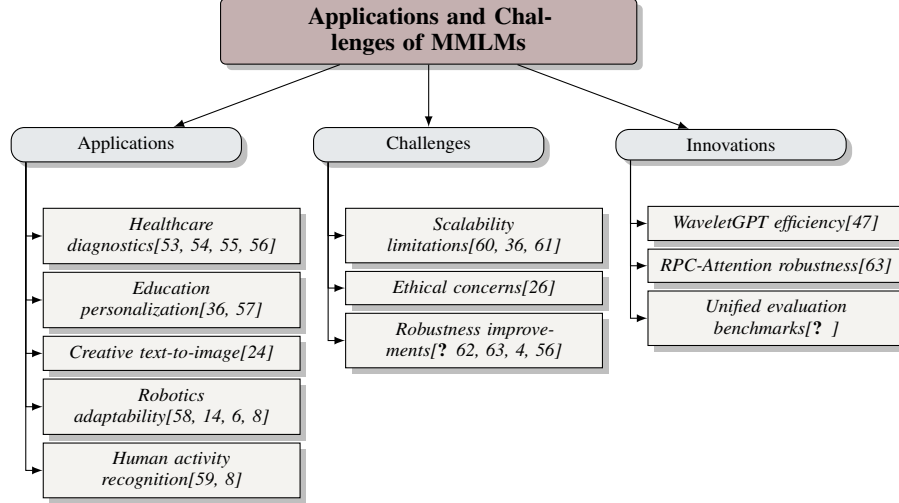


Figure 3: This figure illustrates the applications, challenges, and innovations related to Multimodal Large Language Models (MMLMs). The applications span diverse domains such as healthcare, education, robotics, and creative industries, while the challenges include scalability, ethical concerns, and robustness. Innovations like WaveletGPT and RPC-Attention address these challenges, enhancing the efficiency and evaluation of MMLMs.

Feature	Transformer Architectures	Neural Network Methodologies	Encoder-Decoder Frameworks
Integration Strategy	Attention Mechanisms	Feature Extraction	Image-caption Alignment
Challenges Addressed	Semantic Alignment	Cross-modal Representation	Contextual Coherence
Applications	Text Classification	Language Model Enhancement	Text Generation

Table 1: The table provides a comparative overview of methodologies employed in multimodal language models (MMLMs), focusing on transformer architectures, neural network methodologies, and encoder-decoder frameworks. It highlights their integration strategies, challenges addressed, and applications, emphasizing their role in advancing natural language processing tasks and multimodal data integration.

4 Large Language Models

In exploring the transformative impact of Large Language Models (LLMs) on various domains, it is essential to delve into their specific capabilities and applications. This examination reveals how LLMs have not only advanced natural language processing but have also facilitated innovative solutions across diverse fields. The following subsection will provide a detailed overview of the capabilities and applications of LLMs, highlighting their versatility and the significant contributions they have made in areas such as creative industries, healthcare, robotics, and multilingual tasks.

4.1 Capabilities and Applications

Large Language Models (LLMs) have revolutionized natural language processing (NLP) by demonstrating advanced capabilities across a wide range of applications, including text generation, classification, summarization, recommendation systems, conversational AI, and multimodal reasoning. These advanced models utilize intricate architectures and large-scale datasets to replicate human-like text comprehension and generation, leading to significant breakthroughs in various fields such as natural language processing, information systems, and text analytics. By addressing key challenges such as bias, reasoning, and coherence, these models not only enhance applications like summarization, translation, and question answering but also pave the way for innovative research avenues and improved methodologies in detecting machine-generated content, ultimately transforming business practices and societal interactions. [16, 4, 20]

The versatility of LLMs is particularly evident in the creative industries, where they have showcased their ability to generate coherent and stylistically nuanced outputs. For example, the XMeCap

framework employs supervised fine-tuning and reinforcement learning to incorporate both global and local similarities between visuals and text, significantly advancing tasks such as meme caption generation [24]. This highlights the adaptability of LLMs in generating creative and contextually appropriate outputs. Moreover, prompt engineering techniques have been utilized to guide LLMs in emulating specific language styles, thereby enhancing their utility in tasks requiring stylistic fidelity and creative expression. The integration of such techniques demonstrates the potential of LLMs to not only produce novel content but also to engage with specific cultural and contextual nuances, which is increasingly important in today's diverse media landscape.

In the healthcare sector, large language models (LLMs) have been increasingly integrated into diagnostic frameworks, demonstrating their potential to enhance clinical decision-making processes and improve patient outcomes. Recent studies indicate that adapted LLMs can outperform medical experts in summarizing clinical texts, such as radiology reports and patient communications, thereby alleviating the documentation burden on clinicians. Additionally, the development of benchmarks like CLIBENCH allows for a comprehensive evaluation of LLMs across a wide range of clinical tasks, including treatment procedure identification and lab test ordering. These advancements highlight the capacity of LLMs to support patient-specific decision-making in real-world clinical settings, ultimately aiming to increase the efficiency and accessibility of medical care [64, 56]. Their ability to process and synthesize multimodal data, including textual and visual inputs, has been instrumental in tasks such as clinical report generation and the classification of causes of death from verbal autopsy reports. By enabling nuanced reasoning across modalities, LLMs contribute to the development of robust AI-driven healthcare solutions.

The integration of LLMs into robotic systems has significantly enhanced human-robot interaction, particularly in complex, multi-step dialogue scenarios. For instance, embodied robotic systems leverage LLMs to interpret natural language instructions and respond effectively based on visual perception, exemplifying their transformative impact on multimodal reasoning and robotics. These advancements highlight the significant potential of Large Language Models (LLMs) to enhance task execution and adaptability in dynamic environments, particularly through their advanced natural language processing and data analysis capabilities. By leveraging iterative data augmentation strategies, such as LLM2LLM, these models can improve performance even in low-data scenarios, addressing challenges like bias and interpretability while fostering more efficient and transformative solutions [65, 36, 66].

Efficiency and scalability are pivotal challenges in the development of large language models (LLMs), yet innovative architectures like WaveletGPT are making significant strides in overcoming these obstacles. WaveletGPT employs wavelet-inspired techniques during pre-training, allowing it to leverage the multi-scale structure inherent in various data types, such as text, audio, and images. This approach enables the model to achieve nearly double the pre-training speed without adding any extra parameters to the existing GPT architecture, resulting in substantial performance gains comparable to those of larger models. Furthermore, it enhances the model's ability to process input across different temporal resolutions, thus improving its adaptability and efficiency. In addition, strategies like LLM2LLM provide targeted data augmentation, further enhancing performance in low-data scenarios, reducing reliance on extensive data curation, and facilitating the deployment of scalable LLM solutions. Together, these advancements pave the way for more efficient and effective LLMs in various applications [67, 68, 69, 66, 65]. This demonstrates the potential for creating more computationally efficient and accessible systems, thereby broadening the applicability of LLMs across diverse contexts.

The capabilities of LLMs extend significantly into multilingual applications, enhancing the ability to generate and understand diverse language content while also revolutionizing data augmentation strategies by providing innovative methods to create varied training examples that improve model performance without requiring additional data collection [66, 17]. Their ability to process and generate text in multiple languages has improved tasks such as local news classification and multilingual detection, enhancing precision and recall in linguistically diverse scenarios. These capabilities highlight the role of LLMs in fostering inclusivity and accessibility in NLP applications.

Through their diverse capabilities, LLMs have redefined the landscape of natural language processing, offering innovative solutions in creative generation, healthcare, robotics, multilingual tasks, and efficiency optimization. Their adaptability and advancements underscore their pivotal role in enhancing human-machine collaboration and addressing real-world challenges [24].

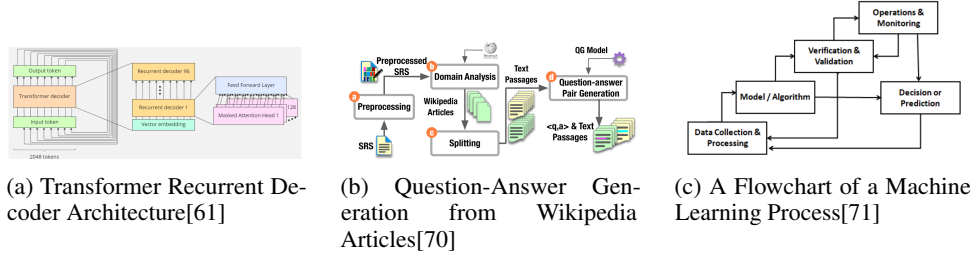


Figure 4: Examples of Capabilities and Applications

As shown in Figure 4, Large Language Models (LLMs) have revolutionized the field of artificial intelligence by offering advanced capabilities and a wide array of applications. The examples provided in the figure showcase the versatility and potential of these models across different domains. The first example illustrates a sophisticated Transformer Recurrent Decoder Architecture, which integrates multiple layers to efficiently process and generate tokens, highlighting the complexity and depth of modern LLMs. The second example demonstrates the application of LLMs in generating question-answer pairs from Wikipedia articles, showcasing their ability to comprehend and synthesize information from vast text sources. This process involves meticulous preprocessing and domain analysis, underscoring the model’s capability to handle structured data and derive meaningful insights. Lastly, the flowchart of a machine learning process encapsulates the comprehensive stages involved in developing and deploying LLMs, from data collection to operations. These examples collectively underscore the transformative impact of LLMs in enhancing computational processes, improving data analysis, and enabling innovative applications across various fields. [61, 70, 71]

4.2 Architectures and Training Processes

The architectures and training processes of large language models (LLMs) have seen substantial advancements, primarily through the development of transformer-based frameworks and innovative optimization techniques. Recent advancements in Large Language Models (LLMs) have significantly enhanced their performance in natural language processing (NLP) tasks by effectively capturing intricate contextual relationships and improving training efficiency. For instance, innovative methods such as retrieval-augmented generation (RAG) and iterative data augmentation strategies like LLM2LLM have demonstrated the ability to automate literature reviews and enhance model training, respectively. These techniques allow LLMs to better learn from complex datasets, even in low-data scenarios, thereby achieving higher accuracy in tasks such as document understanding and information extraction. Additionally, the development of tools like GigaCheck highlights the growing capability of LLMs to generate human-like text, while also addressing the need for effective detection methods to mitigate misuse. Overall, these advancements underscore the transformative impact of LLMs in automating and streamlining various NLP applications [69, 51, 37, 66, 65].

Transformer models, as detailed in foundational works, serve as the backbone of contemporary LLM architectures. These models utilize self-attention mechanisms to process input sequences in parallel, facilitating the efficient handling of long-range dependencies in text. For instance, the Double Feature Transformer introduces a modified architecture that enhances the model’s ability to process and integrate diverse data types, thereby improving performance in multimodal contexts [72]. The effectiveness of transformer models in handling multi-modal data is further exemplified in vision-language pre-training (VLP) methodologies, which employ both encoder-only and encoder-decoder frameworks to process and integrate visual and textual inputs [25].

Architectural innovations have been pivotal in refining the design of LLMs. The incorporation of discrete wavelet transforms into intermediate layers, as seen in WaveletGPT, allows for efficient multi-scale representation learning, enhancing the model’s ability to capture hierarchical structures in data [68]. Additionally, the RPC-Attention mechanism minimizes the impact of corrupted data by recovering principal components through a robust optimization framework, demonstrating its effectiveness in improving model robustness and interpretability [73].

Training processes have evolved to address challenges related to scalability and robustness. Techniques such as employing prompts to guide LLMs eliminate the need for extensive training datasets,

reducing costs and enhancing the personalization of digital assistants [74]. Furthermore, the use of task instructions to prompt models and evaluate their ability to generate correct responses without prior training on specific datasets underscores the role of task-specific fine-tuning in optimizing model outputs. The integration of quantum-classical algorithms, such as simulated annealing for generating contextually relevant sentences, offers an innovative approach to exploring sentence space and enhancing context-aware generation.

The continuous evolution of transformer-based architectures and training methodologies has positioned LLMs as transformative tools in NLP. Recent advancements in Large Language Models (LLMs) significantly enhance their ability to tackle intricate linguistic challenges, thereby driving innovation across a variety of applications, such as automated literature reviews, ethical AI use in research, and improved document understanding. These developments not only highlight the crucial role of LLMs in enhancing human-machine interaction but also raise important considerations regarding ethical implications, bias, and the need for responsible deployment. As LLMs continue to evolve, they promise to transform fields like education and research while necessitating ongoing dialogue among stakeholders to ensure their benefits are maximized and risks mitigated [69, 37, 36, 66].




Figure 5: Examples of Architectures and Training Processes

The insights provided in Figure 5 illustrate the intricate architectures and training processes that underpin large language models, shedding light on their operational mechanisms. The first example depicts a neural network model designed for generating context vectors from input sequences, which is a fundamental process in language models that transforms words or tokens into meaningful representations through attention mechanisms. The second image representing the memory processing cycle emphasizes cognitive processes that inspire neural network functionalities, highlighting stages such as sensory input and encoding, which parallel data processing in artificial intelligence. Lastly, the comparison of methods for semantic segmentation in images showcases the versatility of these models, demonstrating various approaches to segmenting fish and birds and underscoring the robustness of large language model architectures in handling complex tasks across different domains. Collectively, these examples provide a comprehensive view of the sophisticated frameworks and methodologies that drive advancements in LLMs, paving the way for innovative applications in artificial intelligence [19, 75, 76].

4.3 Optimization Strategies

The optimization of large language models (LLMs) is a crucial aspect of their development, aiming to enhance efficiency and reduce computational costs while maintaining high performance across diverse applications. Recent advancements in methodologies, hardware, and algorithms have led to the development of innovative techniques aimed at overcoming the computational challenges associated with Large Language Models (LLMs). For instance, new frameworks like GigaCheck enhance the detection of LLM-generated content through refined classification methods, while automated literature review systems utilize retrieval-augmented generation to efficiently process vast amounts of research data. Additionally, iterative data enhancement strategies, such as LLM2LLM, leverage teacher-student model dynamics to improve performance in low-data scenarios, demonstrating significant gains in various NLP tasks. These advancements collectively contribute to the ongoing evolution of LLM capabilities and their applications across diverse fields [37, 66, 65, 36, 77].

Transformer-based models have significantly benefited from optimization strategies that reduce redundancy among attention heads and improve representation capabilities. The Transformer-MGK model exemplifies these advancements by achieving lower computational costs and enhanced efficiency

in processing complex linguistic tasks [78]. These improvements are pivotal in enabling LLMs to handle large-scale data efficiently, thereby broadening their applicability in real-world scenarios.

Model-based approaches, such as LUNA, have introduced semantic binding and multiple quality metrics to differentiate themselves from existing analysis methods, offering a structured framework for optimizing LLMs. By incorporating these metrics, LUNA enhances the interpretability and efficiency of language models, making them more adaptable to diverse linguistic tasks [79]. This integration of semantic metrics is essential for refining model outputs and ensuring consistency in performance across applications.

Quantum-inspired models, like MSFF-QDConv, have demonstrated reduced parameter counts and computational complexity, offering improved accuracy in text classification tasks compared to classical models. These quantum models leverage multiscale feature fusion techniques to optimize LLMs, showcasing their potential in achieving high efficiency without compromising performance [80]. The adoption of quantum methodologies represents a promising direction for enhancing the scalability and efficiency of language models.

Prototypical network-based frameworks, such as Proto-lm, have achieved inherent interpretability in LLMs while maintaining competitive performance across various NLP tasks. By providing faithful explanations of model behavior, Proto-lm enhances the transparency and trustworthiness of AI systems, demonstrating that optimization can be achieved without sacrificing interpretability [81]. This approach is crucial for developing models that are both efficient and reliable in decision-making processes.

Despite advancements in efficiency, some models, like Stanza, prioritize accuracy, which can lead to longer processing times compared to other toolkits [40]. Addressing this trade-off between accuracy and computational efficiency remains a key focus in optimizing LLMs, as researchers strive to balance these aspects to enhance model utility in practical applications.

Deep reinforcement learning frameworks have also contributed to optimization strategies by offering comprehensive insights into model behavior, allowing researchers to identify patterns and trends in decision-making. This global view is essential for building trust in AI systems and ensuring that optimization efforts align with ethical and performance standards [82].

5 Artificial Intelligence and Machine Learning

Category	Feature	Method
Table 2: Table summarizing the methodologies and features associated with key algorithms in artificial intelligence and machine learning, specifically focusing on attention mechanisms, transfer learning, and embedding models. This table provides a comprehensive overview of the techniques that have significantly contributed to advancements in natural language processing and their applications across various domains.		

The convergence of artificial intelligence (AI) and machine learning (ML) has significantly advanced natural language processing (NLP). This section examines the foundational algorithms and techniques that have propelled intelligent language systems, focusing on attention mechanisms, transfer learning, and embedding models. Table 2 provides a detailed summary of the key methodologies and features associated with attention mechanisms, transfer learning, and embedding models, highlighting their transformative impact on natural language processing. These innovations not only enhance model performance and scalability but also have broad implications for domains such as healthcare, finance, and education. The next subsection delves into these key algorithms and techniques, illustrating their transformative impact on NLP.

5.1 Key Algorithms and Techniques

The evolution of intelligent language systems has been markedly influenced by attention mechanisms, transfer learning, and embedding models. Attention mechanisms, integral to transformer language models (TLMs) like BERT, have improved NLP tasks such as text analysis, translation, and dialogue systems by enabling models to focus on pertinent input data. Transfer learning facilitates the adaptation of pre-trained models to specific tasks, enhancing performance across applications and

enabling multilingual data processing. These methodologies allow systems to dynamically capture semantic relationships, process multimodal data, and tackle domain-specific challenges.

Attention mechanisms are crucial in modern language models, allowing systems to dynamically focus on relevant input elements. Their effectiveness is demonstrated in frameworks like RadCLIP, which align multimodal inputs such as radiologic images and text, enhancing interpretability in complex reasoning tasks. This versatility extends to real-time applications, such as COVID-19 misinformation detection, where attention mechanisms capture nuanced data relationships. Kernel-based frameworks further highlight attention mechanisms' role in linking functional analysis with transformer architectures, providing robust theoretical foundations.

Transfer learning has revolutionized NLP by enabling models to generalize across tasks and domains. Pretraining large models and fine-tuning them on specific tasks have significantly improved performance, as seen in legal text processing, where domain-specific adaptations are essential. Insights from neural collapse suggest a principled, parameter-efficient fine-tuning method, enhancing adaptability with minimal training. This is particularly valuable in rapidly evolving fields requiring timely information processing, shaping more sophisticated and responsive language models.

Embedding models provide structured linguistic data representations, facilitating task-specific adaptations. The proposed framework based on sufficiency and informativeness offers a task-agnostic evaluation metric, underscoring the need for robust ranking mechanisms. Metrics like GR and GRIM offer granular views of model performance, improving evaluation accuracy. Advancements in embedding techniques deepen language semantics understanding and improve models' ability to handle complex linguistic phenomena, fostering a comprehensive approach to language processing.

In entailment tasks, models like e-INFERSENT achieve competitive accuracy while generating coherent explanations, integrating explanations without sacrificing performance. This underscores algorithmic innovations' importance in advancing intelligent language systems. By integrating attention mechanisms, transfer learning, and embedding models, researchers enhance NLP systems' performance and scalability, enabling applications across healthcare, misinformation detection, and legal analysis. These algorithms drive innovation and interdisciplinary collaboration, leading to more effective data processing and insights.

5.2 Addressing Bias and Ethical Concerns

AI-driven language models must address biases and ethical considerations to ensure trust and reliability. Bias in training data and algorithms can perpetuate stereotypes and lead to unfair outcomes. For instance, biases in language identification tools necessitate more inclusive NLP models that account for linguistic diversity. Fairness-aware algorithms are essential to ensure equitable outcomes across domains. However, the lack of comprehensive datasets and standardized evaluation metrics complicates bias mitigation, highlighting the need for ongoing research to develop effective assessment methodologies.

Transparency and interpretability are critical for ethical AI systems, enabling stakeholders to understand decision-making processes. This is crucial for addressing bias, misinformation, and ethical implications in research and education. Discussions among stakeholders are necessary to establish guidelines for interpretability methods, ensuring responsible AI use. The black-box nature of deep learning models, particularly Multi-Layer Perceptrons (MLPs), challenges explainability. Developing explainable fact-checking methods that provide understandable explanations is crucial for enhancing trust in AI systems. Ethical considerations surrounding automated UI agents highlight potential misuse risks, especially for marginalized communities, emphasizing responsible AI development.

Ethical concerns in AI and LLMs include data privacy risks and misinformation spread. Robust policies and regulations prioritizing transparency and accountability are necessary. Systems like ChatGPT raise concerns over hallucinated outputs, complicating fairness efforts. Addressing these issues is critical in applications like autonomous robots, where data privacy and safety are paramount.

Bias mitigation in NLP involves balancing bias reduction with model performance. Addressing biases such as gender bias can lead to fairer outcomes but may compromise accuracy and generalizability. Debiasing methods often have trade-offs, necessitating careful consideration of model effectiveness. The computational costs of training large models and the need for extensive labeled datasets further

constrain bias mitigation. A standardized AI framework can enhance collaboration and integrate bias mitigation methodologies.

Explainability aligns AI systems with ethical standards. Humor generation systems using models like GPT-3 benefit from explainable mechanisms aligning outputs with human comedic expectations. Modeling human comedy theories and employing cognitive distance strategies enhance user understanding. Humor style classification and computational techniques like incongruity detection and sentiment analysis improve humor resonance. This bridges the gap between machine-generated content and human comprehension, addressing humor generation challenges. Non-parametric approaches in few-shot learning aid in identifying and mitigating task-specific biases.

Addressing these challenges enhances AI-driven language models' ethical deployment, ensuring fair, reliable, and beneficial development. Bias mitigation, transparency, data accountability, privacy protection, and explainability are critical for fostering trust and aligning AI systems with societal values.

6 Multimodal Data Integration

The integration of multimodal data is a transformative area in artificial intelligence, enhancing language models by combining text, images, audio, and video. This capability not only improves model performance but also addresses challenges in processing diverse data types. Multimodal integration is crucial for applications ranging from natural language processing to robotics, improving robustness and contextual understanding.

6.1 Importance of Multimodal Data Integration

Multimodal data integration enhances the semantic richness and task-specific performance of language models by synthesizing complementary information from text, images, audio, and video. This approach addresses semantic alignment, contextual coherence, and robustness, which are vital for AI systems. Visual data integration, for instance, enhances Visual Question Answering (VQA) systems in ophthalmology, leading to more precise responses [27]. Similarly, the WALL-E method combines language and visual inputs to improve robotic system interactions [11].

In dialogue systems, Emotional Attribution Encoding (EAE) modules, such as HCAN, capture emotional attributions using IA-attention mechanisms, surpassing traditional methods in understanding conversational contexts [10]. Integration of visual context with language processing is crucial for understanding ambiguous instructions [83]. Additionally, weak supervision and advanced NLP techniques improve news content understanding, enhancing local relevance identification [28].

Multimodal integration also contributes to the robustness and adaptability of NLP systems, addressing challenges with word and document embeddings [1]. Techniques extracting visually grounded phrases enrich semantic understanding [9]. In NL2VIS, nvBench provides a dataset that strengthens natural language and visualization task integration [23].

Moreover, integrating diverse data types enhances AI interpretability and trustworthiness. Benchmarks like SNLI promote interpretable models with human-understandable justifications [26]. Continuous exploration of multimodal integration techniques is essential for addressing complex challenges and unlocking transformative potential across domains.

6.2 Strategies for Combining Multimodal Data

Advanced methodologies and frameworks enable effective multimodal data integration, processing diverse data types into unified representations. These strategies support applications like multimodal sentiment analysis, where neural networks combine visual and textual features to infer emotions from social media posts. They also facilitate long-form summarization of complex documents like financial reports, addressing challenges such as numeric hallucination and position bias [6, 8].

Structured Optimal Transport (SOT) enhances multimodal integration by optimizing a structured cost function to reflect dependencies among mappings [84]. This method improves semantic coherence and interpretability. Task-Agnostic Dialect Adapters (TADA) use contrastive and morphosyntactic losses to adapt dialects, enhancing model adaptability in diverse linguistic contexts [85]. TaskMatrix.AI's

architecture, with its modular components, facilitates seamless multimodal data integration for complex tasks [86].

TextConvNet uses convolutional neural networks (CNNs) for feature extraction, capturing spatial relationships within multimodal data [87]. This is beneficial for tasks requiring spatial and temporal information integration. The Topic Segmentation and Labeling Framework uses graph-based methods to process asynchronous data streams, enhancing contextual coherence [88]. The Blender framework integrates features from different modalities, maintaining task-specific performance [89].

The iParaphrasing method involves clustering algorithms for multimodal data integration, emphasizing visual grounding [9]. Future research could explore kernel applicability and Transformer architecture relationships [90]. These strategies enhance multimodal data integration, improving applications in text-to-image and text-to-video generation, visual question answering, and image captioning [15, 9, 58, 91, 4].

6.3 Challenges in Multimodal Data Integration

Multimodal data integration presents technical and practical challenges, including scalability, robustness, and domain applicability. Controllable data augmentation, model interaction complexity, and resource demands for optimizing large language models (LLMs) are significant hurdles. Hallucination in LLMs complicates result evaluation, necessitating refined methodologies for effective multimodal applications [47, 17]. Challenges include data synchronization, scalability, interpretability, alignment complexities, and memory management.

Data synchronization involves aligning heterogeneous modalities into coherent representations, a critical challenge due to semantic disparities. Insufficient multimodal datasets exacerbate this issue, especially in tasks like video summarization [14]. Effective utilization of knowledge beyond training datasets remains a barrier, as models struggle to generalize contextually [15].

Scalability challenges arise from the computational complexity of processing large-scale multimodal datasets, requiring substantial resources [15]. Developing scalable architectures and optimized training strategies is crucial for broader applicability.

Interpretability is hindered by the complexity of underlying architectures and variability in embedding strategies, complicating model behavior interpretation. Enhancing interpretability requires principled frameworks for transparency in decision-making processes [11].

Alignment complexities involve achieving semantic coherence across diverse modalities, necessitating domain-specific knowledge and advanced methodologies [15]. Memory management and multi-user understanding introduce practical challenges in real-world applications, such as embodied robotic systems [11]. Addressing these challenges requires robust memory architectures and adaptive mechanisms for effective operation in complex environments.

7 Deep Learning in Language Models

In deep learning, neural network architecture is crucial for the effectiveness of language models. This section explores key aspects of architecture design and efficiency, examining how different frameworks such as transformers, recursive autoencoders, and convolutional networks influence model performance and interpretability. By analyzing these architectures, we aim to elucidate their operational mechanisms and implications in natural language processing.

7.1 Neural Network Architectures and Efficiency

Neural network architecture significantly impacts language model accuracy and efficiency. Transformers, with their self-attention mechanisms, revolutionize natural language processing by capturing long-range dependencies and enhancing computational efficiency and scalability. This is particularly evident in tasks like Named Entity Recognition (NER), where fine-tuned transformers outperform traditional models, especially in low-resource languages, and improve outcomes through document structure understanding [92, 93]. Recursive Autoencoders (RAEs) enhance interpretability by focusing on structure and performance, aligning with the goal of developing transparent models for complex linguistic patterns [94]. Proto-lm, integrating interpretable prototypes, improves model

transparency and performance, highlighting the importance of incorporating interpretability into neural network design [81].

Convolutional neural networks (CNNs) like TextConvoNet excel in text classification by extracting inter-sentence features using two-dimensional convolutional filters, demonstrating CNNs' adaptability in processing diverse data types [87]. Methods minimizing sentence vector distances of paraphrases enhance semantic coherence, contributing to improved linguistic understanding [95]. The continuous evolution of architectures such as transformers, RAEs, prototypical networks, and CNNs underscores their impact on language model accuracy and efficiency, driving advancements in natural language processing across diverse applications [1].

7.2 Attention Mechanisms and Interpretability

Attention mechanisms enhance language model performance and interpretability by focusing on relevant input data portions, facilitating nuanced reasoning and transparent decision-making processes. These mechanisms address semantic alignment and contextual understanding challenges, enabling interpretable outputs while maintaining high performance [52]. Attention weights offer insights into input component importance, providing a structured approach to analyzing model behavior compared to traditional parameters [96]. Frameworks like Latent Attention Network (LAN) further enhance this analysis by generating attention masks to measure input component replaceability [97].

Innovative methodologies improve interpretability through attention mechanisms, such as iAdvT-Text generating adversarial perturbations linked to interpretable linguistic changes [98], and the Interpretation Quality Score (IQS) providing a standardized framework for evaluating interpretability methods [99]. Integrating attention with symbolic knowledge could enhance interpretability by combining neural mechanisms with human-understandable reasoning processes [19]. Tools like Inseq facilitate accessible interpretability analyses, promoting the adoption of interpretability practices [100].

Attention mechanisms also enhance reinforcement learning model interpretability, as demonstrated by the Adversarial Inverse Reinforcement Learning (AIRL) framework, which generates discriminators capturing decision-making patterns [82]. This versatility across machine learning paradigms underscores attention mechanisms' vital role in bridging model performance and interpretability, shaping the future of language models and their applications.

8 Computational Linguistics

The intersection of computational linguistics and its applications provides valuable insights into domains such as hate speech detection. As the field evolves, leveraging linguistic principles to enhance detection methodologies becomes increasingly critical. Integrating linguistic insights into computational models enables researchers to devise more nuanced strategies for identifying and mitigating hate speech. This section examines the role of linguistic knowledge in advancing hate speech detection, emphasizing the challenges and opportunities arising from this integration.

8.1 Linguistic Knowledge in Hate Speech Detection

Computational linguistics plays a pivotal role in hate speech detection by addressing biases and promoting fairness in language models. However, the lack of a universally accepted definition of hate speech creates inconsistencies in data annotation and research outcomes, complicating efforts to standardize detection methodologies across languages and cultural contexts [101]. This ambiguity affects dataset labeling and impacts the training of machine learning models, which rely on high-quality annotations. Developing linguistic frameworks to clarify and unify hate speech definitions is essential for advancing this research area.

Feature-attribution methods have been instrumental in providing insights into the importance of linguistic features in hate speech detection. However, existing benchmarks, such as the Interpretation Quality Score (IQS), are limited in their applicability to non-feature-attribution approaches, necessitating the development of more comprehensive evaluation frameworks. Such frameworks would enable more accurate assessments of model performance and facilitate comparisons across methodologies, thereby advancing the field [99].

Incorporating linguistic knowledge into model development can address challenges such as bias, hallucinations, and inconsistencies in text generation tasks [102, 4]. By analyzing syntactic and semantic patterns, researchers can enhance model fairness and accuracy. For instance, accounting for dialectal variations and contextual nuances can mitigate biases and ensure equitable treatment of diverse linguistic expressions. Furthermore, task-specific benchmarks and annotation guidelines informed by linguistic expertise can improve the consistency and reliability of hate speech detection systems.

8.2 Semantic Coherence and Sentence Embeddings

Semantic coherence and sentence embeddings are fundamental to advancing language models, as they enable contextually relevant and meaningful text understanding and generation. Sentence embeddings provide compact, informative representations of sentences, capturing both syntactic and semantic nuances. These embeddings enhance performance across tasks such as sentiment analysis, machine translation, and information retrieval. Recent innovations, such as transition matrices, refine embeddings to better capture latent semantic meanings, improving tasks like sentence classification and summarization. Meta-embeddings, which combine multiple pre-trained word embeddings, address challenges across NLP applications, including text mining and question-answering, by leveraging the strengths of individual representations [95, 103, 104].

Semantic coherence significantly improves the interpretability and reliability of language models. Ensuring logical and contextually appropriate sequences in generated text enhances alignment with human communication. This is particularly critical in applications like cognitive impairment identification, where linguistic context improves diagnostic accuracy [105]. Maintaining semantic coherence not only enhances interpretability but also ensures user interactions with language technologies are more effective and aligned with expectations.

The development of sentence embeddings has also advanced semantic coherence by representing sentences in continuous vector spaces, capturing intricate relationships between words and phrases. This capability enhances tasks such as information retrieval, text classification, and semantic similarity by enabling nuanced connections between terms, even when distantly related in text [1, 103, 48, 104]. Tasks like paraphrase detection and text summarization benefit from these embeddings, as they allow models to achieve higher semantic coherence, leading to more accurate and interpretable outputs. Continued refinement of sentence embeddings will play a critical role in the future development of advanced NLP applications.

9 Conclusion

9.1 Future Directions and Emerging Trends

The trajectory of multimodal language models (MMLMs) is set to be influenced by advancements in architectural optimization, interdisciplinary approaches, and ethical frameworks. As artificial intelligence progresses, optimizing model architectures and training methodologies becomes crucial for enhancing scalability and applicability across various domains. This focus will enable faster processing and the handling of complex tasks, broadening the utility of MMLMs in real-world applications.

Emerging trends point to an increasing focus on multi-task learning frameworks, particularly for visually grounded paraphrase extraction. Extending this approach to other datasets could expand the scope of MMLMs, fostering more robust models capable of addressing complex multimodal tasks. By integrating diverse learning objectives, researchers can create models that are efficient and adaptable to a wide array of inputs and contexts.

In sentiment analysis and creative generation, future research should aim to improve model performance in positive sentiment categories and explore the cross-cultural applicability of humor in generated content. This could lead to culturally sensitive models, expanding the impact of MMLMs in creative industries. Understanding cultural nuances can enable these models to resonate with diverse audiences, increasing their relevance and effectiveness.

Developing advanced models that leverage explanations for improved performance is a promising research avenue. Investigating the role of explanations in natural language understanding tasks

could enhance the interpretability and efficacy of language models, improving their utility in diverse applications. Efforts to enhance explainability will bolster user trust and facilitate the integration of MMLMs into critical domains like healthcare and education.

Expanding existing benchmarks to include more tasks, visualization types, and support for conversational and underspecified queries is crucial for advancing MMLM capabilities. Such expansions will drive innovation in multimodal data integration and processing, ensuring rigorous evaluation and continuous improvement of MMLMs. Future research should focus on refining transfer learning metrics, exploring interdisciplinary collaborations, and addressing ethical considerations to ensure responsible MMLM deployment.

9.2 Explainability and Evaluation Metrics

Benchmark	Size	Domain	Task Format	Metric
MMNERD[22]	42,908	Named Entity Recognition	Named Entity Recognition	F1-score
TRAVEL-BM[106]	10,500	Travel	Question Answering	E2E, Ragas
ICMS[107]	51,906	Construction Cost Management	Multi-class Text Classification	Accuracy, F1-score
SDAAP[77]	20,000	Spectral Analysis	Question Answering	BLEU, ROUGE
CLLMS[56]	198,000	Clinical Text Summarization	Summarization	BLEU, ROUGE-L
AraSum[108]	4,000	Clinical Documentation	Summarization	F1 Score, PDQL-9
Bios[109]	400,000	Text Classification	Occupation Prediction	GP, TPR
SciTLDR[37]	5,400	Literature Review	Summarization	ROUGE-1, ROUGE-2

Table 3: Overview of representative benchmarks used for evaluating language models across diverse domains, task formats, and evaluation metrics. The table highlights benchmark sizes, associated tasks, and commonly used performance metrics, providing a comprehensive summary of resources for model assessment and comparison.

Explainability and robust evaluation metrics are essential for the development and deployment of language models, ensuring reliability, transparency, and alignment with societal values. As language models advance, improved interpretability becomes critical, particularly for conversational agents and educational assessments. The role of AI in enhancing conversational capabilities highlights the need for models that are understandable and trusted by users, fostering engagement and effective AI deployment.

Addressing multilingual performance challenges, especially in low-resource languages, requires further research to enhance model efficacy across diverse linguistic contexts. This involves adopting standardized benchmarks and evaluation protocols for reliable model assessment. Table 3 presents a detailed summary of representative benchmarks that are critical for evaluating language model performance across various domains and task formats. Such benchmarks ensure integrity in various applications, including educational assessments, allowing for effective performance comparison and improvement identification.

Improving detector interpretability and developing robust evaluation metrics are crucial for assessing language model performance, particularly in tasks involving automatic detection of machine-generated content. Integrating human-in-the-loop approaches can enhance bias mitigation, promoting transparency and collaboration in NLP model development. This aligns with the need for transparency in AI systems, ensuring adherence to ethical standards and societal expectations.

The NLP community is encouraged to adopt diverse statistical methods, including Bayesian techniques, to improve empirical claim rigor and clarity. This enhances evaluation metric robustness, providing a comprehensive framework for assessing model performance and reliability. By diversifying methodological approaches, researchers can gain deeper insights into model strengths and weaknesses, leading to more effective and trustworthy language technologies.

References

- [1] Siwei Lai. Word and document embeddings based on neural network approaches, 2016.
- [2] Sagar Gubbi Venkatesh, Partha Talukdar, and Srini Narayanan. Ugif: Ui grounded instruction following, 2023.
- [3] Tiantian Feng, Digbalay Bose, Tuo Zhang, Rajat Hebbar, Anil Ramakrishna, Rahul Gupta, Mi Zhang, Salman Avestimehr, and Shrikanth Narayanan. Fedmultimodal: A benchmark for multimodal federated learning, 2023.
- [4] Jonas Becker, Jan Philip Wahle, Bela Gipp, and Terry Ruas. Text generation: A systematic literature review of tasks, evaluation, and challenges, 2024.
- [5] Zhiyuan Liu, Chuanzheng Sun, Yuxin Jiang, Shiqi Jiang, and Mei Ming. Multi-modal application: Image memes generation, 2021.
- [6] Anthony Hu and Seth Flaxman. Multimodal sentiment analysis to explore the structure of emotions, 2018.
- [7] Dainis Boumber, Rakesh M. Verma, and Fatima Zahra Qachfar. A roadmap for multilingual, multimodal domain independent deception detection, 2024.
- [8] Tianyu Cao, Natraj Raman, Danial Dervovic, and Chenhao Tan. Characterizing multimodal long-form summarization: A case study on financial reports, 2024.
- [9] Chenhui Chu, Mayu Otani, and Yuta Nakashima. iparaphrasing: Extracting visually grounded paraphrases via an image, 2018.
- [10] Shanglin Lei, Xiaoping Wang, Guanting Dong, Jiang Li, and Yingjian Liu. Watch the speakers: A hybrid continuous attribution network for emotion recognition in conversation with emotion disentanglement, 2023.
- [11] Tianyu Wang, Yifan Li, Haitao Lin, Xiangyang Xue, and Yanwei Fu. Wall-e: Embodied robotic waiter load lifting with large language model, 2023.
- [12] Jan Philip Wahle, Terry Ruas, Mohamed Abdalla, Bela Gipp, and Saif M. Mohammad. We are who we cite: Bridges of influence between natural language processing and other academic fields, 2024.
- [13] Yunfan Jiang, Agrim Gupta, Zichen Zhang, Guanzhi Wang, Yongqiang Dou, Yanjun Chen, Li Fei-Fei, Anima Anandkumar, Yuke Zhu, and Linxi Fan. Vima: General robot manipulation with multimodal prompts, 2023.
- [14] Toqa Alaa, Ahmad Mongy, Assem Bakr, Mariam Diab, and Walid Gomaa. Video summarization techniques: A comprehensive review, 2024.
- [15] Feng Li, Hao Zhang, Yi-Fan Zhang, Shilong Liu, Jian Guo, Lionel M. Ni, PengChuan Zhang, and Lei Zhang. Vision-language intelligence: Tasks, representation learning, and large models, 2022.
- [16] Ross Gruetzemacher and David Paradise. Deep transfer learning beyond: Transformer language models in information systems research, 2021.
- [17] Bosheng Ding, Chengwei Qin, Ruochen Zhao, Tianze Luo, Xinze Li, Guizhen Chen, Wenhan Xia, Junjie Hu, Anh Tuan Luu, and Shafiq Joty. Data augmentation using large language models: Data perspectives, learning paradigms and challenges, 2024.
- [18] Luran Wang, Mark Gales, and Vatsal Raina. An information-theoretic approach to analyze nlp classification tasks, 2024.
- [19] Andrea Galassi, Marco Lippi, and Paolo Torrioni. Attention in natural language processing, 2021.
- [20] Ganesh Jawahar, Muhammad Abdul-Mageed, and Laks V. S. Lakshmanan. Automatic detection of machine generated text: A critical survey, 2020.

-
- [21] Dan Sun, Yaxin Liang, Yining Yang, Yuhua Ma, Qishi Zhan, and Erdi Gao. Research on optimization of natural language processing model based on multimodal deep learning, 2024.
 - [22] Dongsheng Wang, Xiaoqin Feng, Zeming Liu, and Chuan Wang. 2m-ner: Contrastive learning for multilingual and multimodal ner with language and modal fusion, 2024.
 - [23] Yuyu Luo, Jiawei Tang, and Guoliang Li. nvbench: A large-scale synthesized dataset for cross-domain natural language to visualization task, 2021.
 - [24] Yuyan Chen, Songzhou Yan, Zhihong Zhu, Zhixu Li, and Yanghua Xiao. Xmecap: Meme caption generation with sub-image adaptability, 2024.
 - [25] Feilong Chen, Duzhen Zhang, Minglun Han, Xiuyi Chen, Jing Shi, Shuang Xu, and Bo Xu. Vlp: A survey on vision-language pre-training, 2022.
 - [26] Oana-Maria Camburu, Tim Rocktäschel, Thomas Lukasiewicz, and Phil Blunsom. e-snli: Natural language inference with natural language explanations. *Advances in Neural Information Processing Systems*, 31, 2018.
 - [27] Xiaolan Chen, Ruoyu Chen, Pusheng Xu, Weiye Zhang, Xianwen Shang, Mingguang He, and Danli Shi. Visual question answering in ophthalmology: A progressive and practical perspective, 2024.
 - [28] Deven Santosh Shah, Shiyang He, Gosuddin Kamaruddin Siddiqi, and Radhika Bansal. What’s happening in your neighborhood? a weakly supervised approach to detect local news, 2024.
 - [29] Hao Kang and Chenyan Xiong. Researcharena: Benchmarking large language models’ ability to collect and organize information as research agents, 2025.
 - [30] Bonan Min, Hayley Ross, Elmor Sulem, Amir Pouran Ben Veyseh, Thien Huu Nguyen, Oscar Sainz, Eneko Agirre, Ilana Heintz, and Dan Roth. Recent advances in natural language processing via large pre-trained language models: A survey. *ACM Computing Surveys*, 56(2):1–40, 2023.
 - [31] David M. W. Powers. What the f-measure doesn’t measure: Features, flaws, fallacies and fixes, 2019.
 - [32] Yongjun Zhang. Generative ai has lowered the barriers to computational social sciences, 2025.
 - [33] Subhash Nerella, Sabyasachi Bandyopadhyay, Jiaqing Zhang, Miguel Contreras, Scott Siegel, Aysegul Bumin, Brandon Silva, Jessica Sena, Benjamin Shickel, Azra Bihorac, Kia Khezeli, and Parisa Rashidi. Transformers in healthcare: A survey, 2023.
 - [34] Tianqi Shang, Shu Yang, Weiqing He, Tianhua Zhai, Dawei Li, Bojian Hou, Tianlong Chen, Jason H. Moore, Marylyn D. Ritchie, and Li Shen. Leveraging social determinants of health in alzheimer’s research using llm-augmented literature mining and knowledge graphs, 2025.
 - [35] Maria T. Baldassarre, Danilo Caivano, Berenice Fernandez Nieto, Domenico Gigante, and Azzurra Ragone. The social impact of generative ai: An analysis on chatgpt, 2024.
 - [36] Ahmed S. BaHammam, Khaled Trabelsi, Seithikurippu R. Pandi-Perumal, and Hiatham Jahrami. Adapting to the impact of ai in scientific writing: Balancing benefits and drawbacks while developing policies and regulations, 2023.
 - [37] Nurshat Fateh Ali, Md. Mahdi Mohtasim, Shakil Mosharrof, and T. Gopi Krishna. Automated literature review using nlp techniques and llm-based retrieval-augmented generation, 2024.
 - [38] Tom Young, Devamanyu Hazarika, Soujanya Poria, and Erik Cambria. Recent trends in deep learning based natural language processing. *ieee Computational intelligence magazine*, 13(3):55–75, 2018.
 - [39] Akshansh Mishra, Vijaykumar S Jatti, Vaishnavi More, Anish Dasgupta, Devarrishi Dixit, and Eyob Messele Sefene. Performance prediction of data-driven knowledge summarization of high entropy alloys (heas) literature implementing natural language processing algorithms, 2023.

-
- [40] Peng Qi, Yuhao Zhang, Yuhui Zhang, Jason Bolton, and Christopher D Manning. Stanza: A python natural language processing toolkit for many human languages. *arXiv preprint arXiv:2003.07082*, 2020.
 - [41] Mohd Halim Mohd Noor and Ayokunle Olalekan Ige. A survey on state-of-the-art deep learning applications and challenges, 2025.
 - [42] Da Song, Zhijie Wang, Yuheng Huang, Lei Ma, and Tianyi Zhang. Deeplens: Interactive out-of-distribution data detection in nlp models, 2023.
 - [43] Ye Zhang, Md Mustafizur Rahman, Alex Braylan, Brandon Dang, Heng-Lu Chang, Henna Kim, Quinten McNamara, Aaron Angert, Edward Banner, Vivek Khetan, Tyler McDonnell, An Thanh Nguyen, Dan Xu, Byron C. Wallace, and Matthew Lease. Neural information retrieval: A literature review, 2017.
 - [44] Shashidhar Reddy Javaji, Haoran Hu, Sai Sameer Vennam, and Vijaya Gajanan Buddhavarapu. A comparative and experimental study on automatic question answering systems and its robustness against word jumbling, 2023.
 - [45] Michael Kamfonas and Gabriel Alon. What can secondary predictions tell us? an exploration on question-answering with squad-v2.0, 2022.
 - [46] Yajing Wang and Zongwei Luo. Enhance multi-domain sentiment analysis of review texts through prompting strategies, 2024.
 - [47] Sandeep Chataut, Tuyen Do, Bichar Dip Shrestha Gurung, Shiva Aryal, Anup Khanal, Carol Lushbough, and Etienne Gnimpieba. Comparative study of domain driven terms extraction using large language models, 2024.
 - [48] Alfredo Silva and Marcelo Mendoza. A data-driven strategy to combine word embeddings in information retrieval, 2021.
 - [49] Yuqi Si, Jingqi Wang, Hua Xu, and Kirk Roberts. Enhancing clinical concept extraction with contextual embeddings, 2019.
 - [50] William Hogan. An overview of distant supervision for relation extraction with a focus on denoising and pre-training methods, 2022.
 - [51] Jiabang He, Lei Wang, Yi Hu, Ning Liu, Hui Liu, Xing Xu, and Heng Tao Shen. Icl-d3ie: In-context learning with diverse demonstrations updating for document information extraction, 2023.
 - [52] Ting-Yun Chang, Ta-Chung Chi, Shang-Chi Tsai, and Yun-Nung Chen. xsense: Learning sense-separated sparse representations and textual definitions for explainable word sense networks, 2018.
 - [53] Shibin Wu, Bang Yang, Zhiyu Ye, Haoqian Wang, Hairong Zheng, and Tong Zhang. Improving medical report generation with adapter tuning and knowledge enhancement in vision-language foundation models, 2023.
 - [54] Cheng Su, Jinbo Wen, Jiawen Kang, Yonghua Wang, Yuanjia Su, Hudan Pan, Zishao Zhong, and M. Shamim Hossain. Hybrid rag-empowered multi-modal llm for secure data management in internet of medical things: A diffusion-based contract approach, 2024.
 - [55] Yuting He, Fuxiang Huang, Xinrui Jiang, Yuxiang Nie, Minghao Wang, Jiguang Wang, and Hao Chen. Foundation model for advancing healthcare: Challenges, opportunities, and future directions, 2024.
 - [56] Dave Van Veen, Cara Van Uden, Louis Blankemeier, Jean-Benoit Delbrouck, Asad Aali, Christian Bluethgen, Anuj Pareek, Malgorzata Polacin, Eduardo Pontes Reis, Anna Seehofnerova, Nidhi Rohatgi, Poonam Hosamani, William Collins, Neera Ahuja, Curtis P. Langlotz, Jason Hom, Sergios Gatidis, John Pauly, and Akshay S. Chaudhari. Adapted large language models can outperform medical experts in clinical text summarization, 2024.

-
- [57] Baihan Lin. Reinforcement learning and bandits for speech and language processing: Tutorial, review and outlook, 2023.
- [58] Aditi Singh. A survey of ai text-to-image and ai text-to-video generators, 2023.
- [59] Hrishikesh Singh, Aarti Sharma, and Millie Pant. Pixels to prose: Understanding the art of image captioning, 2024.
- [60] Sukhpal Singh Gill and Rupinder Kaur. Chatgpt: Vision and challenges, 2023.
- [61] Frank Joublin, Antonello Ceravola, Joerg Deigmoeller, Michael Gienger, Mathias Franzius, and Julian Eggert. A glimpse in chatgpt capabilities and its impact for ai research, 2023.
- [62] Timothy R. McIntosh, Teo Susnjak, Nalin Arachchilage, Tong Liu, Paul Watters, and Malka N. Halgamuge. Inadequacies of large language model benchmarks in the era of generative artificial intelligence, 2024.
- [63] Wenxuan Zhang, Sharifah Mahani Aljunied, Chang Gao, Yew Ken Chia, and Lidong Bing. M3exam: A multilingual, multimodal, multilevel benchmark for examining large language models, 2023.
- [64] Mingyu Derek Ma, Chenchen Ye, Yu Yan, Xiaoxuan Wang, Peipei Ping, Timothy S Chang, and Wei Wang. Clibench: A multifaceted and multigranular evaluation of large language models for clinical decision making, 2024.
- [65] Nicholas Lee, Thanakul Wattanawong, Sehoon Kim, Karttikeya Mangalam, Sheng Shen, Gopala Anumanchipalli, Michael W. Mahoney, Kurt Keutzer, and Amir Gholami. Llm2llm: Boosting llms with novel iterative data enhancement, 2024.
- [66] Irina Tolstykh, Aleksandra Tsybina, Sergey Yakubson, Aleksandr Gordeev, Vladimir Dokholyan, and Maksim Kuprashevich. Gigacheck: Detecting llm-generated content, 2024.
- [67] Yiheng Liu, Tianle Han, Siyuan Ma, Jiayue Zhang, Yuanyuan Yang, Jiaming Tian, Hao He, Antong Li, Mengshen He, Zhengliang Liu, Zihao Wu, Lin Zhao, Dajiang Zhu, Xiang Li, Ning Qiang, Dingang Shen, Tianming Liu, and Bao Ge. Summary of chatgpt-related research and perspective towards the future of large language models, 2023.
- [68] Prateek Verma. Wavelet gpt: Wavelet inspired large language models, 2025.
- [69] Marcel Lamott and Muhammad Armaghan Shakir. Leveraging distillation techniques for document understanding: A case study with flan-t5, 2024.
- [70] Saad Ezzini, Sallam Abualhaija, Chetan Arora, and Mehrdad Sabetzadeh. Ai-based question answering assistance for analyzing natural-language requirements, 2023.
- [71] Erik Blasch, James Sung, Tao Nguyen, Chandra P. Daniel, and Alisa P. Mason. Artificial intelligence strategies for national security and safety standards, 2019.
- [72] Edward Vendrow and Ethan Schonfeld. Understanding transfer learning for chest radiograph clinical report generation with modified transformer architectures, 2022.
- [73] Rachel S. Y. Teo and Tan M. Nguyen. Unveiling the hidden structure of self-attention via kernel principal component analysis, 2024.
- [74] Ziyang Chen and Stylios Moscholios. Using prompts to guide large language models in imitating a real person’s language style, 2024.
- [75] Savya Khosla, Zhen Zhu, and Yifei He. Survey on memory-augmented neural networks: Cognitive insights to ai applications, 2023.
- [76] Jianghui Wang, Yang Chen, Xingyu Xie, Cong Fang, and Zhouchen Lin. Task-robust pre-training for worst-case downstream adaptation, 2023.
- [77] Jiheng Liang, Ziru Yu, Zujie Xie, and Xiangyang Yu. A quick, trustworthy spectral knowledge qa system leveraging retrieval-augmented generation on llm, 2024.

-
- [78] Tam Nguyen, Tan M. Nguyen, Dung D. Le, Duy Khuong Nguyen, Viet-Anh Tran, Richard G. Baraniuk, Nhat Ho, and Stanley J. Osher. Improving transformers with probabilistic attention keys, 2022.
- [79] Da Song, Xuan Xie, Jiayang Song, Derui Zhu, Yuheng Huang, Felix Juefei-Xu, and Lei Ma. Luna: A model-based universal analysis framework for large language models, 2024.
- [80] Yixiong Chen and Weichuan Fang. Multi-scale feature fusion quantum depthwise convolutional neural networks for text classification, 2024.
- [81] Sean Xie, Soroush Vosoughi, and Saeed Hassanpour. Proto-lm: A prototypical network-based framework for built-in interpretability in large language models, 2023.
- [82] Sean Xie, Soroush Vosoughi, and Saeed Hassanpour. Towards interpretable deep reinforcement learning models via inverse reinforcement learning, 2024.
- [83] Minheng Ni, Yutao Fan, Lei Zhang, and Wangmeng Zuo. Visual-o1: Understanding ambiguous instructions via multi-modal multi-turn chain-of-thoughts reasoning, 2024.
- [84] David Alvarez-Melis, Tommi S. Jaakkola, and Stefanie Jegelka. Structured optimal transport, 2017.
- [85] Will Held, Caleb Ziems, and Diyi Yang. Tada: Task-agnostic dialect adapters for english, 2023.
- [86] Yaobo Liang, Chenfei Wu, Ting Song, Wenshan Wu, Yan Xia, Yu Liu, Yang Ou, Shuai Lu, Lei Ji, Shaoguang Mao, Yun Wang, Linjun Shou, Ming Gong, and Nan Duan. Taskmatrix.ai: Completing tasks by connecting foundation models with millions of apis, 2023.
- [87] Sanskar Soni, Satyendra Singh Chouhan, and Santosh Singh Rathore. Textconvonet:a convolutional neural network based architecture for text classification, 2022.
- [88] Shafiq Rayhan Joty, Giuseppe Carenini, and Raymond T Ng. Topic segmentation and labeling in asynchronous conversations, 2014.
- [89] Minxing Zhang, Ahmed Salem, Michael Backes, and Yang Zhang. Vera verito: Multimodal hijacking attack, 2024.
- [90] Matthew A. Wright and Joseph E. Gonzalez. Transformers are deep infinite-dimensional non-mercer binary kernel machines, 2021.
- [91] Lochan Basyal and Mihir Sanghvi. Text summarization using large language models: A comparative study of mpt-7b-instruct, falcon-7b-instruct, and openai chat-gpt models, 2023.
- [92] Ridewaan Hanslo. Deep learning transformer architecture for named entity recognition on low resourced languages: State of the art results, 2022.
- [93] Jan Buchmann, Max Eichler, Jan-Micha Bodensohn, Ilia Kuznetsov, and Iryna Gurevych. Document structure in long document transformers, 2024.
- [94] Christian Scheible and Hinrich Schuetze. Cutting recursive autoencoder trees, 2013.
- [95] Myeongjun Jang and Pilsung Kang. Sentence transition matrix: An efficient approach that preserves sentence semantics, 2019.
- [96] Julia El Zini and Mariette Awad. On the explainability of natural language processing deep models, 2022.
- [97] Christopher Grimm, Dilip Arumugam, Siddharth Karamcheti, David Abel, Lawson L. S. Wong, and Michael L. Littman. Modeling latent attention within neural networks, 2017.
- [98] Motoki Sato, Jun Suzuki, Hiroyuki Shindo, and Yuji Matsumoto. Interpretable adversarial perturbation in input embedding space for text, 2018.
- [99] Sean Xie, Soroush Vosoughi, and Saeed Hassanpour. Interpretation quality score for measuring the quality of interpretability methods, 2022.

-
- [100] Gabriele Sarti, Nils Feldhus, Ludwig Sickert, Oskar van der Wal, Malvina Nissim, and Arianna Bisazza. Inseq: An interpretability toolkit for sequence generation models, 2023.
- [101] Anna Schmidt and Michael Wiegand. A survey on hate speech detection using natural language processing. In *Proceedings of the fifth international workshop on natural language processing for social media*, pages 1–10, 2017.
- [102] Michael Grohs, Luka Abb, Nourhan Elsayed, and Jana-Rebecca Rehse. Large language models can accomplish business process management tasks, 2023.
- [103] Laura V. C. Quispe, Jorge A. V. Tohalino, and Diego R. Amancio. Using word embeddings to improve the discriminability of co-occurrence text networks, 2020.
- [104] Shree Charran R and Rahul Kumar Dubey. Meta-embeddings for natural language inference and semantic similarity tasks, 2020.
- [105] Tanish Tyagi, Colin G. Magdamo, Ayush Noori, Zhaozhi Li, Xiao Liu, Mayuresh Deodhar, Zhuoqiao Hong, Wendong Ge, Elissa M. Ye, Yi han Sheu, Haitham Alabsi, Laura Brenner, Gregory K. Robbins, Sahar Zafar, Nicole Benson, Lidia Moura, John Hsu, Alberto Serrano-Pozo, Dimitry Prokopenko, Rudolph E. Tanzi, Bradley T. Hyman, Deborah Blacker, Shibani S. Mukerji, M. Brandon Westover, and Sudeshna Das. Using deep learning to identify patients with cognitive impairment in electronic health records, 2021.
- [106] Sonia Meyer, Shreya Singh, Bertha Tam, Christopher Ton, and Angel Ren. A comparison of llm finetuning methods evaluation metrics with travel chatbot use case, 2024.
- [107] J. Ignacio Deza, Hisham Ihshaish, and Lamine Mahdjoubi. A machine learning approach to classifying construction cost documents into the international construction measurement standard, 2022.
- [108] Chanseo Lee, Sonu Kumar, Kimon A. Vogt, Sam Meraj, and Antonia Vogt. Advancing complex medical communication in arabic with sporo arasum: Surpassing existing large language models, 2024.
- [109] Fanny Jourdan, Laurent Risser, Jean-Michel Loubes, and Nicholas Asher. Are fairness metric scores enough to assess discrimination biases in machine learning?, 2023.

Disclaimer:

SurveyX is an AI-powered system designed to automate the generation of surveys. While it aims to produce high-quality, coherent, and comprehensive surveys with accurate citations, the final output is derived from the AI's synthesis of pre-processed materials, which may contain limitations or inaccuracies. As such, the generated content should not be used for academic publication or formal submissions and must be independently reviewed and verified. The developers of SurveyX do not assume responsibility for any errors or consequences arising from the use of the generated surveys.