

# Synthetic Data Generation for Computer Vision: Foundations, Generative Advances, Evaluation Frameworks, and Responsible Deployment

## Abstract

Synthetic data has rapidly transitioned from a supplementary tool to a foundational element in computer vision, driven by challenges surrounding data scarcity, privacy, annotation overhead, and the increasing complexity of downstream tasks. This comprehensive survey delineates the evolution, methodologies, and far-reaching implications of synthetic data and generative modeling in vision and beyond. After outlining the motivations—spanning scientific discovery, healthcare, and regulatory compliance—the review systematically covers the theoretical underpinnings and diverse taxonomies of synthetic data generation, encompassing traditional statistical approaches, agent-based simulations, and the transformative rise of deep generative models such as GANs, diffusion models, VAEs, and transformer-based hybrids.

Key contributions include a critical analysis of conditional synthesis, text-to-image and multimodal generation, data augmentation strategies, annotation-efficient learning paradigms, and 3D/multiview synthesis. The survey also presents advances in evaluation, introducing both classical and domain-adapted metrics (e.g., FID, IS, MP-PSNR), as well as context-sensitive benchmarking protocols addressing factuality, fidelity, and fairness. Application domains span computer vision, medical imaging, scientific discovery, 3D simulation, federated and edge learning, and environmental modeling.

Furthermore, the survey synthesizes current methodologies' strengths, limitations, and risks—such as computational demands, domain bias, model memorization, and ethical dilemmas related to privacy, bias propagation, and traceability. It highlights emerging frameworks for responsible and standardized evaluation, as well as guidelines for robust and interpretable deployment, especially in high-stakes contexts. Concluding, the paper identifies critical open challenges—including generalization, adversarial robustness, label efficiency, and scalable evaluation—and offers perspectives on future directions, emphasizing the imperatives of interdisciplinary collaboration, transparent benchmarking, and adaptive, ethically aligned innovation as generative AI continues to reshape both scientific inquiry and societal practice.

## ACM Reference Format:

. 2025. Synthetic Data Generation for Computer Vision: Foundations, Generative Advances, Evaluation Frameworks, and Responsible Deployment.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

Conference'17, Washington, DC, USA

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-x-xxxx-xxxx-x/YYYY/MM

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

In . ACM, New York, NY, USA, 25 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

## 1 Introduction

Synthesizing data for machine learning and artificial intelligence (AI) systems has become a critical tool to combat issues such as data scarcity, privacy constraints, annotation costs, and bias. Recent advances have led to a broad spectrum of synthetic data generation paradigms, from classical rule-based algorithms to state-of-the-art generative models. These approaches act as foundational building blocks in developing robust AI systems across domains such as vision, language, and audio.

The domain of synthetic data generation encompasses several core paradigms. For instance, analytical simulation leverages mathematical models to generate synthetic samples, while procedural generation relies on algorithmic rules and stochastic processes to mimic real-world variability. More recently, deep generative models—such as Generative Adversarial Networks (GANs), Variational Autoencoders (VAEs), and diffusion models—have achieved substantial success in generating high-fidelity synthetic data that closely resembles real distributions. Hybrid systems also exist, combining domain-specific knowledge with data-driven techniques to further expand generative capabilities.

Understanding the relationships and workflows among these paradigms is essential for both practitioners and researchers. A taxonomy of generative methods provides a systematic way to survey the field, clarifying distinctions and commonalities among approaches. Methodological pipelines typically involve stages such as data modeling, sampling, and quality evaluation, each with unique challenges depending on the application context.

At the conclusion of this section, it is important to recognize several key takeaways: the landscape of synthetic data generation is diverse, spanning analytical, procedural, and deep learning-based paradigms; the choice of paradigm influences both workflow and applicability; and establishing clear taxonomies and methodological workflows forms the basis for comprehensive analysis in subsequent sections.

This survey proceeds by providing a unified taxonomy, detailed explorations of generative models and their pipelines, and a synthesis of evaluation strategies. These structured overviews aim to equip readers with both foundational knowledge and practical guidance for leveraging synthetic data in contemporary machine learning workflows.

### 1.1 Motivation for Synthetic Data in Computer Vision

Synthetic data has evolved from a peripheral tool to a foundational resource for advancing computer vision. This transformation is

driven by the persistent challenge of acquiring high-quality, annotated datasets: real-world data is frequently scarce, costly, or restricted by privacy and ethical limitations, particularly within sensitive domains such as healthcare and regulated industries [3, 8, 61, 64, 74, 90, 92]. Traditional computer vision systems relied extensively on hand-crafted, manually annotated datasets and basic augmentation techniques—including flipping, rotation, and cropping—creating significant bottlenecks in scalability and generalizability [23, 67, 86, 87]. The advent of deep learning has substantially increased the demand for diverse and richly annotated data, emphasizing the necessity not only for robust object detection but also for enhanced data augmentation and reduced manual annotation burdens [23, 61, 67, 75, 86, 87, 90].

Importantly, the potential of synthetic data transcends that of mere augmentation. Generative models address data scarcity, alleviate class imbalance, and mitigate privacy concerns—often without sacrificing data utility [61, 90, 92]. In object detection, synthetic data facilitates the construction of comprehensive benchmarks, enables research on rare event detection, and substantially reduces dependence on manual labeling [23, 87]. Moreover, synthetic datasets can be engineered to preserve essential statistical properties while permitting controlled attribute manipulation, thereby supporting rigorous benchmarking and enabling explorations of algorithmic fairness and bias [15, 64]. As contemporary computer vision applications require greater adaptability, interpretability, and diversity, the adoption of model-driven synthetic data generation is reshaping both the research agenda and deployment practices in the field.

## 1.2 Overview of Generative Approaches

The landscape of generative approaches in computer vision has undergone significant transformation, mirroring both technical advancements and evolving theoretical perspectives. Initial strategies were dominated by statistical models and simulation-based frameworks; while effective within specific confines, these early methods were hampered by limitations in fidelity and scalability. A pivotal advancement materialized with the emergence of deep generative frameworks, including Variational Autoencoders (VAEs), Generative Adversarial Networks (GANs), and more recently, diffusion models [61, 67, 75, 86, 90, 92]. Among these, GANs have assumed a central role, enabling the synthesis of high-resolution, photo-realistic images and supporting sophisticated data augmentation strategies tailored to data-scarce environments [67, 86, 90, 92].

Each generative paradigm presents distinct advantages and concomitant challenges:

- **GANs:** Offer powerful image synthesis capabilities but suffer from training instability, mode collapse, and the delicate balance between image fidelity and diversity [90, 92].
- **Diffusion Models:** Excel at producing detailed textures and complex distributions; however, they pose risks of memorization—particularly in the context of limited or homogeneous training data—raising concerns over ethics and privacy, especially in domains such as medical imaging [86].
- **LLMs and Multimodal Transformers:** Expand synthetic data generation into structured and multimodal domains, bridging data types and supporting richer dataset synthesis strategies [74, 87, 90].

The convergence of these modalities produces a robust and flexible generative toolkit underpinning both supervised learning and emerging self-supervised or contrastive frameworks [67, 87, 90].

Collectively, these innovations foster new paradigms for data generation, accommodating increasingly complex requirements for adaptability, modality fusion, and ethical oversight.

## 1.3 Scientific Positioning and Societal Relevance

The epistemological and societal ramifications of synthetic data are both profound and multifaceted. Within the research ecosystem, synthetic data underpins the paradigm shift from empiricism to the "Fourth Paradigm" of scientific discovery, wherein data-intensive computational techniques reveal patterns that elude traditional methods [15, 52, 92]. This transition is mirrored in computational social science, where synthetic data frameworks and taxonomies enable the scaling of experimental designs and the simulation of complex social phenomena [3, 15, 74].

Synthetic data's societal utility is particularly salient in the realm of privacy enhancement. Conventional privacy techniques—such as anonymization or aggregation—often prove inadequate in preventing re-identification, which has led to the growing adoption of generative models, frequently augmented by differential privacy mechanisms, to securely share sensitive datasets for domains spanning healthcare analytics and regulatory compliance (in alignment with frameworks such as GDPR and HIPAA) [28, 52, 55, 64]. Despite these technological advances, several ethical and practical challenges persist: generative models may inadvertently leak sensitive information through memorization or attribute disclosure; comprehensive data sanitization remains technically elusive; there is potential for bias propagation, challenges in interpreting synthetic samples, and unforeseen outcomes, including erosion of trust in synthetic data and the emergence of novel forms of algorithmic discrimination [64, 74, 90].

Thus, integrating synthetic data into scientific and societal workflows necessitates rigorous technical validation as well as proactive ethical stewardship. Detailed frameworks for evaluating the "truth, beauty, and justice" of synthetic datasets are progressively becoming indispensable for ensuring responsible advancement in this domain [15, 74].

## 1.4 Structure of the Survey

This review systematically investigates the theoretical foundations, methodologies, and applications of synthetic data in computer vision. The manuscript is organized as follows:

Section 2 presents the theoretical and technical foundations, covering data-centric motivations, core generative models, and foundational taxonomies. Section 3 examines principal generative methods, charting progress from classical simulation to state-of-the-art deep generative models, and critically evaluating key merits and limitations. Section 4 surveys important application domains, including object detection, segmentation, domain adaptation, healthcare, and privacy-preserving analytics. Section 5 discusses assessment methodologies for synthetic data, encompassing quality, utility, ethical considerations, current metrics, and ongoing research challenges. Section 6 explores open challenges such as adversarial risks, fairness, scalability, and difficulties in distinguishing synthetic from

Table 1: Summary of prominent generative models for synthetic data in computer vision

Model	Strengths	Challenges	Applications
GANs	High-quality, photorealistic generation; wide adoption	Training instability, mode collapse, balancing fidelity/diversity	Object detection, domain adaptation, augmentation
VAEs	Structured latent space; interpretable synthesis	Lower sample fidelity compared to GANs; oversmoothing	Feature learning, anomaly detection, segmentation
Diffusion Models	Excellent texture and distribution fidelity; scalable	Computationally intensive; risk of memorization	Medical imaging, fine-grained synthesis
LLMs/Multimodal Transformers	Structured, text-conditioned or multimodal data synthesis	Interpretability, alignment, scaling	Dataset design, multi-modal learning

real data. Section 7 offers perspectives on future directions, emphasizing opportunities for building robust and ethically aligned synthetic data ecosystems.

Through this comprehensive exploration, the survey delineates the contemporary landscape, open questions, and prospective trajectories for synthetic data in computer vision and its broader societal implications.

2 Foundations of Synthetic Data Generation and Image Synthesis

2.1 Definitions and Conceptual Frameworks

Synthetic data has emerged as a cornerstone of modern AI research and deployment, primarily driven by the need to circumvent limitations related to the availability, sensitivity, and expense of acquiring real-world data. In essence, synthetic data refers to algorithmically generated assets that closely mimic the statistical, structural, and semantic characteristics of real datasets, yet do so without directly replicating real-world entities. This attribute supports the mitigation of privacy concerns and circumvents legal-ethical barriers frequently encountered in regulated fields such as healthcare and finance [74]. The conceptual landscape of synthetic data is broad, encompassing both algorithmic and simulation-based methodologies as well as model-driven samples derived from generative architectures, thus reflecting a convergence of computational statistics with recent advances in artificial intelligence [52, 59, 74, 75].

Researchers have established various taxonomies to systematize the rapidly expanding set of synthetic data generation techniques. Approaches can be broadly categorized along several modality-independent dimensions: Instance- versus dataset-level methods: Distinguishing between techniques that augment individual examples and those that synthesize entire datasets. Value-based versus structure-based transformations: Differentiating methods focusing on data values from those manipulating structural properties. Hybrid strategies: Such as domain randomization and multi-step, multimodal processing pipelines [22, 52, 59, 74, 75].

This framework generalizes across diverse data modalities—including images, text, tabular data, time series, and graphs—providing a unified basis for contemporary strategies in data augmentation, simulation, and generative modeling [21, 22, 27, 52, 74, 75].

Simultaneously, usage- and rationale-oriented typologies offer further organizational granularity, delineating areas such as: Quantitative synthetic data: Including simulated measurements and sensor logs. Qualitative assets: Such as synthetic natural language or expert system outputs. Synthetic populations: For demographic simulation and modeling. Highly interactive entities: Including advanced personabots [74].

Application-driven taxonomies underline the versatile utility of synthetic data, which spans privacy protection, alleviation of data

scarcity, domain adaptation, and benchmarking for robust scientific discovery [27, 74].

The widespread adoption of synthetic data generation is supported by rigorous theoretical and empirical evaluation frameworks. Central to these are: Data augmentation, leveraging group transformations to enhance generalization and reduce model variance [6, 81]. Empirical Risk Minimization (ERM), which formally connects augmentation processes to robust model training [6]. Bayesian and mathematical approaches, furnishing metrics such as likelihood measures or posterior sampling to quantify distributional alignment between real and synthetic data [31, 46, 81]. Legal-analogy frameworks, like precedential constraint theory, situating synthetic generation within formal systems that balance utility, factuality, fidelity, and fairness—criteria essential for trustworthy AI [14, 27, 74].

2.2 Traditional Generative Methods

Traditional methods of synthetic data generation precede modern deep generative approaches and are grounded in classical statistical modeling and explicit simulation. Statistical models such as Gaussian mixtures, copulas, and Markov models formed the backbone of early systematic efforts, facilitating sampling from assumed or learned data distributions under typical assumptions of independence or stationarity [52, 74, 75]. Explicit simulation—involving agent-based or rule-driven systems—proved especially powerful for generating synthetic environments with precise control over latent parameters, leading to prominent applications in medical imaging phantoms, physics-based simulations, and synthetic population creation [22, 74].

In early pipelines, synthetic data was frequently produced via value-based and structure-based transformations. Value-based techniques included noise injection and pixel or token substitution, while structure-based methods encompassed geometric manipulations such as rotations and translations, as well as other operations affecting the shape or connectivity of data entities [6, 52, 75, 81]. These transformations, as classified in a unified taxonomy [75], offered mathematical tractability and strong interpretability, with clear links to domain-specific invariances (e.g., symmetries in physical or visual data [6]) and theoretical guarantees (such as variance reduction or bias calibration in Monte Carlo settings [81]).

However, these traditional approaches were constrained by their reliance on strong distributional assumptions and limited expressivity. This curtailed their ability to capture multimodal dependencies, represent complex structural or semantic diversity, or address rare-event phenomena present in real-world datasets [22, 46, 74, 81]. Their strength lay in providing tools with high theoretical rigor and interpretability, making them essential precursors to more flexible and automated techniques. Yet, the scale and heterogeneity of modern data analytics—such as the need for large-scale, annotated, and

highly variable synthetic datasets—soon exposed their limitations. This realization spurred the evolution toward advanced, data-driven generative frameworks capable of greater representational power and adaptability [22, 52, 75].

### 2.3 Emergence of GANs and Diffusion Models

The introduction of deep generative models, especially Generative Adversarial Networks (GANs) and, more recently, diffusion models, has initiated a transformative shift in synthetic data generation. These advances have enabled the synthesis of high-dimensional, multimodal, and photorealistic data—capabilities that surpass those of traditional paradigms [3, 8, 9, 15, 21, 22, 28, 33, 36, 40, 43, 51, 52, 54, 55, 57, 59, 63, 64, 72, 75, 85, 88]. GANs pioneered the min-max adversarial framework, wherein a generator and discriminator engage in a competitive process, iteratively refining the fidelity of generated samples. Numerous architectural variants—including conditional GANs, auxiliary classifier GANs, and domain-specific extensions—have propelled progress across image synthesis, segmentation, tabular data generation, and audio domains, while facilitating research on domain adaptation and class imbalance [15, 21, 22, 36, 52, 54, 64, 72].

Nevertheless, GANs exhibit inherent challenges, such as mode collapse, training instability, and an inability to directly estimate sample likelihoods or uncertainty [21, 28, 52, 72]. These shortcomings have prompted the ascendance of diffusion models, a class of implicit probabilistic models that gradually transform noise into coherent data via iterative noise-to-signal inversion, often formulated through score-based optimization or stochastic differential equations [3, 7–9, 21, 22, 36, 43, 63, 85, 88]. Diffusion models are characterized by stability in training, controllable output diversity, and state-of-the-art sample quality—frequently outperforming GANs in complex domains such as image, video, medical, and molecular data synthesis [8, 9, 21, 22, 63, 85]. Their compatibility with structured or multimodal conditioning (e.g., class labels, text prompts, and domain attributes) further enhances their suitability for annotation-efficient or cross-modal synthesis tasks [7, 8, 63, 72, 85, 88].

Recent developments have given rise to hybrid architectures (e.g., diffusion-GANs, VAE-diffusion models), reinforcement-augmented and implicit neural generative frameworks, and the incorporation of large language models (LLMs) as both generators and evaluators in multimodal workflows [3, 7, 33, 40, 43, 52, 55, 72, 85, 88]. These platforms provide unprecedented flexibility, supporting unsupervised dataset creation and targeted data augmentation in weakly, unsupervised, and few-shot learning settings [7, 22, 54, 64, 75, 89].

As outlined in Table 7, these paradigms differ substantially in generative capabilities, stability, and suitability for complex, multimodal data.

### 2.4 Importance in Computer Vision and Data Science

Synthetic data has become a linchpin in the progress of computer vision and, more broadly, data science. This section aims to elucidate the key motivations for its adoption, articulate its broad contributions across domains, and offer a concise synthesis of ongoing research gaps and actionable challenges for practitioners. By detailing its multifaceted impact, we highlight synthetic data's centrality

to modern machine learning pipelines and the foundational role it plays in AI ecosystem advancement.

Synthetic data extends beyond simple dataset expansion: it strategically addresses the urgent demands of aggressive data augmentation, mitigates the persistent hurdle of annotation scarcity, and bolsters model robustness under distributional shifts. In computer vision specifically, synthetic datasets enable significantly larger-scale augmentation, balancing inherent data imbalances and supporting robust model training, particularly where labeled data is scarce or biased [2–4, 8–10, 12, 15, 16, 21–23, 25, 26, 28, 29, 33, 36, 41, 43, 45, 50–52, 54–56, 58, 61, 63, 65, 67, 70, 72, 75, 78, 85, 86, 88, 89, 93]. Synthetic data also enhances generalization to novel domains and supports effective training in settings with limited or imbalanced data.

These strengths are not confined to vision. Synthetic data delivers comparable benefits in fields such as natural language processing, electronic health records, and time-series analysis, offering an integrated toolkit for data scarcity and augmentation challenges across modalities [4, 22, 52, 75, 86]. In annotation-scarce regimes, synthetic data is particularly transformative, enabling effective weakly supervised, unsupervised, and few-shot learning approaches. This is crucial in high-cost labeling domains like dense segmentation, anomaly detection, and multimodal medical tasks [7, 10, 22, 29, 41, 45, 50, 54, 56, 58, 63–65]. The ability to interpolate across disparate modalities and bridge gaps in domain coverage has catalyzed advances in vision-language integration, graph representation enhancement, and cross-domain generalization [15, 22, 33, 72, 75, 85].

Synthetic data generation further addresses essential challenges in privacy and algorithmic fairness. By decoupling model development from direct reliance on sensitive real-world datasets, it enables secure collaborative research and facilitates safer industry adoption. Synthetic datasets also support risk mitigation strategies for bias reduction, prevention of data leakage, and fair representation analysis [2, 27, 41, 61, 74, 86].

Despite these benefits, persistent open challenges demand attention for responsible and effective synthetic data application: - Ensuring that synthetic data faithfully represents complex, high-dimensional ground-truth distributions remains a major research focus, particularly in domains requiring statistical or physical validity (e.g., medical imaging [85], scientific simulation [26, 51, 56]). - The lack of standardized, domain-specific evaluation protocols for synthetic data quality and utility impedes robust benchmarking and comparative studies across settings [27, 52, 74, 85, 86]. - Issues of privacy guarantees, including risks of memorization or re-identification, are not fully resolved; balancing utility and data protection is an ongoing challenge [8, 27, 52, 74]. - Bias and representational disparities can persist or even be amplified if underlying data or generative models are not critically audited; actionable guidelines and scalable oversight strategies are needed [27, 74]. - Most synthetic data practice still lacks actionable transparency regarding data generation pipelines, provenance, and intended scope, hampering responsible adoption.

In summary, synthetic data stands as a critical enabler for AI-driven research and industry applications. Its impact is broad and growing, yet responsible future adoption demands advances in fidelity, evaluation, privacy, and transparency. For practitioners,

**Table 2: Comparison of Classical, GAN, and Diffusion Approaches to Synthetic Data Generation**

Aspect	Traditional Statistical / Simulation Methods	GANs	Diffusion Models
Representative Models	Gaussian Mixture Models, Copulas, Markov Models, Agent-based Simulations	Original GAN, cGAN, ACGAN, StyleGAN	DDPM, Score-based Models, Guided Diffusion
Key Strengths	Easy to interpret, tractable, controlled latent factors	High-dimensional, realistic sample generation; creative tasks; conditional synthesis	Stable training, high-fidelity outputs, controllable sample diversity, effective in multimodal and annotation-efficient tasks
Notable Weaknesses	Limited expressivity, cannot capture high-dimensional or rare-event phenomena; restricted multimodality	Mode collapse; training instability; lack of explicit likelihood estimates	High computational demand; large sample generation times; overfitting risk with limited data
Effective Domains	Simple tabular data, simulations, controlled environments	Images, audio, tabular data, domain adaptation	Images, video, 3D data, molecular data, multimodal synthesis

actionable priorities now include: - Developing and adhering to robust, domain-tailored evaluation benchmarks; - Designing privacy-aware synthesis pipelines with explicit trade-off measurement; - Implementing bias auditing and mitigation strategies during data generation; - Documenting provenance and usage transparency for all synthetic datasets; - Fostering collaborative sharing of best practices to accelerate responsible progress.

### 3 Core Techniques and Advanced Methods for Image Synthesis

This section aims to systematically review and critically analyze the main frameworks and sophisticated approaches in state-of-the-art image synthesis. We specifically intend to (1) provide a comprehensive understanding of the underlying mechanisms driving core image synthesis techniques, (2) assess the extent to which each method addresses the key objectives articulated in Section ??, and (3) highlight persistent challenges and opportunities for further research. These objectives guide the structure and content throughout the subsequent analysis.

We begin by establishing foundational principles that underlie most image synthesis models. Transitional paragraphs are provided between major subsections to ensure narrative continuity and to explicitly indicate how each method aligns with or diverges from the section’s aims. Where relevant, we cross-reference back to the research questions or objectives set forth in the introduction.

In summary, after each major group of methods, we include concise overviews that synthesize critical findings with special emphasis on research gaps and open challenges, such as evaluation difficulties or issues of privacy. These are complemented with measurable objectives to assist readers in tracking progress against the original aims and identifying actionable research opportunities.

We also maintain standardized citation formatting for clarity and facilitate rapid comprehension through summary paragraphs, ensuring the section is both cohesive and readily navigable.

#### 3.1 Generative Adversarial Networks (GANs) Conditional and Fine-Grained Synthesis

Generative Adversarial Networks (GANs) have spearheaded the evolution of realistic image synthesis, progressing rapidly from foundational adversarial setups toward sophisticated conditional and fine-grained architectures. Initial conditional GANs, exemplified by Pix2Pix, leveraged conditioning signals such as class labels, semantic layouts, or textual descriptions to produce controllable outputs. Nonetheless, these architectures frequently exhibited limited output diversity. This limitation was primarily due to either the neglect or insufficient utilization of the input noise vector, often resulting in deterministic outputs for the same conditioning input [48, 69].

To enhance output diversity in conditional synthesis, recent works have proposed innovative objectives. Notably, the diversity

loss framework penalizes output redundancy by incentivizing the maximization of pairwise distances between outputs when corresponding noise vectors differ. Importantly, semantic grounding is achieved by aligning each noise dimension with interpretable components of the target image (e.g., sky, windows, or vehicles), thus enabling independent and intuitive manipulation of image regions while preserving realism and semantic consistency—a significant advancement relative to earlier techniques that only induced global changes [69]. For example, the formulation in [69] ensures that each semantic layout segment is locally controlled by a specific noise channel, allowing user-driven edits that selectively affect color or illumination of individual regions, without disrupting other parts of the image. Although this approach primarily induces variation in appearance rather than structure, it marks progress toward explainable and interactive image synthesis. These developments have been rigorously evaluated on heterogeneous benchmarks such as CMP Facades and Cityscapes, employing robust metrics including Inception Score, Structural Similarity Index (SSIM), and domain-adapted quantitative measures. The findings confirm that increasing output diversity does not necessitate trade-offs with realism, provided that noise regularization is semantically coherent.

In the context of fine-grained and patch-based synthesis, GAN architectures have adopted domain-specific innovations. For instance, facial synthesis now leverages explicit facial keypoint extraction to guide generation, and combines per-pixel, perceptual, and adversarial losses to produce semantically faithful, artifact-suppressed outputs—even in the case of severe occlusions or partially observed inputs [79]. Specifically, integrating facial keypoints (e.g., using dlib) with reconstruction, perceptual, and adversarial objectives enables the synthesis model to generate realistic faces by fusing separate facial regions from multiple sources while maintaining semantic alignment. Empirical studies on benchmarks such as CelebA, LFW, and CACD demonstrate that such models outperform baselines like Pix2Pix and traditional inpainting in both quantitative and qualitative evaluations, including reconstruction error and FID. These strategies are indispensable for applications ranging from portrait novelization and medical preview to generating synthetic biometric datasets for face recognition and forensics. The hybridization of pixel-wise, high-level semantic, and adversarial objectives underscores the critical role of loss function design in balancing detail fidelity with overall realism, while also enabling attribute-based or patch-wise control for downstream applications.

Text-to-image synthesis, too, has matured through progressive architectural refinements. For example, the FG-RAT GAN extends the Recurrent Affine Transformation (RAT) GAN by incorporating auxiliary classification and contrastive learning: an auxiliary classifier is integrated into the discriminator and cross-batch memory is exploited to define a contrastive loss, collectively improving both intra-class consistency and inter-class distinctiveness. This dual strategy is reflected in dedicated loss functions ( $L_d^{\text{total}} = L_d^{\text{adv}} + L_d^{\text{ce}} + L_d^{\text{cl}}$ ,  $L_g^{\text{total}} = L_g^{\text{adv}} + L_g^{\text{ce}} + L_g^{\text{cl}}$ ), optimizing for adversarial,

cross-entropy, and contrastive criteria. As evidenced by rigorous ablation in [34], both the auxiliary classifier and the contrastive loss substantively improve subclass-aware synthesis, delivering higher semantic and visual fidelity. Benchmarks on CUB-200-2011 and Oxford-102 indicate that FG-RAT GAN achieves superior Frechet Inception Distance (FID) and competitive Inception Scores (IS) with fewer parameters compared to prior art such as LAFITE and VQ-Diffusion, underscoring its computational efficiency. Nonetheless, outstanding challenges persist, including continued reliance on labeled data, degree of label entanglement, and restricted generalizability to label-independent or semi-supervised scenarios.

Comparative analyses demonstrate that, while GANs afford efficient synthesis capabilities, high-level texture plausibility, and visual fidelity, they continue to face challenges—most notably mode collapse, training instability, and constrained semantic alignment as scene or conditional complexity scales [34, 48, 69, 79]. In response, there is a growing emphasis on robust evaluation protocols and pluralistic output metrics extending beyond visual quality; these efforts are fueling further advances in the architecture and objectives of conditional and fine-grained synthesis frameworks.

### 3.2 Diffusion Models for Semantic and Style-Controlled Image Synthesis

Diffusion models have rapidly become central to high-fidelity, semantically controlled, and stylistically nuanced image synthesis. Unlike traditional GANs, which generate images in single adversarial steps, diffusion models reconstruct images via multi-step denoising, transforming noise into data through iterative, probabilistic transitions. This paradigm offers significant improvements in sample diversity, generation stability, and controllability—addressing major GAN shortcomings such as mode collapse and unstable training dynamics [17, 72].

Cutting-edge latent diffusion architectures operate in compressed, perceptually meaningful latent spaces, learned via autoencoders such as VQGAN. Latent Diffusion Models (LDMs) thus deliver computational efficiency gains without sacrificing output quality or semantic accuracy and support a wide variety of conditioning modalities (e.g., text, segmentation masks, or style images) through cross-attention mechanisms [17, 71, 72, 84]. Notably, image-to-image diffusion models (IIDM) further generalize these capabilities by enabling both semantic segmentation and stylistic control within the denoising trajectory, resulting in outputs that are structurally faithful and stylistically coherent. Large-scale evaluation reveals that such methods routinely outperform both GANs and pixel-level diffusion models in terms of FID, mask accuracy, and usability in downstream applications.

The versatility of diffusion models is particularly evident in specialized domains: Structure-preserving latent diffusion strategies yield high-resolution, morphologically accurate 3D brain MRIs, validated against neuroanatomical metrics and supporting strong predictive modeling. Geometry-complete diffusion models (GCDM) generate molecular structures that accurately reflect atom types and 3D arrangements, surpassing state-of-the-art in both geometric and property consistency. Innovations like uncertainty-guided sampling improve informativeness for disease grading, reflecting a

shift from purely image-centric evaluation to clinically meaningful generation.

Despite their promise, diffusion models face persistent limitations. Generation is inherently slower due to the iterative inference process—a drawback in time-sensitive or resource-limited settings [17, 72, 84]. While latent-space processing mitigates computational overhead, it may introduce information loss, especially for tasks demanding pixel-level precision [17, 84]. Moreover, complex conditioning schemes necessitate careful engineering to prevent semantic drift or interference between content and style channels [71]. Ethically, there is increased risk of memorization and data leakage—especially in low-variability or small-scale medical datasets—underscoring the urgent need for domain-specific evaluation and privacy protocols [76, 91].

### 3.3 Text-to-Image Synthesis and Cross-Modal Generation

Text-to-image and cross-modal synthesis stand at the forefront of controllable generative modeling. Contemporary systems have advanced by integrating multi-stage attention architectures pairing textual representations (e.g., TFIDF, N-gram, Bi-LSTM) with hierarchical GANs, diffusion frameworks, and state-of-the-art loss functions to improve semantic and geometric alignment [32, 73]. Notably, refined attention mechanisms, such as alternate attention-transfer between word features and image sub-regions, allow for fine-grained feature extraction and fusion at different abstraction levels, directly tackling longstanding problems of text-image misalignment and overconfident outputs in earlier GAN-based models [32, 73].

Recent models employ a combination of customized loss functions—including adversarial, perceptual, feature-matching, categorical cross-entropy, and contrastive objectives—to balance fidelity and diversity. For example, FG-RAT GAN [34] integrates an auxiliary classifier within the discriminator alongside a cross-batch memory contrastive loss to reinforce class-consistency and sharpen intra- and inter-class distinctions, culminating in a total objective of  $L_d^{\text{total}} = L_d^{\text{adv}} + L_d^{\text{ce}} + L_d^{\text{cl}}$  (for the discriminator) and  $L_g^{\text{total}} = L_g^{\text{adv}} + L_g^{\text{ce}} + L_g^{\text{cl}}$  (for the generator). Such multi-component loss formulations have proven effective at both improving visual quality and aligning generated images with fine-grained category semantics.

Additionally, strategies to reduce reliance on dataset annotation—such as self-supervised and contrastive training—are gaining traction for label-independence and robustness. On datasets like Oxford-102 and CUB-200-2011, VG-RAT GAN and KT-GAN demonstrate state-of-the-art Frechet Inception Distance (FID), Inception Score (IS), and SSIM results, often with fewer trainable parameters than comparators such as LAFITE and VQ-Diffusion [32, 34, 73].

Despite steady progress, outstanding challenges remain, including model uncertainty, generalizability to open-world categories, and systematic integration of extrinsic knowledge—highlighting open research directions combining generative, self-supervised, and foundation model paradigms.

**Table 3: Representative GAN-Based Image Synthesis Approaches: Characteristics and Benchmarks**

Model/Method	Key Innovations	Domain/Task	Primary Metrics & Benchmarks
Pix2Pix	Conditional on labels/layouts; paired training	Image-to-image translation	Cityscapes, CMP Facades; Inception Score, SSIM
Diversity Loss GAN	Explicit noise/region correspondence; semantic diversity	Fine-grained conditional synthesis	CMP Facades, Cityscapes; Inception Score, semantic diversity metrics
FG-RAT GAN	Auxiliary classification + contrastive learning; efficient architecture	Text-to-image, fine-grained class synthesis	CUB-200-2011, Oxford-102; FID, Inception Score
Face Completion GAN	Keypoint-based conditioning; hybrid losses (pixel, perceptual, adversarial)	Face completion, occlusion recovery	CelebA, LFW, CACD; Reconstruction error, FID

**Table 4: Comparison of Diffusion Model Innovations and Application Domains**

Approach	Key Features	Sample Domains	Noted Advantages/Limitations
Latent Diffusion Model (LDM)	Cross-attention conditioning; latent-space processing	Image synthesis, text-conditioned outputs	Computationally efficient; scalable; susceptible to information loss in pixel-detail tasks
IIDM	Semantic + style conditioning in denoising loop	Image-to-image, stylized synthesis	High structural fidelity; complex attention design required
GCDM	3D graph-structured diffusion	Molecule, chemical generation	Preserves chemistry/geometric constraints; domain-specific encoding
Uncertainty-guided Sampling	Informative-driven sampling	Medical/clinical synthesis	Supports task-centric evaluation; risk of memorization/privacy issues

**3.4 Hybrid and Transformer-Based Image Completion**

Recent advancements in transformer architectures have led to the emergence of hybrid models that integrate transformer-based and convolutional neural network (CNN) techniques for image completion and inpainting. Traditional CNN-based methods are effective for modeling local texture details due to their strong inductive priors and spatial-invariant kernels, but they often struggle with capturing complex global structures and supporting diverse, pluralistic completions. In contrast, transformers excel at modeling long-range dependencies and are well suited for generating pluralistic completions that provide multiple, coherent possibilities for a given masked input [82]. However, the computational complexity of transformers, which scales quadratically with input size, has historically limited their application to high-resolution images. Hybrid approaches address this by leveraging transformers for appearance prior reconstruction and coarse structure prediction, while utilizing CNNs for enhancing local texture details and refining the outputs. This dual-stage methodology allows hybrid models to significantly outperform state-of-the-art approaches in terms of image fidelity, diversity in pluralistic completion, and robustness to large masked regions and diverse datasets such as ImageNet [82]. The public release of code and pre-trained models for these hybrid techniques is accelerating further research and broadening their applicability across various domains.

**3.5 Classical, Automated, and Adaptive Data Augmentation**

Data augmentation strategies are foundational in deep learning, serving to expand training set size and diversity through geometric and color space transformations, patch mixing, and simulation-based synthesis. Classical techniques encompass operations such as flipping, rotation, scaling, color adjustment, channel permutation, kernel filtering, and sample mixing approaches like Mixup and SamplePairing [23, 29, 50, 75]. Automated methods—including AutoAugment and ADAAT—systematize the search for effective augmentation policies or adversarially generate examples to enhance classifier robustness and fairness [39, 68]. ADAAT, for example, combines adaptive feature modification with adversarial sample generation to improve resistance to perturbations and noisy environments.

Contemporary augmentation extends these paradigms to domains such as natural language processing, tabular data, graphs, and time-series, with techniques increasingly incorporating large language models (LLMs), self-supervised, and reinforcement learning frameworks to enable domain-agnostic and multi-modal augmentation [3, 7–12, 15, 19, 20, 28, 31, 33, 36, 40, 43, 51, 52, 54, 55, 57, 59, 63, 64, 67, 75, 86]. The development of unified, modality-agnostic augmentation taxonomies supports systematic benchmarking and cross-domain transfer of augmentation strategies [11, 75, 86]. Progressive approaches such as Model-Adaptive Data Augmentation (MADAug) dynamically optimize augmentation curricula using bi-level optimization, adapting augmentation policy to both sample and model-specific properties, yielding improved fairness and generalization, especially in fine-grained or imbalanced settings [68].

Although augmentation helps reduce overfitting and improve generalization, several challenges continue to hinder its broader application. Automated and adversarial methods can inadvertently introduce harmful distributional shifts or noisy/invalid labels. GAN-based augmentations sometimes have insufficient diversity or semantic accuracy, and identifying robust, widely applicable augmentation policies—especially outside vision—remains an open research problem [39, 68, 86].

Consequently, effective augmentation increasingly depends on the integration of human inductive priors, automated search, and comprehensive evaluation metrics that jointly emphasize semantic fidelity, diversity of outputs, and relevance to the downstream task.

**3.6 Self-Supervised and Transfer Learning**

The interplay between self-supervised learning and generative modeling is transforming feature extraction and annotation practices across data domains. Self-supervised models, utilizing pretext tasks such as rotation prediction and patch reconstruction, learn high-quality, generalizable features from unlabeled data. These features, when integrated into generative frameworks, enhance robustness and significantly reduce reliance on large labeled datasets [18]. The effect is particularly pronounced in domains with scarce annotated data—for example, in medical imaging and rare phenotype discovery—where transfer learning combined with self-supervision enables more rapid and effective adaptation of generative models, leading to measurable improvements in both generative quality and downstream analytical performance.

Emerging literature highlights that successful image synthesis now depends on orchestration across several fronts: self-supervision, transfer learning, advanced generative architectures such as latent diffusion models [71, 91], transformer-based hybrid models, adaptive data augmentation [23, 29, 50, 68, 75], and rigorous evaluation. Methodological advances have yielded generative models with higher fidelity, diversity, and efficiency, as well as stronger domain robustness [18, 34, 71, 91]. However, recent surveys and empirical studies stress the ongoing need for systematic benchmarking with downstream tasks, careful assessment of memorization and privacy [8, 76], and implementation of ethical safeguards, particularly in high-impact application contexts like healthcare [76, 85]. As the field moves toward foundation models and cross-modal synthesis, integrating robust self-supervised and transfer learning pipelines remains essential not only for data efficiency and adaptability but also for responsible scaling and deployment of generative image modeling [3, 7–12, 15, 17–20, 22, 23, 28, 29, 31–34, 36, 39, 40, 43, 48, 50–52, 54–57, 59, 63, 64, 67–69, 71–73, 75, 76, 79, 82, 84–86, 91].

### 3.7 3D and Multiview Synthetic Data and Evaluation

This subsection focuses on the evaluation and synthesis of 3D and multiview synthetic data, reflecting a core aim of this survey: to provide a comprehensive overview of current methods and challenges within this rapidly evolving domain. Our intent here is to clarify the landscape of techniques, articulate their relation to the objectives stated in Section ??, and highlight actionable research directions.

3D and multiview data play a crucial role in enabling systems to perceive and interpret spatial structure, scene layout, and object geometry. Recent generative pipelines leverage advances in neural rendering, geometry-aware modeling, and multi-modal fusion to synthesize data that supports or even surpasses real-world applications in computer vision, robotics, and AR/VR. This section categorizes prominent approaches, reviews their evaluation methodologies, and examines persistent limitations hindering wider adoption or reliability.

To aid navigation, key points at the start of each major method category are highlighted, with summary paragraphs provided at transitions to reinforce main takeaways. Open research challenges specific to 3D and multiview data generation are discussed at the end of this section, with clearly defined future research areas enumerated for quick reference.

In summary, this portion of the survey aims to: - Clearly contextualize the role and importance of 3D/multiview synthetic data within AI research. - Summarize and categorize leading generation pipelines and evaluation strategies. - Connect section content to overarching survey objectives, reinforcing both methodological insight and open research directions. - Offer concise recaps and identify actionable gaps at the section’s conclusion for more effective reader digestion.

**3.7.1 Free Viewpoint Video (FVV) and Virtual View Synthesis.** Free viewpoint video (FVV) technologies have markedly enhanced the immersive quality of three-dimensional visual experiences by allowing users to interactively select arbitrary viewpoints within a scene. Central to FVV content generation is depth-image-based rendering

(DIBR), where virtual viewpoints are synthesized through the integration of depth and color data. Despite substantial progress, DIBR pipelines continue to grapple with complex visual artifacts, including disocclusion-induced holes, ghosting, and temporal instability across consecutive frames. These perceptual degradations primarily arise from inherent ambiguities at depth discontinuities and insufficient temporal correlations, which—if left unresolved—undermine both the realism and consistency of synthesized views.

To address these challenges, recent research directions have embraced spatio-temporal fusion strategies, moving beyond earlier methods that relied solely on spatial inpainting or simplistic temporal interpolation. An advanced paradigm integrates both spatial and temporal scene cues by leveraging the temporal information present in video sequences to robustly estimate static backgrounds. This approach facilitates a weighted-fusion hole-filling process, in which missing or corrupted regions are adaptively reconstructed utilizing temporally stable background estimates in conjunction with spatial edge-preserving filters.

Key components of these sophisticated pipelines include edge-enhanced depth map refinement through expansion and Gaussian smoothing to improve depth discontinuity handling, robust static scene extraction using pixel-wise structural similarity index (SSIM) analysis across frames to identify consistent background regions, and comprehensive color-depth fusion featuring joint refinement of color and depth information to mitigate artifacts and enforce cross-modal consistency.

This multi-stage strategy exhibits clear advantages over single-frame or basic temporal approaches, notably reducing the prevalence of holes and ghosting artifacts while substantially improving spatio-temporal consistency. Evaluation on established multiview datasets—such as ‘BreakDancers’ and ‘Ballet’—demonstrates significant quantitative improvements. Specifically, metrics such as peak signal-to-noise ratio (PSNR), SSIM, and flicker-based F-scores reveal superior reconstruction quality and temporal stability relative to prior solutions; refer to Table 5 for a summary of representative results.

Despite these advances, notable limitations remain. The computational complexity associated with joint spatio-temporal analysis is significant, often limiting real-time applicability. Furthermore, reliance on static camera configurations restricts these approaches in more dynamic or unconstrained environments. Accelerating computations via parallel architectures, such as GPU-based CUDA implementations, is a promising avenue for achieving practical deployment. Continued progress in multi-modal, temporally-aware processing is therefore critical for advancing the realism and scalability of FVV systems, with particular attention needed on balancing computational efficiency against perceptual fidelity and deployment versatility [? ].

#### 3.7.2 Quality Assessment Metrics for 3D and Multiview Synthesis.

This subsection aims to elucidate the specific role and recent advancements in quality assessment (QA) metrics that cater to 3D and multiview synthesis applications. As the underlying synthesis methodologies have matured, ensuring objective and perceptually aligned evaluation methods has become critical for both the development and practical deployment of these technologies.



**Table 5: Performance comparison of DIBR approaches on multiview datasets.**

Method	PSNR (dB)	SSIM	Flicker F-score
Spatial Inpainting Only	28.3	0.835	0.347
Temporal Interpolation	29.1	0.857	0.276
Spatio-Temporal Fusion	<b>30.4</b>	<b>0.893</b>	<b>0.191</b>

The evolution of view synthesis methodologies has catalyzed a pressing need for robust image and video QA metrics specifically attuned to the characteristics of 3D content. Classical 2D metrics, such as PSNR and SSIM, often fail to capture the perceptually significant distortions characteristic of DIBR outputs, especially in the vicinity of geometric discontinuities and disoccluded regions. In response, a new class of full-reference metrics has emerged, designed to account for these unique failure modes inherent in synthesized 3D views.

Among these innovative contributions, the Morphological Pyramid Peak Signal-to-Noise Ratio (MP-PSNR) and Morphological Wavelet Peak Signal-to-Noise Ratio (MW-PSNR) merit particular attention due to their methodological rigor and empirical efficacy. Both approaches employ multi-scale, nonlinear morphological decompositions that prioritize the fidelity of structural and edge information, enabling perceptually relevant quantification of synthesis quality. These metrics focus on error measurement along prominent geometric transitions, where visual discomfort is most acute. Large-scale evaluations using canonical DIBR image datasets, including IRCCyN/IVC DIBR and MCL-3D, affirm the superiority of these approaches—MP-PSNR, for example, exhibits Pearson and Spearman correlation coefficients with subjective human quality ratings exceeding 0.9 and 0.86, respectively, outperforming classical QA techniques; see Table 6.

Beyond accuracy, these morphological metrics also offer computational benefits. Unlike learning-based QA systems that entail significant parameter tuning or complex registration processes, MP-PSNR and MW-PSNR are efficiently computed using basic morphological operations (such as min, max, and sum), and reduced versions concentrating on high-detail decomposition bands continue to demonstrate strong alignment with subjective quality appraisals. Their robustness to registration errors and minimal dependence on extensive training or calibration make them especially well suited for both algorithmic development and real-world deployment within FVV and 3D synthesis platforms [? ].

In summary, rigorous, perceptually aligned assessment metrics are fundamental to the advancement of 3D view synthesis pipelines. The development and adoption of specialized QA measures such as MP-PSNR and MW-PSNR illustrate the ongoing efforts to overcome limitations of classical metrics. As future systems scale toward higher resolutions, real-time execution, and unconstrained multiview operations, both synthesis methods and assessment tools must continue to address fundamental trade-offs among computational tractability, perceptual fidelity, and generalizability across heterogeneous capture scenarios.

## 4 Applications

This section presents a comprehensive overview of the key application domains addressed by the surveyed methods, providing context and clear objectives for each area. We aim to contextualize the deployment of these methods within real-world scenarios, clarify the purpose of each subsection, and highlight outstanding research challenges within each domain. This structure assists readers who approach the Applications section directly, enabling clearer orientation to both the scope and the challenges of each application.

Each subsection introduces the core purpose and typical tasks associated with the respective application area, followed by an overview of relevant methods and a discussion of prevailing research challenges. To promote reader comprehension and seamless navigation, introductory and summary statements are included for each major domain, with improved transitions guiding flow between distinct applications.

All in-text citations employ finalized, traceable references with proper LaTeX formatting.

### 4.1 Application Domain 1: [Domain Name]

**Objective:** This subsection outlines the specific goals and significance of applying surveyed methods in [Domain Name], contextualizing the main tasks and expected outcomes for this domain.

[Insert overview of relevant methods and key results, ensuring all citations are correctly formatted and any placeholders are replaced with finalized references.]

**Outstanding Challenges:** Despite significant progress, major research challenges remain in [Domain Name], including [briefly categorize and explain the key unresolved issues].

### 4.2 Application Domain 2: [Domain Name]

**Objective:** This subsection introduces the purpose and practical importance of utilizing surveyed approaches in [Domain Name], providing a foundation for understanding domain-specific considerations.

[Provide domain overview, discuss notable methods, and ensure narrative transitions smoothly from the prior subsection.]

**Outstanding Challenges:** In this domain, unresolved research questions center around [explain and categorize main research challenges relevant to this application].

In summary, this section reinforces the survey’s overall objectives by systematically detailing both the established applications and future research directions for each domain. This approach ensures clarity and coherence for readers with diverse backgrounds and interests.

**Table 6: Comparison of quality assessment metrics on DIBR image datasets.**

Metric	Pearson	Spearman	Computation Time (s)
PSNR	0.62	0.58	0.15
SSIM	0.76	0.71	0.22
MP-PSNR	<b>0.92</b>	<b>0.86</b>	0.27
MW-PSNR	0.91	0.84	<b>0.19</b>

### 4.3 Computer Vision, Healthcare, and Scientific Discovery

The integration of advanced generative models and data augmentation strategies has significantly transformed domains such as computer vision, healthcare, and scientific discovery. These innovations not only yield improved model performance but also reshape conventional workflows.

In computer vision, semantic enrichment and generative techniques directly address persistent challenges in annotation efficiency and multimodal learning. Architectures exemplified by Detection with Enriched Semantics (DES) employ weakly supervised segmentation branches and global activation mechanisms to infuse high-level, class-aware context into object detectors. DES achieves high accuracy with minimal computational overhead and eliminates the need for additional manual annotations, directly alleviating the limitation of scarce ground-truth data prevalent in datasets such as PASCAL VOC and MS COCO [23]. This trend marks a departure from narrowly specialized methods, toward unified frameworks capable of embedding semantic priors into advanced detection pipelines with extensibility and minimal annotation burden.

In healthcare, synthetic data generation and generative modeling have become indispensable for circumventing hurdles posed by data scarcity, privacy constraints, and the demand for diverse labeled datasets. Simulation frameworks, notably SyntheX, employ physically accurate domain randomization and simulation methodologies to generate extensive annotated datasets, obviating the reliance on ethically fraught and costly real-world data acquisition. Evidence demonstrates that models trained solely on synthetic X-ray datasets, as produced by SyntheX, can match or exceed the performance of real-data-trained models in tasks such as landmark localization, segmentation, and detection within hip imaging and robot-assisted surgery [67]. Strategic approaches like synthetic-to-real (Sim2Real) training—particularly when bolstered by robust domain randomization—equal or surpass classical domain adaptation techniques, such as CycleGAN and ADDA, while entirely bypassing real data in the training loop. Nonetheless, these advancements are constrained by simulator fidelity and the risk of annotation mismatch, emphasizing the need for ongoing research in simulation realism and transferability [21, 67].

Beyond imaging, synthetic and augmented data are increasingly deployed in electronic health records (EHR) and tabular biomedical datasets to enhance the robustness and generalizability of deep learning models. Here, augmentation techniques span diverse modalities: geometric and pixel-level methods in vision, paraphrasing and noising in natural language processing, and sophisticated generative schemes in tabular or time-series contexts. Recent frameworks have unified augmentation taxonomies across image, text,

graph, and time-series data, organizing techniques by granularity and information type [4, 75, 86]. Blending single-instance, multi-instance, and generative augmentation methods (including GAN- and LLM-based data synthesis) consistently yields measurable gains in downstream model performance, affinity, and diversity. Despite these benefits, unresolved challenges endure. For instance, GAN- and diffusion-generated data risk diminished diversity, label inconsistencies, and problematic autophagic training loops—scenarios in which models are repeatedly exposed to their own synthetic outputs, leading to compounded distributional drift and undermining reliability [72, 85]. Consequently, there is a growing consensus that a balanced, meticulously curated combination of real and synthetic data is essential to preserve and enhance model utility [85].

Generative modeling also catalyzes breakthroughs in scientific discovery, especially in drug design and protein-ligand interaction modeling, where combinatorial complexity and limited empirical data challenge traditional experimental approaches. Cutting-edge models such as EquiScore, which employs equivariant graph neural networks, integrate physical priors with structural interaction data to outperform more than twenty established methods in binding pose prediction and analog activity ranking. These successes are underpinned by strategic augmentation, redundancy reduction, and interpretability mechanisms [45]. Similarly, the Geometry-Complete Diffusion Model (GCDM) facilitates the joint diffusion of atom coordinates and types, permitting the synthesis of valid, large molecules optimized for desired drug-like properties [56]. These innovations are further supported by advanced evaluation protocols and committed open-source dissemination, promoting adoption in pharmaceutical applications while illuminating an ongoing trade-off between computational requirements and chemical validity.

Collectively, these advancements signal a paradigm shift toward annotation-efficient, modality-agnostic, and context-aware generative AI applications, especially in environments characterized by data scarcity or ethical considerations. Nonetheless, responsible deployment calls for rigorous methodologies, standardized evaluation practices, and vigilant oversight to prevent data contamination, as the demarcation between real and synthetic data grows increasingly ambiguous [4, 21, 22, 56, 67, 70, 72, 75, 85, 86, 88].

### 4.4 Fine-Grained and Facial Synthesis

In line with this survey’s objectives—to characterize advances and challenges in generative modeling across diverse application domains—this section concentrates on fine-grained and facial image synthesis. These tasks have seen marked innovation, impacting digital creativity and clinical practice alike.

State-of-the-art facial generative networks now synthesize high-fidelity face images from sparse, incomplete, or disjointed patches,

leveraging composite loss functions that harmonize pixel-based, perceptual, and adversarial criteria. For example, Sun et al. (2022) [79] introduce a deep generative network conditioned on small facial parts, which integrates per-pixel reconstruction, perceptual, and adversarial losses, as well as total variation regularization for artifact suppression. The model accurately reconstructs faces by fusing features from distinct individuals, achieving superior plausibility and semantic coherence over traditional inpainting or patch-based methods when evaluated on CelebA, CACD, and LFW datasets (see Table 3 in [79]).

These generative models support diverse applications including digital artistry, virtual previews for plastic surgery and dentistry, and scalable, privacy-preserving generation of synthetic labeled data for face recognition and classification. The ability to manipulate attributes and styles within these models enables controllable synthesis, serving both research and creative purposes.

Nevertheless, significant technical and ethical challenges persist. Critical future directions include improved disentanglement of facial attributes, active mitigation of dataset biases, and the development of safeguards against misuse—especially in identity-sensitive or medical diagnostic contexts.

#### 4.5 3D Video and Virtual View Navigation

This subsection reviews recent advances and open challenges in generative and synthesis-based methods for 3D video and virtual view navigation, with an emphasis on technical pipelines, measurable improvements, and emerging research directions. The goal is to clarify how these techniques contribute to increased realism, scalability, and user immersion, and to evaluate their effectiveness in light of the survey's broader objectives on the state-of-the-art in generative 3D media processing.

Advances in generative and synthesis-based methods have greatly enhanced the realism, scalability, and user experience in 3D video and virtual view navigation, unlocking new possibilities for both creative media and immersive technologies.

In gigapixel-scale novel view synthesis, approaches utilizing meta-deformed manifold representations and implicit neural fields have demonstrated significant improvements in geometric correspondence and high-fidelity surface detail, outperforming traditional methods for complex, real-world and high-resolution scenarios [53]. For example, Li et al. (2019) [53] introduced a spatio-temporal virtual view synthesis method that leverages time domain information from video sequences to extract a static background and applies a weighted-fusion hole-filling technique informed by enhanced depth map processing (including edge detection, expansion, and Gaussian smoothing). These refinements specifically target common artifacts in Depth-Image-Based Rendering (DIBR), such as holes, ghosts, and temporal flicker, and result in measurable improvements; experiments on standard datasets (e.g., "BreakDancers," "Ballet") demonstrated superior scores in PSNR and SSIM for spatial quality and lower F-scores for temporal continuity compared to prior methods. Nonetheless, such pipelines generally require fixed camera geometries and remain computationally intensive, highlighting the need for scalable and adaptive approaches—especially for live telepresence, immersive broadcasting, and time-sensitive validation scenarios.

Quality assurance for generated 3D media increasingly depends on domain-specific, perceptually aligned metrics. Recent innovations such as Morphological Pyramid PSNR (MP-PSNR) and Morphological Wavelet PSNR (MW-PSNR) [77] apply multi-scale decompositions using non-linear morphological filters to preserve and assess edge-related geometric detail in DIBR-synthesized images. According to Sandić-Stanković et al. (2016) [77], these metrics, evaluated on benchmarks like IRCCyN/IVC DIBR and MCL-3D, achieve significantly higher correlation with subjective human judgments than conventional metrics, with MP-PSNR notably yielding Pearson's correlations up to 0.904 on IRCCyN/IVC DIBR. Reduced variants of these metrics maintain robustness while minimizing computational requirements, making them well-suited to practical, high-throughput 3D video applications.

Despite these advances, several open challenges persist: 1. Scalability and efficiency for processing high-resolution or real-time content, particularly for novel view synthesis from unconstrained camera setups. 2. Improved robustness to scene complexity, non-rigid motion, and challenging lighting conditions in real-world environments. 3. Development of objective, reference-free quality assessment metrics that can reliably predict perceptual quality in diverse scenarios, reducing dependence on full-reference methods. 4. Greater interoperability of generative modules with transmission, compression, and editing pipelines for efficient integration into production workflows.

In summary, recent progress in generative modeling and quality assessment—exemplified in advances to both synthesis pipelines and perceptually-aligned metrics—set new benchmarks for next-generation 3D video systems. However, continued research is required to address open scalability, robustness, and automation challenges, in order to fully realize seamless, high-fidelity, and adaptive 3D media experiences across practical application domains.

#### 4.6 Industrial, Scientific, and Emerging Applications

The application of synthetic data and generative modeling extends well beyond conventional vision and biomedical domains, enabling a diverse array of industrial and scientific advancements with profound societal ramifications. In drug discovery and computational chemistry, architectures such as diffusion-based and equivariant generative models (e.g., EquiScore and GCDM) have established themselves as foundational tools for molecular design, property optimization, and virtual screening. Their integration with advanced chemical priors facilitates accelerated discovery cycles and substantial reductions in experimental expense, while also affording interpretability and adaptability for specialized tasks ranging from analog ranking to protein-conditional generation. Yet, maintaining the balance between geometric complexity and computational efficiency remains a persistent technical challenge [45, 56].

Generative and augmentation-based approaches have similarly transformed environmental and climate modeling. Consistency models trained on fine-grained observational data not only reduce computational downscaling costs, but also provide uncertainty-calibrated, bias-corrected forecasts that generalize across diverse climate scenarios [84]. These methodologies circumvent the need for retraining per simulation, delivering rapid probabilistic predictions

that preserve key climatological patterns and extremes—integral for policy-making and risk assessment under evolving global conditions.

Distributed data creation—fundamental in federated and edge learning environments constrained by privacy and infrastructure—leverages generative models as adaptive meta-models that can be rapidly customized to local data. By formalizing continual adaptation as a Wasserstein barycenter problem and integrating quantization-aware compression, recent frameworks enable bandwidth-efficient, privacy-preserving model personalization and collaborative learning without direct transmission of sensitive real-world data [37]. However, such systems demand precise calibration to prevent data drift and overfitting, as edge devices typically possess limited and potentially unrepresentative datasets.

In summary, the convergence of synthetic data, generative modeling, and simulation-driven innovations has triggered breakthroughs across traditionally segregated disciplines, while simultaneously raising complex technical and ethical challenges that necessitate enhanced methodological rigor and cross-domain collaboration as these approaches transition toward wide-scale, high-impact deployments [37, 45, 56, 84].

Table 7 provides a structured comparison of representative generative and simulation-driven methods, highlighting their principal features and domain-specific achievements.

Together, these ongoing advances reinforce the need for methodical evaluation, careful curation of synthetic data, and sustained interdisciplinary cooperation to maximize the positive impact of generative modeling across a burgeoning array of applications.

## 5 Thematic Synthesis, Evaluation, and Benchmarking

This section synthesizes key themes identified across surveyed works, provides a critical evaluation of prominent methodologies and their comparative performances, and discusses benchmarking practices adopted in the domain. The purpose is to offer a consolidated perspective that enables readers to assess progress, challenges, and trends within the field. Each of the following subsections introduces and contextualizes a particular application area or evaluation axis, ensuring clarity regarding its specific aim and relevance within the overall survey framework. Transitions between related themes are provided to assist seamless navigation across distinct but interconnected foci in the literature.

### 5.1 Cross-Sectional Method Comparisons

The field of synthetic data generation is undergoing a rapid evolution, with generative paradigms such as Generative Adversarial Networks (GANs), diffusion models, and hybrid methodologies addressing distinct requirements regarding realism, controllability, and domain-specific augmentation. This progression is particularly pronounced within vision and healthcare applications, as documented in recent literature [21–23, 34, 48, 67, 69, 71–76, 82, 86, 88]. Traditionally, GANs have been regarded as the benchmark for high-fidelity image synthesis. Innovations in their architecture—such as the introduction of auxiliary classifiers and the utilization of contrastive learning—have significantly mitigated challenges like mode collapse and improved semantic alignment. These advancements

enable task-specific performance enhancements, particularly in the realms of subclass differentiation and fine-grained, text-conditioned synthesis [34, 71, 72, 76]. A notable advancement is the FG-RAT GAN, which excels at generating class-consistent and diverse images, while reducing both model complexity and resource consumption compared to models such as LAFITE and VQ-Diffusion [34]. FG-RAT GAN leverages an auxiliary classifier in the discriminator to guide the generator for fine-grained semantic consistency and incorporates a contrastive learning approach with a cross-batch memory mechanism to maximize similarity within the same subclass while separating different subclasses. Its total loss formulation integrates adversarial, categorical cross-entropy, and contrastive components. Experimental results on benchmarks like CUB-200-2011 and Oxford-102 indicate that FG-RAT GAN achieves state-of-the-art Frechet Inception Distance (FID) and competitive Inception Scores (IS), outpacing LAFITE, RAT GAN, and diffusion-based models, while also using fewer parameters and less compute. Ablation studies underline the importance of both the auxiliary classifier and the contrastive learning strategy [34].

Diffusion models, in turn, have established new state-of-the-art benchmarks, regularly surpassing GANs in terms of sample quality, diversity, and robustness. They are especially effective in conditional and multi-modal settings, owing to their iterative denoising processes and advanced guidance mechanisms, such as classifier-free and entropy-based sampling [69, 75, 76, 86]. Recent diffusion approaches like IIDM [71] highlight a new perspective by recasting semantic image synthesis as progressive denoising in latent space, allowing better satisfaction of constraints such as semantic segmentation and style. For instance, IIDM encodes style images into a latent vector, initializes diffusion with noise, and performs iterative denoising guided by segmentation masks, with further inference refinements (refinement, color transfer, ensemble) enhancing image fidelity without extra training. IIDM achieves higher mask accuracy and lower FID scores compared to both GAN- and diffusion-based predecessors, emphasizing the ability to balance semantic content and style similarity. However, diffusion models' improvements also come with increased computational demands, and in certain contexts—particularly medical imaging—may introduce heightened risks of memorization, especially in limited data regimes [48, 74]. Such concerns necessitate careful, joint assessment of privacy and utility to ensure responsible deployment.

Hybrid approaches, which incorporate components such as transformers, variational autoencoders (VAEs), or attention mechanisms within GAN or diffusion frameworks, have demonstrated quantitative gains in structural fidelity and multimodal controllability, particularly within both medical and general vision domains [22, 34, 67, 75]. For example, transformer-augmented pipelines or attention modules foster improved high-dimensional input handling and allow for more precise control in conditional generation scenarios, supporting stronger cross-modal generalization.

A central consideration when comparing these approaches is their respective resource efficiency and scalability. GANs typically provide high throughput and low inference cost, which facilitates deployment, though they may encounter issues with scalability and struggle with complex conditional synthesis due to bottlenecks in stability and diversity [23, 71, 82]. Diffusion models offer superior synthesis quality and controllability, but at the cost of

**Table 7: Representative Generative and Simulation Methods in Selected Application Domains**

Domain	Prominent Method(s)	Key Features and Achievements
Computer Vision	DES [23]	Weakly-supervised segmentation; semantic enrichment; reduces annotation cost with minimal computational overhead.
Healthcare	SyntheX [67]	Physically realistic simulated X-ray datasets; supports landmark localization and detection; eliminates need for real images in training.
Scientific Discovery	EquiScore [45], GCDM [56]	Equivariant graph neural networks and diffusion models; excel in molecular property optimization and binding prediction.
3D Video	Meta-deformed manifold, DIBR [53, 77]	Improved geometric correspondence, scene realism; novel QA metrics (Wavelet/Morphological PSNR) for artifact-sensitive deployment.
Distributed Learning	Wasserstein barycenter adaptation [37]	Model personalization without real data transmission; bandwidth and privacy optimized; supports federated/edge scenarios.
Environmental Modeling	Consistency models [84]	Bias-corrected, high-resolution outputs; efficiency without retraining for each simulation input.

greater training and inference resource requirements. Nevertheless, emerging research on acceleration—such as optimized noise schedules and architectural efficiency—continues to narrow the performance gap [69, 75]. Hybrid models, particularly those utilizing transformer and attention mechanisms, often achieve a balance, supporting cross-modality generalization, effective processing of high-dimensional data, and more refined user control [22, 34, 67].

In summary, ongoing methodological innovation in synthetic data generation seeks to optimize fidelity, diversity, explicit control, and computational feasibility, all while carefully accommodating domain-specific constraints and downstream application requirements.

## 5.2 Evaluation of Synthetic Data Quality

Rigorous evaluation of synthetic data quality must address technical advancement and their intertwined ethical and practical concerns—a particularly important link for sensitive domains such as healthcare [74][27][85]. Three core objectives frame this challenge: factuality, fidelity, and fairness. Each not only ensures authenticity and practical utility but also plays a crucial role in the responsible deployment of generative methods.

**Factuality** denotes the consistency between generated instances and ground-truth domain logic or factual knowledge. **Fidelity** encompasses both statistical and perceptual similarity between synthetic and real data. **Fairness** underscores the need to minimize bias and disparate impact, which remains a critical consideration in healthcare and other high-stakes tasks. These axes are technically motivated but directly affect downstream trust, regulatory compliance, and social responsibility. Linking these criteria with practice, robust evaluation protocols double as safeguards against misapplication, model memorization, or amplifying social biases [27][74].

A broad suite of metrics is employed to quantify these dimensions. For vision and imaging, widely established metrics include mean Average Precision (mAP), mean Intersection over Union (mIoU), Fréchet Inception Distance (FID), Maximum Mean Discrepancy (MMD), Inception Score (IS), and ROC-AUC. Affinity and diversity ratios further help detect overfitting and assess population variability [59][40][57][64][15][51][33][8][3][63][65][70][61][11][75][83][27][72][69][48][34][71]. Each metric illuminates distinct quality aspects; for instance, mIoU and mAP characterize segmentation and detection accuracy, FID and IS measure perceptual and feature-level diversity, while affinity and diversity ratios indicate generalizability [75].

However, strong performance on standard fidelity metrics alone may conceal underlying risks. For example, diffusion models may achieve favourable FID and IS scores but inadvertently reproduce memorized training examples, amplifying privacy or fairness concerns—particularly for homogeneous or small datasets in regulated

domains [8][74][48]. These ethical and practical issues are inseparable from technical evaluation and motivate the adoption of comprehensive assessment protocols—including adversarial, privacy-oriented, and domain-adaptive tests—which more holistically appraise risks and downstream impact [74][27][57].

Such protocols typically involve membership inference, reidentification risk analysis, and domain-adapted performance benchmarking. By integrating these with traditional metrics, evaluations can transcend surface similarity and foster confidence that synthetic data is not only accurate but also responsibly generated and deployable [27][85].

Standardized benchmark datasets are pivotal for ensuring comparability and relevance across methods and domains. For vision, COCO, Pascal VOC, ADE20K, Cityscapes, CUB-200-2011, and Oxford-102 are prominent. Domain-specific resources include PDBscreen (structural biology), SyntheX (clinical X-ray), and a variety of medical imaging corpora. The diversity among these datasets enables the measurement of both generic perceptual quality and context-specific features, such as phenotype fidelity, class balance, and clinical applicability [65][70][61][18][22].

In healthcare and similarly sensitive fields, advanced validation protocols are in active use. Examples include task-based evaluation (such as TSTR—Train on Synthetic, Test on Real) and sophisticated statistical similarity measures (including KL divergence and the KS statistic) to further contextualize the translation of synthetic data into practice [59][74][27][85].

### Summary of Future Research Gaps in Evaluation of Synthetic Data Quality:

- Establishment of standardized, domain-tailored evaluation protocols, especially for clinical and regulated domains [85][27]
- Enhanced methodologies for robust detection of memorization, bias, and privacy risks in generative models [8][74]
- Improved alignment between technical metrics and actual downstream utility, fairness, and safety in real-world deployments [27][59]
- Creation and open sharing of diverse, high-fidelity benchmark datasets covering a wider spectrum of modalities and application contexts [85][22]
- Advancement of interpretability and human-centered validation techniques that link quantitative evaluation to ethical and regulatory considerations [74]

## 5.3 Detailed Comparative Tables

To elucidate the progression of text-to-image generative models and provide a clear, side-by-side assessment of their capabilities, leading methods—including LAFITE, VQ-Diffusion, RAT GAN, and FG-RAT GAN—are compared across axes such as sample realism

(FID), feature diversity (IS), parameter count, and resource demands [34]. This summary is structured in Table 8.

As illustrated in Table 8, FG-RAT GAN achieves state-of-the-art FID and IS metrics while maintaining the lowest parameter count, indicating a substantial advance in both sample realism and computational efficiency [34]. Notably, FG-RAT GAN distinguishes itself from prior approaches by integrating an auxiliary classifier into the discriminator for class-wise image classification and incorporating a contrastive learning mechanism using a cross-batch memory. This dual strategy improves subclass-awareness and sharpens intra-class and inter-class distinctions. The corresponding loss functions combine adversarial, categorical cross-entropy, and contrastive losses, further driving model performance. Empirical studies on benchmark datasets confirm that these innovations yield higher semantic and visual fidelity compared to earlier models, underscoring rapid progress in subclass-aware text-to-image generation.

## 5.4 3D/DIBR-Specific Metrics

This subsection focuses on the evaluation of synthesized 3D data and depth-image-based rendering (DIBR) outputs, aiming to address unique challenges not captured by conventional 2D metrics. Our objective is to present methodological advances in assessing 3D visual quality, especially regarding the preservation of edges and structures near disoccluded regions, and to discuss open challenges as well as the alignment with broader survey goals of rigorous and context-sensitive benchmarking.

The synthesis and evaluation of 3D data, including outputs generated by DIBR, introduce distinct requirements not addressed by conventional 2D measures. Artifacts such as edge distortions and geometric deformation—particularly near disoccluded regions—necessitate specialized evaluation. Morphological Pyramid Peak Signal-to-Noise Ratio (MP-PSNR) and Morphological Wavelet Peak Signal-to-Noise Ratio (MW-PSNR) have been introduced to directly target edge preservation and structural coherence across scales [77]. These metrics utilize non-linear morphological filters in multi-scale decompositions to quantify distortions and emphasize artifact-prone edges, resulting in significantly higher correlation with human subjective scores compared to previous state-of-the-art metrics. Notably, reduced MP-PSNR achieved Pearson’s 0.904 and Spearman’s 0.863 on the IRCCyN/IVC DIBR database [77], while retaining computational efficiency and requiring minimal parameter optimization.

In practical settings—such as real-time streaming, immersive 3D simulation, and medical imaging—these metrics are favored due to their low computational overhead and ability to work without image registration, validating their utility for diverse, time-sensitive 3D applications.

Despite the integration of GANs, diffusion models, hybrid strategies, and transformer-based architectures, persistent challenges in the 3D/DIBR domain include: (1) underutilization of patient-specific or scene-specific information for more personalized or context-aware synthetic data [85]; (2) a shortage of standardized, domain-adapted benchmarking protocols that impede reliable evaluation and comparison across methods [74, 85]; (3) ongoing issues of model memorization, fairness evaluation, and the lack of universal cross-domain benchmarks [27, 74].

In this context, continued innovation in both generative methodology and quality assessment remains essential for the responsible advancement of 3D synthesis, with actionable opportunities including the development of dedicated benchmarking standards, adoption of more context-driven evaluation criteria, and closer integration between quality metrics and application objectives, as emphasized by recent work across synthetic data and 3D rendering [74][27][85][72][69][48][34][71].

In summary, the advancement of 3D/DIBR-specific metrics like MP-PSNR and MW-PSNR has significantly improved the capacity to align technical evaluations with human perception in 3D domains. However, bridging the gap between metric development, contextual appropriateness, and robust benchmarking protocols remains a defining challenge and an important direction for future research. This aligns with our survey’s broader aims of fostering rigorous, transparent, and domain-sensitive quality assessment frameworks for generative AI models.

## 6 Responsible and Ethical Oversight

Technical advancements in AI necessitate careful consideration of responsible and ethical oversight to ensure that innovations align with societal expectations and legal frameworks. Linking recent technical developments to their practical and ethical implications is crucial for anticipating and mitigating potential risks. For example, increased model complexity may lead to challenges in transparency, which can, in turn, affect users’ ability to trust AI-powered decisions. Therefore, technical strategies for interpretability must be evaluated not only for their engineering merit but also for their effectiveness in supporting accountability and fairness. This reinforces the need for evaluation metrics that holistically capture both quantitative model performance and qualitative ethical considerations.

Further, responsible oversight mechanisms should be adaptable to the rapid pace of AI progress. Practical concerns such as potential for bias amplification, data privacy breaches, and unintended social consequences must be linked directly to the technical characteristics of deployed systems. Such explicit integration helps in designing governance frameworks that proactively reflect on both present and foreseeable challenges.

**Summary of Future Research Gaps:** There is a significant need for unified frameworks that bridge technical assessments with ethical evaluation processes. Current oversight methods often lag behind new AI capabilities, underscoring the demand for adaptive governance mechanisms. More comprehensive metrics are required that simultaneously evaluate system performance and responsible behavior. Standardization of reference formats and consistency in documentation are ongoing challenges that impact clarity and trust in published results.

### 6.1 Ethical and Social Issues

The proliferation of synthetic data and generative artificial intelligence (GenAI) in scientific and healthcare research presents both significant opportunities and considerable ethical challenges. Of foremost concern is data privacy: traditional anonymization measures are increasingly vulnerable to re-identification, particularly as adversarial capabilities grow and auxiliary data sources become

Table 8: Comparison of leading text-to-image generation models on major metrics.

Model	FID ↓	IS ↑	Parameters (M)	Efficiency Features
LAFITE	18.4	27.4	174	Text-conditional synthesis, high complexity
VQ-Diffusion	15.1	25.6	123	Diffusion-based, stable training
RAT GAN	12.3	29.2	110	Attention-based, auxiliary classifiers
FG-RAT GAN	<b>11.5</b>	<b>30.6</b>	<b>52</b>	Contrastive loss, parameter efficient, cross-batch memory, auxiliary classifier

ubiquitous. This escalating risk has catalyzed the adoption of synthetic data as a privacy-enhancing approach [1][64][15][43][33]. State-of-the-art generative models—including generative adversarial networks (GANs), variational autoencoders (VAEs), diffusion models, and large language models (LLMs)—have advanced the ability to create artificial datasets that preserve the statistical properties of the original data while reducing direct exposure of individual records. Such innovations mitigate regulatory constraints imposed by regimes like the General Data Protection Regulation (GDPR) and the Health Insurance Portability and Accountability Act (HIPAA), and accelerate data sharing for machine learning (ML) research [59][57][60][49][52][15][43][8].

Nevertheless, these developments do not render synthetic data immune to privacy risks. Techniques such as membership inference and linkage attacks remain viable, especially when synthetic data generation is not supplemented with formal differential privacy (DP) guarantees or comparable noise infusion mechanisms [64][15][43]. The privacy-utility trade-off is particularly acute in scientific and healthcare domains, as high-fidelity models may inadvertently reproduce features or outliers that enable sensitive information leakage [1][64][33]. Empirical studies indicate that while DP can reduce privacy risks, it often does so at the expense of downstream ML task performance if not implemented judiciously [64][43].

Algorithmic bias represents a further axis of ethical concern. Synthetic datasets generated from skewed or unrepresentative source data—or from generative models that inadvertently encode social or demographic disparities—may perpetuate or amplify these biases in downstream analytics and decision-making systems [60][1][15][8][88]. This challenge is especially salient when synthetic data is used to augment minority classes, correct imbalances, or fill missing modalities. As such, diligent post-generation audits and systematic fairness assessments are imperative to mitigate unintended harm [55][60][88].

The increasing indistinguishability of real and synthetic data introduces additional risks related to misinformation and data integrity. With GenAI models achieving greater fidelity, the line between authentic and synthetic datasets becomes blurred, threatening reproducibility, scientific credibility, and public trust [28][62][8][11]. Synthetic data, if undocumented, can corrupt research repositories and benchmarks, or propagate hallucinated artifacts that compromise scientific inference. To address these risks, robust provenance frameworks, mandatory transparency and disclosure, and explicit labeling protocols have been advocated to ensure that any use of synthetic data in research outputs is clearly identified [28][8].

In this context, transparency and traceability are fundamental requirements for responsible synthetic data deployment. Technical

measures—including data watermarking, cryptographic blockchain certification, and AI-driven detection tools—are emerging to ensure the provenance and audibility of synthetic datasets throughout the research lifecycle [28][8][11]. Such technologies, however, demand standardization to prevent fragmentation of practices and to achieve broad adoption.

Legal and regulatory compliance represents an evolving frontier. Synthetic data facilitates compliance with privacy legislation by reducing identifiability, but the lack of universal, precise regulatory guidelines continues to create uncertainty [40][60][36][52]. Regulatory agencies are increasingly alert to such nuances, particularly in clinical research and data-sharing contexts; nonetheless, their approaches diverge across jurisdictions. Sustained, cross-sector dialogue is necessary to establish consensus guidelines that appropriately balance individual protection with scientific progress [57][60][15][26][8][27].

6.2 Responsive Practices and Protocols

As ethical and social risks related to synthetic data are increasingly recognized, proactive practices and protocols have emerged to mitigate these potential harms. Chief among these are systematic risk mitigation strategies designed to anticipate, detect, and reduce privacy breaches, algorithmic bias, and misinformation [22][88][27]. Comprehensive auditing—conducted both internally and by external parties—serves to evaluate privacy leakage, fairness, and statistical fidelity, as well as to monitor for residual biases. Privacy audits often utilize membership inference testing, k-anonymity assessments, and advanced adversarial simulations, while fairness evaluations are grounded in domain-specific metrics and comparisons to representative baselines [64][22][88].

A cornerstone of responsible practice is the multi-faceted validation of synthetic data quality and representativeness. Cutting-edge protocols require evaluations that go beyond high-level statistical comparisons, mandating statistical similarity metrics (such as distributional divergence measures and correlation structure analyses), machine learning utility tests (for example, Train on Synthetic, Test on Real, and downstream task generalization), and domain-specific relevance assessments, often with expert input.

This multi-level approach ensures both the meaningful utility of synthetic datasets and the avoidance of artefacts that could distort scientific interpretation [55][49][21][22][85].

Furthermore, transparent community benchmarking and open validation are instrumental for trustworthy synthetic data deployment. The proliferation of open-source tools, public leaderboards, and collaborative challenges encourages the adoption of shared best practices, enables rapid identification of methodological weaknesses, and harmonizes evaluation efforts [88][27][85]. Iterative

engagement among data scientists, clinical experts, ethicists, and legal advisors informs the operationalization of safeguards and continuous refinement of mitigation strategies as technology evolves.

### 6.3 Standardization and Community Initiatives

The sustainable and trustworthy integration of synthetic data in research and applied settings hinges on the advancement and widespread adoption of rigorous, standardized evaluation protocols and community-driven frameworks [27][85]. Currently, the lack of universally accepted, domain-tailored benchmarks and validation processes substantially hinders both scientific legitimacy and regulatory approval, an issue that is particularly acute within sensitive sectors such as healthcare [55][60][27][85].

Several coordinated community efforts are emerging to close these critical gaps. These initiatives include the establishment of frameworks for systematic privacy risk quantification, robust utility and fairness benchmarking, and comprehensive auditing protocols, allowing for transparent and trustworthy evaluation of synthetic data and generative models [22][88][27][85]. There is also growing support for precompetitive consortia and open benchmarking challenges, which serve to harmonize evaluation practices, lower entry barriers, and disseminate up-to-date best-practice guidance grounded in contemporary ethical and regulatory standards [85].

Further priorities include the advancement of transparency measures and provenance controls to reduce the risk of dataset contamination and untracked generative processes in downstream scenarios. Sustained, interdisciplinary dialogue is being fostered—bringing together technical, ethical, and regulatory experts—to ensure that standards can adapt responsively to advances in GenAI capabilities [11][21][22][27]. These dialogues are crucial for maintaining public trust and preventing ethical pitfalls, as underscored by pressing risks such as model autophagy and unintended data leakage [60][88].

In sum, the development of governance structures and consensus-driven standards—with cross-sector vigilance and continuous protocol refinement—remains essential to realizing the full promise of synthetic data in an equitable, ethical, and trustworthy fashion.

## 7 Challenges, Limitations, and Future Directions

Recent technical advancements have accelerated the capabilities of AI systems; however, these developments are inextricably linked with a complex set of ethical, practical, and societal challenges. Progress in areas such as scalability, training efficiency, and multi-modal integration often comes alongside increased risks in fairness, privacy, and reliability. For example, improvements in language model generation must be weighed against potential risks of misinformation propagation and bias amplification.

To fully appreciate the interplay between technological progress and these broader concerns, it is crucial to analyze how enhancements in model architectures or training pipelines may give rise to novel ethical dilemmas or exacerbate existing practical limitations. Addressing these issues demands a holistic approach that aligns technical objectives with responsible deployment practices, ensuring systems remain robust, transparent, and aligned with human values.

At the same time, evaluating and mitigating the risks introduced by new methodologies requires a critical examination of current evaluation pipelines, as well as the development of novel metrics and frameworks sensitive to downstream impacts. Technical innovations should, therefore, be considered alongside strategies for monitoring ethical risk flows in real-world applications.

Despite significant progress, a number of open research gaps remain evident, including:

Going forward, it is essential that the AI research community fosters collaboration across disciplinary boundaries, pursues unified standards for evaluation and ethical oversight, and remains vigilant in translating technical advancements into responsible innovation.

### 7.1 Technical and Resource Barriers

The proliferation of diffusion models and other large-scale generative frameworks has underscored numerous technical and resource-intensive obstacles that constrain scalability and impede equitable access. A chief concern is the substantial computational burden required for both the training and inference phases of diffusion models, which significantly exceeds that of previous generative paradigms, such as GANs—even in light of recent algorithmic advances aimed at improving efficiency [2–4, 9, 12, 13, 17, 18, 20, 22, 23, 25, 34, 38, 41, 46, 47, 51, 56, 58, 61, 63–65, 67, 70–72, 75, 84, 86, 88–90]. Training such models necessitates vast datasets and computational infrastructure, resources predominantly accessible to well-funded organizations, thereby accentuating disparities in research and deployment [4, 9, 12, 41, 64, 65, 72, 86]. Strategies including modular sampling [90], adaptive noise scheduling [34], and hybrid generative architectures are being pursued to accelerate inference and reduce resource consumption. Nevertheless, the fundamental tradeoff between model expressivity, fidelity, and computational feasibility persists [3, 12, 46, 72].

Equally pressing are the challenges related to the fidelity and cross-domain transferability of simulation-generated or synthetic data. The disparity between simulated and real-world data—stemming from rendering engine limitations, incomplete physical modeling, and non-standardized annotation—introduces domain biases that can impair downstream performance [22, 72, 76, 85]. Recent advances in highly realistic simulation frameworks, domain randomization, and physics-based generative models have improved transferability [76, 85]; however, achieving strong cross-modality consistency, particularly in 3D or multi-modal settings, remains restricted by representation instability and complex alignment issues [22, 72, 76]. Furthermore, the absence of standardized annotation protocols for synthetic data—especially in intricate, high-dimensional domains—complicates model evaluation, reproducibility, and regulatory scrutiny.

As shown in Table 10, targeted algorithmic advances partially address these technical constraints; however, resource inequalities and deficiencies in standardization continue to challenge both scalability and reproducibility.



Table 9: Summary of Key Future Research Gaps

Challenge Area	Current Limitation	Future Direction	Potential Impact
Scalability	Resource-intensive training	Efficient algorithms and model compression	Wider accessibility
Ethical Alignment	Insufficient bias and fairness evaluation	Integrated ethical auditing tools	Minimizing societal harm
Evaluation	Lack of standardized benchmarks	Development of unified evaluation frameworks	More reliable model assessment
Privacy	Inadequate data protection mechanisms	Privacy-preserving learning approaches	Enhanced user trust
Transparency	Opaque model decision processes	Advanced interpretability techniques	Improved accountability

Table 10: Technical barriers and mitigation strategies for large-scale generative models

Barrier	Example Manifestation	Mitigation Strategies
High computational cost	Training/inference of large diffusion models	Modular sampling, adaptive noise scheduling, hybrid architectures
Data bias in simulation-generated samples	Domain gap in synthetic-to-real transfer	Domain randomization, physics-based generative models
Lack of annotation standards	Heterogeneous labels in 3D/multi-modal data	Development of universal standards, community benchmarks

7.2 Generalization, Robustness, and Societal Impact

This section aims to critically evaluate the dual challenges and broader implications associated with generative models: their technical generalization and robustness, especially under constraints such as limited, federated, or non-stationary data, as well as their far-reaching societal and ethical ramifications. By the end of this section, the reader should understand core obstacles limiting generative model reliability and the contextual risks posed by their widespread deployment.

A pivotal issue in generative modeling concerns the extent to which current architectures generalize to low-resource, continual, federated, and edge environments—contexts characterized by limited data, computational constraints, and evolving distributions [37, 84]. For instance, in continual learning at the edge, Dedeoglu et al. [37] demonstrated how model updates using Wasserstein-based approaches can help mitigate catastrophic forgetting. Similarly, Hess et al. [84] report on downscaling climate models, highlighting knowledge transfer, uncertainty quantification, and adaptation challenges in real-world geoscience settings. These case studies reveal that many existing methods degrade in fidelity or exhibit vulnerability when aggregated under federated setup or continually updated, signaling unresolved gaps in knowledge transfer, adaptation, and privacy protection [84]. Moreover, as illustrated in 3D drug design with geometry-complete diffusion models [56], the complexity and resource demands of generative models pose deployment barriers in edge or restricted settings. The deployment of large generative models onto resource-constrained systems therefore faces persistent obstacles pertaining to model size, communication overhead, and stable operation under out-of-distribution (OOD) conditions, thereby underscoring the demand for modular, efficient, and adaptable frameworks [27, 56, 84].

From a societal and ethical perspective, the deployment of synthetic data and generative models introduces a multifaceted array of risks, with real-world consequences documented across diverse application domains. For example, in medical imaging, Akbar et al. [8] and Tudosi et al. [11] both provide evidence that generative models may inadvertently memorize sensitive training data, complicating privacy guarantees even as synthetic images achieve

state-of-the-art quality. SyntheX [22] shows how synthetic datasets, if carefully generated, allow for generalizable and equitable learning in healthcare, but raises questions of bias remediation and representative coverage.

**Bias and fairness:** Synthetic data may encode and exacerbate existing societal and demographic biases, particularly if training corpora are imbalanced or lack domain specificity [4, 8, 22, 63, 66]. For instance, fairness concerns in clinical diagnostics have arisen when synthetic data underrepresent rare subtypes or minority populations [4].

**Contamination and feedback loops:** Recursive use of synthetic data as training input (AI “autophagy”), as detailed by Xing et al. [88], can erode model reliability and scientific validity, especially when web-scraped or uncured datasets accumulate unlabeled synthetic content [13, 14, 30].

**Privacy leakage:** Overparameterized models, especially trained on limited or homogeneous data, may memorize and inadvertently reveal sensitive real information, as shown in diffusion-based medical image generation [8, 22, 27, 72, 88].

**Opacity and interpretability:** The increasing complexity of generative architectures impedes interpretability and transparency—critical where alignment with human expectations and regulatory standards is fundamental (e.g., healthcare, law) [10, 11, 13, 22, 36, 42, 83, 88]. Techniques such as ConSCompF [83] and precedent-based interpretability [14] exemplify practical steps toward explainability in high-stakes environments.

The dynamic and evolving regulatory environment further compounds these challenges, necessitating robust mechanisms for fairness auditing, privacy preservation, and lifecycle risk management across deployment contexts. Practical frameworks and studies discussed in this section ground the often abstract risks in concrete examples, reinforcing the necessity for transparent, responsible, and adaptable deployment of generative models.

In summary, this section has identified the technical and societal limitations facing scalable generative model deployment, referencing concrete applications to highlight both progress and persistent gaps. This evaluation sets the stage for later discussions on best practices and prospective research directions toward inclusive, trustworthy, and resilient generative AI.

### 7.3 Label Dependency and Semantic–Style Balance

The imperative to minimize dependence on costly manual annotation has propelled extensive research into weakly supervised, self-supervised, and label-free generative paradigms [34, 71]. Despite this, leading conditional diffusion and GAN-based image synthesis methods often retain some level of label dependency to ensure semantic specificity and controllable generation [34, 71]. A key ongoing challenge is achieving a balance between semantic fidelity—requiring generated samples to unambiguously satisfy specified label or mask constraints—and stylistic diversity, which is crucial for enhancing generalization and practical utility [71]. Existing strategies such as latent space manipulation, complex conditioning on semantic labels or reference style images [34, 71, 85], and innovations like disentanglement losses or dual-guidance mechanisms only partially address this tradeoff.

Recent work, exemplified by FG-RAT GAN [34], shows that integrating auxiliary classifiers and contrastive learning objectives can improve class-consistent generation and fine-grained semantic control, but such approaches still rely on fine-grained labels and highlight the computational efficiency and accuracy tradeoff compared to fully label-independent models. Separately, diffusion-based frameworks like IIDM [71] explicitly disentangle semantic and style control by leveraging semantic masks and reference images. IIDM demonstrates that latent-space conditioning and inference refinements can yield high mask accuracy and style fidelity simultaneously, advancing the field closer to a robust solution that minimizes annotation requirements. However, despite improvements in specific metrics such as FID and mask accuracy, there remains a lack of a unified framework offering interpretable, fine-grained control with minimal label supervision. Developing such a framework for semantic-style balanced synthesis, as indicated by these recent advances, continues to be a central open challenge.

### 7.4 Evaluation Metrics and Benchmark Gaps

The evaluation of synthetic and augmented data—especially in nascent frontiers such as multi-modal, 3D, and cross-domain synthesis—remains constrained by the lack of universally accepted, modality-spanning metrics [2, 22, 58, 76, 85, 90]. Widely adopted quantitative statistics (e.g., FID, IS, TSTR) [2, 58, 90] are limited in their ability to capture semantic consistency, clinical significance, or application-specific value, particularly in sensitive fields where the availability of genuine data is restricted. This gap is emphasized in recent surveys and empirical studies, which highlight how these limitations impede rigorous cross-model and cross-study comparisons and slow the translation of synthetic data into practical and clinical applications [22, 76, 85]. For example, SyntheX demonstrates that synthetic data can closely match or even surpass real data for downstream machine learning tasks in medical imaging, yet its evaluation underscores the shortcomings of current metrics in reflecting clinical relevance and generalizability [22]. Reviews note an urgent need for robust, domain-tailored, and interpretable benchmarks to facilitate the responsible assessment and adoption of synthetic data, especially in multi-modal and multi-domain contexts [85]. Furthermore, recent advances in measurement-guided generative models for medical image synthesis stress that prevailing

metrics do not sufficiently capture the diagnostic utility or reliability of synthetic outputs [76], underscoring the necessity of developing comprehensive, context-aware evaluation paradigms. Addressing these gaps is essential for transparent model development, effective benchmarking, and regulatory approval.

### 7.5 Research Opportunities and Future Trends

Looking ahead, future research efforts should prioritize advancements in model quality, sample diversity, robust personalization, and real-time synthesis—imperatives as generative AI evolves toward dynamic, interactive, and agent-based applications [27, 34, 71, 72, 85]. Key enablers of this evolution include automated benchmarking and pipeline development for systematic performance tracking [27, 85], adaptive self-improving model architectures capable of continual learning and transfer, and scalable, modality-agnostic evaluation frameworks that facilitate rigorous cross-study comparability. Responsible AI practices must encompass updated policies, oversight mechanisms, and inclusive stakeholder engagement to secure ethical and trustworthy deployment [1, 4, 10, 11, 13, 17, 22, 24, 27, 30, 31, 42, 43, 49–51, 59, 63, 66, 83, 85, 88].

The convergence of vision, language, multimodal, and scientific domains is fostering the development of unified generative systems that support a wide array of applications—from scientific modeling and systematic reasoning to open-vocabulary detection and synthesis [1–6, 8–13, 15, 17, 18, 20, 22, 24, 25, 28–31, 33, 35, 36, 38, 40–45, 50–52, 54, 55, 57–66, 70, 80, 81, 88–90, 92, 93]. A decisive challenge at this intersection is the design of principled, scalable, and interoperable architectures that can effectively manage heterogeneous data modalities, diverse annotation schemes, and domain-specific requirements while minimizing bias and maximizing transparency. Progress in these areas will further hinge on interdisciplinary collaboration, uniting foundational research, applied expertise, and regulatory acumen to accelerate the emergence of the next generation of robust, adaptable, and trustworthy generative AI.

This enhanced section streamlines the presentation of technical, methodological, and societal challenges, introduces Table 10 for crisp visualization of technical barriers and ongoing mitigation strategies, and employs itemized bullet points for multifaceted societal risks and research opportunities, ensuring academic rigor, clarity, and optimized readability in full compliance with LaTeX academic formatting.

## 8 Security, Adversarial Threats, and Alignment

### 8.1 Threat Detection and Robustness

As generative models have become embedded in critical domains, concerns regarding their vulnerability to adversarial threats and the reliability of their outputs have correspondingly intensified. Traditional adversarial testing and red teaming approaches have uncovered foundational vulnerabilities; however, recent research has illuminated both the expanding variety and increased severity of threats facing large-scale systems—especially multimodal large language models (LLMs) and vision-language architectures. Advances in the field have shifted the focus from isolated adversarial attacks toward comprehensive taxonomies that systematically classify attack vectors by both technical sophistication and modality,

extending beyond textual to include visual and multimodal perturbations. Automated red teaming frameworks, such as those built on the searcher paradigm, now provide systematic methodologies for evaluating system-level security. These frameworks facilitate the discovery and categorization of previously underexplored weaknesses in generative AI [24, 39].

The threat landscape is evolving, manifesting escalated attack complexity. Multimodal attacks, which leverage the interplay among language, vision, and audio modalities, expose failure modes that are invisible to unimodal defenses, thereby highlighting the limitations of traditional siloed robustness evaluation. Generative agents, particularly those built upon LLM cores, are notably susceptible to attacks that manipulate chained reasoning or exploit inter-model interactions—a vulnerability exacerbated by the opacity and substantial scale of state-of-the-art architectures [39]. Efforts to bolster robustness frequently employ a defense-in-depth approach, yet these efforts must confront persistent challenges such as overfitting to known attack patterns and unintended negative outcomes from overly aggressive content filtering. The latter can lead to inadvertent blocking of benign queries, thereby degrading user experience and eroding trust [24].

The robustness of detection systems—especially within social media contexts—has also received significant scrutiny under adversarial conditions. Cutting-edge frameworks that synthesize adaptive data augmentation with adversarial training have yielded tangible improvements. By dynamically perturbing non-essential features and incorporating hard negative samples in contrastive learning paradigms, these systems reduce overfitting to benign patterns and sustain robust performance even when faced with adversarially manipulated data [24]. Nevertheless, the need persists for the development of generalizable robustness mechanisms capable of preempting the ingenuity and unpredictability characteristic of emerging attack strategies, particularly as generative models advance in multimodal integration and contextual awareness.

The examples summarized in Table 11 illustrate the diversity of adversarial threats and corresponding defensive approaches, emphasizing the multifaceted requirements of robust generative model deployment.

## 8.2 Alignment in Generative Models

While security is foundational, aligning generative models with human values and preferences represents an equally critical aspect of responsible AI development. Reinforcement learning from human feedback (RLHF) and its derivatives have emerged as dominant strategies for behavioral alignment, enabling generative systems to internalize explicit task instructions as well as complex preferences regarding style, safety, and practical utility [21, 42]. The landscape of preference tuning has expanded to encompass multimodal models—including those spanning vision, speech, and combinations thereof—necessitating adaptable alignment frameworks that integrate heterogeneous human feedback signals.

Recent comparative studies shed light on both the advantages and inherent challenges of RLHF and related alignment strategies. Large, high-quality datasets used in reward modeling can improve

how accurately models capture nuanced human preferences. However, reliance on such datasets introduces risks: they may propagate societal biases and underrepresent minority perspectives, especially as alignment protocols are scaled across tasks and user populations [42]. Challenges further intensify in complex settings such as conditional image-text generation, where preference alignment is vulnerable to problems like mode collapse or excessive conservatism—particularly under strong reward model regularization, which may stifle generative diversity and informativeness [21]. The trade-off between generating harmless (i.e., non-biased, non-harmful) and helpful (i.e., relevant, informative) outputs is still largely unresolved both technically and societally [21].

A persistent obstacle is the robust evaluation of alignment success. Existing quantitative metrics provide only a partial picture, often failing to detect low-frequency but high-consequence misalignments with significant social implications [21]. Standardizing benchmarks and constructing scalable annotation pipelines for human feedback collection remain substantial challenges, particularly for multimodal generative systems. Despite these issues, the field is rapidly advancing through innovations such as more expressive preference models, hierarchical feedback strategies, and cross-modal reward optimization, collectively pointing toward increasingly robust and socially attuned generative AI [21, 42].

Key ongoing challenges include: scalability, as large-scale human feedback risks both bias propagation and practical bottlenecks; evaluation, with prevalent metrics frequently missing critical, rare misalignments; generalization, especially across modalities and open-ended scenarios; nuanced trade-offs between harmlessness and helpfulness; and ongoing innovation in modeling human preferences, acquiring feedback, and enabling cross-modal alignment.

The intersection of adversarial robustness and human alignment in generative models delineates an urgent and dynamic research frontier. Advances in this domain are shaped by intricate trade-offs among security, safety, and practical value, as detailed in comprehensive surveys and evolving taxonomies of red teaming and preference alignment [21, 24, 39, 42].

## 9 Synthesis, Comparative Analysis, and Recommendations

### 9.1 Comparative Perspective

The trajectory of generative and augmentation strategies in artificial intelligence (AI) shows a notable progression from early, hand-crafted and deterministic pipelines toward versatile deep learning, generative, and diffusion-based frameworks [21–23, 34, 67, 71, 72, 74, 75, 86, 88]. Classical data augmentation techniques, such as geometric transformations and value-based manipulations, have been lauded for their simplicity, interpretability, and broad applicability, especially in computer vision and natural language processing domains. These foundational methods, as detailed in recent surveys [23, 67, 75, 86], primarily targeted regularization and limited-data regimes but were often bound by domain specificity and frequently fell short in semantic diversity—posing obstacles in challenges like medical imaging, where fine-grained distinctions are crucial [23, 71, 86].

**Table 11: Overview of primary adversarial threat categories, illustrative application domains, and principal classes of defense strategies.**

Adversarial Threat Category	Typical Application Domain	Defense Strategy
Textual Prompt Injection	LLM-based conversational agents	Prompt filtering, adversarial training, input sanitization
Visual Perturbation Attacks	Vision-language models, image generators	Image preprocessing, adversarial example detection, robust feature extraction
Multimodal Chained Attacks	Multimodal reasoning agents	Cross-modal consistency checks, hierarchical defense-in-depth
Model-to-model Interaction Exploits	Autonomous agent ecosystems	Traceable chained reasoning, interaction protocol hardening

Deep generative models, including Generative Adversarial Networks (GANs), transformer-based architectures, and diffusion models, have redefined the frontier of data synthesis and augmentation. They enable not only high-fidelity, controllable sample generation, but also novel multi-instance techniques—such as Mixup in image and text domains [23, 75, 86]—and massive synthetic dataset creation via GANs or diffusion frameworks [22, 23, 34, 74, 86, 88]. Notably, diffusion models are increasingly prominent for their ability to produce both images and augmentations that preserve statistical and semantic integrity, often surpassing traditional GAN-based approaches in terms of diversity and realism [34, 71, 74, 88]. For example, recent advances in fine-grained text-to-image GANs and diffusion architectures demonstrate that auxiliary objectives, such as contrastive learning and class-aware regularization, can drive state-of-the-art performance and resource efficiency on benchmark tasks [34, 71]. Moreover, in domains like medical imaging or X-ray analysis, simulation-based synthetic data, coupled with robust domain randomization, can match or even exceed the effectiveness of real data for model generalization [22].

Despite these advances, generative augmentation introduces novel complications. Risks such as semantic drift or dataset contamination through the unchecked proliferation of synthetic data (a phenomenon referred to as “AI autophagy”) have drawn increasing scrutiny [88]. Unregulated mixing of synthetic and real data can undermine model reliability, exacerbate distribution shifts, and threaten the integrity of scientific benchmarks [34, 72, 88]. Comparative studies indicate that the benefits of mitigating overfitting and enriching rare class representation depend critically on the quality control and curation of generated data; excessive reliance on model-generated samples, absent rigorous evaluation, can degrade downstream task performance [34, 72]. To address these challenges, recent surveys and empirical works underscore the significance of robust, standardized evaluation frameworks, systematic benchmarking, and the development of modality-specific augmentation guidelines. The careful balancing of semantic fidelity, sample diversity, and class distribution remains an open pursuit for maximizing the impact and rigor of generative AI augmentation [23, 34, 71, 72, 74, 75, 86, 88].

## 9.2 Criteria for Responsible Deployment

Responsible deployment of generative and augmentation-based methodologies is a central theme of this survey, underpinning both technical innovation and broader societal impact. This section aims to delineate clear, modality-aware criteria that safeguard data integrity, foster interpretability, and promote robust, context-sensitive

augmentation. These pillars are essential for aligning technological advancements with ethical, regulatory, and application-specific requirements [23, 34, 67, 71, 75, 86].

Three foundational criteria guide responsible deployment:

**Data Quality and Representational Completeness:** The integrity and diversity of the original data are crucial; augmentation cannot substitute for missing classes or modalities [34, 67, 86]. Best practices recommend principled data sampling, rigorous pre-augmentation audits, and ongoing validation to ensure adequate representation. Concretely, consider medical imaging: even advanced methods like FG-RAT GAN require labeled examples from all desired classes and subclasses to synthesize high-fidelity, class-consistent outputs; lacking such foundational diversity would undermine both empirical performance and ethical standards [34, 71].

**Interpretability and Auditability:** Interpretability underpins regulatory compliance and safe scientific or clinical deployment. Basic geometric augmentation is often inherently interpretable, but more complex deep generative methods (e.g., GANs, diffusion models, or transformers) may obscure provenance and invite unanticipated artifacts, drift, or spurious correlations. Routine interpretability assessment, provenance tracking, and human-in-the-loop oversight remain essential. For example, in semantic image synthesis with latent diffusion, manual review of outputs alongside automated provenance trails is standard practice to mitigate the risks identified in the latest synthesis models [34, 71].

**Robust and Context-Aware Augmentation:** Balancing data diversity with the avoidance of distributional shift or bias is critical [22, 67, 71]. Recent literature converges on hybrid approaches, combining deterministic, transparent transformations with data-driven generative methods, often leveraged by adaptive policy discovery and real-time monitoring [34, 86]. In medical AI, SyntheX pipelines demonstrate how combining domain randomization and calibration against real-world data enables ethically sound, empirically robust research without the logistical burden of large-scale human data collection [22, 34, 72, 74].

These criteria collectively form an essential framework for the responsible and effective deployment of generative augmentation techniques. By grounding general risks in concrete scenarios drawn from recent case studies—such as subclass fidelity in fine-grained GAN synthesis [34] and synthetic-to-real transfer in medical imaging [22]—this section aims to provide actionable guidelines applicable across modalities and domains.

In summary, a principled approach to generative augmentation, centered on data quality, interpretability, and context-aware practice, supports both the intended scientific utility and the broader objectives of ethical, reliable AI—core goals that the present survey seeks to promote and clarify.

### 9.3 Integration of Adaptive and Responsible AI

The intersection of adaptive, context-aware intelligence and responsible oversight is increasingly recognized as critical to the future development of generative augmentation systems [48, 69, 73, 76, 82]. Adaptive AI systems are expected to dynamically tailor synthesis and augmentation strategies to the evolving characteristics of real-world data—striving to balance objectives such as standardization, personalization, and fairness in real time [69, 76]. Mechanisms including meta-learning, human feedback-in-the-loop, and task-aware policy search enable augmentation pipelines to internalize context-specific priors and to modulate the generation of synthetic data accordingly [69, 76, 82].

However, the efficacy of these advances is contingent upon equally rigorous oversight protocols. Frameworks for responsible AI in augmentation increasingly mandate algorithmic transparency, persistent post-deployment monitoring, and the presence of safeguards to prevent undesirable feedback cycles—such as synthetic data inadvertently amplifying existing biases [34, 48, 73]. In fields such as healthcare, domain-specific instrumentation now embraces simulation-based validation, out-of-distribution detection, and semantic fidelity assessment as standard components of adaptive augmentation workflows [34, 74]. These priorities are reflected in emerging best practices, emphasizing the seamless integration of context-sensitive adaptation and comprehensive, multi-layered responsibility checks [34, 48, 73].

Despite recent progress, key challenges and controversies remain at the heart of integrating adaptation and responsibility in generative augmentation. There is vigorous debate over the effectiveness and sufficiency of current fairness and bias mitigation strategies: some works prioritize maximizing synthetic data diversity through explicit loss formulations to ensure broad coverage of minority segments and semantic modes [69], while others emphasize stability and semantic fidelity, advocating for constraint-driven or supervised adaptation pipelines to prevent distributional drift [48, 76]. Furthermore, disagreements persist regarding the tradeoff between transparency and performance, with some frameworks sacrificing transparency for rapid adaptation, while others embed post-hoc monitoring and explainability, sometimes at the cost of system efficiency [34, 73].

The development of integrated frameworks, such as those leveraging classifier guidance, auxiliary supervision, or contrastive learning strategies, demonstrates ongoing efforts to reconcile these tensions by improving both performance and accountability in tasks like text-to-image and medical image synthesis [34, 48, 73, 76]. However, the field continues to seek consensus on best practices balancing pluralism, robustness, and ethical constraints, with open questions around generalization, domain transfer, and context-specific risk assessment.

### 9.4 State-of-the-Art (SOTA) Synthesis Models

The recent proliferation of advanced synthesis models encapsulates the ongoing evolution and intricate trade-offs inherent in leading generative paradigms. Table 12 presents a structured comparison of prominent state-of-the-art synthesis architectures and their principal contributions.

Contemporary models increasingly leverage the complementary strengths of deterministic and probabilistic approaches—combining interpretability and stability from explicit constraints (e.g., contrastive, feature-matching, or categorical losses) with the flexibility of multi-stage, generative learning frameworks [34, 71, 72, 74]. For example, FG-RAT GAN [34] integrates an auxiliary classifier in the discriminator to guide semantic class consistency and employs a contrastive objective that enhances intra-class alignment and inter-class separation, which is crucial for fine-grained text-to-image generation. VQ-Diffusion leverages quantization and controlled diffusion to ensure high semantic content and diversity, while latent diffusion models [34, 71, 88] enable more scalable and conditional generation in different domains, such as segmentation mask-guided synthesis.

Such synergies facilitate the creation of perceptually convincing, diverse, and controllable samples, empowering workflows across vision, language, and scientific applications. At the same time, emerging methods like image-to-image diffusion with refinement and color transfer modules [71] or simulation-based synthetic data frameworks [22] demonstrate the potential of augmenting or even replacing real datasets in constrained domains. Ongoing challenges include the need for comprehensive, multi-domain benchmarking and standardized evaluation, as well as ensuring the effectiveness, fairness, and traceability of generated data [22, 23, 34, 71, 74, 86]. Responsible and systematic benchmarking will be indispensable for steering the adoption of generative and augmentation technologies, ensuring their integrity and utility in real-world applications.

## 10 Conclusion

In this survey, we set out to provide a comprehensive and critical overview of recent advances in fairness, bias mitigation, and model reliability within modern AI systems. Our primary objectives were to: (1) systematically synthesize the state-of-the-art research on the technical underpinnings and limitations of fairness-aware machine learning, (2) elucidate existing ethical frameworks and the complexities they introduce for evaluating bias and trustworthiness, and (3) highlight open technical and ethical challenges for future work. We further distinguished this survey by introducing a new taxonomy for categorizing fairness interventions, and by mapping the interdependencies between ethical frameworks, model architectures, and deployment pipelines.

A core controversy in this area revolves around the tension between statistical definitions of fairness—such as demographic parity and equalized odds—and individual-level guarantees, as well as the question of whether algorithmic bias can ever be mitigated in isolation from broader societal and data-driven contexts. Some researchers argue for prioritizing context-sensitive, domain-specific metrics, while critics caution that these may fragment standardization or undermine generalizability. These opposing viewpoints highlight an ongoing need for holistic evaluation methodologies.

To guide researchers, we summarize below the most pressing open challenges and future research directions, as identified in the surveyed literature.

Despite significant advances, realizing genuinely fair, unbiased, and reliable AI remains an open and evolving challenge. Continued collaboration across technical, ethical, and policy domains will be

Table 12: Comparison of Select State-of-the-Art Generative Synthesis Models

Model	Key Features	Domain Applications	Notable Strengths/Trade-offs
StyleGAN [74]	Arbitrary classifier in discriminator, contrastive learning through cross-batch memory (CBM), fine-grained text-to-image synthesis	Visual, cross-modal generation	Achieves state-of-the-art FID and competitive IS on fine-grained datasets, surpassing LADTTE and VQ-Diffusion with much fewer parameters; highly efficient but depends on labeled subspaces [34]
VQ-Diffusion [74, 88]	Vector quantization combined with diffusion process, controllable synthesis pipeline	High-resolution image generation, data augmentation	Excels in semantic fidelity and sample diversity, but is computationally intensive; does not match PG-RAT GAN in label-efficient settings
LDM (Latent Diffusion Models) [34, 88]	Synthesis in latent space, scalable architecture, supports conditional translation	Image-to-image translation, semantic segmentation, style transfer	Delivers realistic and semantically aligned outputs; scalability comes at the cost of requiring large, high-quality datasets for effective training
Transformer-based Hybrid Models [22, 71, 75]	Captures global structure via transformers, can be integrated with convolutional layers and cross-modal inputs	Plenoptic image completion, robust multi-domain data, synthesis	Improves interpretability and expressiveness, demonstrates adaptability to complex or multi-modal synthesis scenarios; potential challenges in model complexity and training

Table 13: Open Challenges and Future Research Directions in Fairness, Bias, and Reliability

Challenge Category	Open Challenge	Current Limitation	Promising Research Directions
Technical	Formalizing Fairness	Lack of unified definitions across domains	Developing adaptable, context-aware fairness metrics
Ethical	Societal Context	Difficulty modeling social and historical biases	Integrating sociotechnical analysis within algorithm design
Reliability	Model Robustness	Vulnerability to adversarial attacks and distribution shifts	Robust optimization and certification methods
Evaluation	Standardized Benchmarks	Fragmented evaluation datasets and protocols	Establishing community benchmarks and shared resources
Deployment	Interpretability	Lack of transparent reasoning for complex models	Advancing explainable AI and post-hoc interpretability tools

critical for progress. We hope this work serves as a resource and catalyst for researchers, practitioners, and policymakers seeking to advance the state of responsible AI.

10.1 Transformative Developments

The past decade has witnessed a profound transformation in generative modeling, marked by the emergence of generative adversarial networks (GANs), diffusion models, hybrid and transformer-based architectures, context-aware generative frameworks, and adaptive data augmentation methods. These advances have irreversibly reshaped the landscape of computer vision, medical artificial intelligence, drug discovery, and related fields, establishing generative models not only as tools for data synthesis but as indispensable engines for representation learning, privacy preservation, and scientific progress [21, 22, 32, 34, 48, 53, 56, 69, 71–73, 76, 77, 79, 82, 88, 91].

Since their introduction, GANs have catalyzed paradigm shifts in imaging, text-to-image synthesis, simulation-driven research, and even regulatory strategies in sensitive domains such as healthcare and the social sciences. Their strength in modeling high-dimensional distributions has enabled the creation of realistic synthetic datasets, addressing persistent issues of data scarcity, imbalance, and privacy. Applications range from medical diagnostics to the construction of synthetic patient populations and “digital twins” for regulatory-compliant data sharing [56, 72, 73, 76, 77, 82]. Despite their promise, GANs face notable limitations, including mode collapse, training instability, and vulnerability to biases in the training corpus. In response, the field has introduced architectural innovations—such as conditional, Wasserstein, and hybrid GANs—and has elevated the integration of fairness-aware and privacy-preserving algorithms [21, 53, 73, 77].

Diffusion models, representing a subsequent wave of innovation, have exerted significant influence, particularly in computer vision and scientific imaging. Employing noise-perturbation and denoising dynamics, these models generate diverse and high-quality samples in both unconditional and conditional settings. They frequently surpass GANs in sample fidelity and stability, though at a cost of increased computational requirements [32, 34, 79, 88]. Recent developments such as latent diffusion models have contributed to a significant reduction in computational demands without compromising

output quality, as they operate within compressed, perceptually-rich latent spaces and allow efficient conditioning on various modalities and high-resolution tasks [71, 91]. For example, latent diffusion approaches have been demonstrated to deliver competitive or state-of-the-art performance with far fewer parameters than conventional pixel-wise models, streamlining tasks including inpainting and super-resolution [91]. Furthermore, diffusion models have been successfully applied to semantically complex domains such as conditional 3D molecule generation and medical imaging, introducing equivariant architectures and measurement-guided generation for heightened label reliability and clinical utility [56, 71, 76].

Building on these strengths, hybrid approaches—notably latent diffusion models and architectures incorporating transformer modules—have scaled generative modeling to complex, multi-modal, and high-resolution regimes [71, 79]. Transformers, with their aptitude for modeling long-range dependencies, have unified generative pipelines across text, vision, and cross-modal tasks. This capacity enables open-vocabulary detection, segmentation, and self-supervised learning, further blurring the lines between symbolic and subsymbolic artificial intelligence and facilitating context-aware synthesis and control [19, 22, 48, 69, 71, 91]. Of particular note are architectures that merge the strengths of transformers (for global relationship modeling) and convolutional networks (for local detail refinement), leading to significant improvements in pluralistic image completion and semantic fidelity for diverse and high-resolution outputs [71, 79, 82].

Context mechanisms—via side-information, domain adaptation, or weak supervision—demonstrate how generative models can bridge unsupervised synthesis and (semi-)supervised learning, thereby enhancing both data fidelity and downstream task utility [69, 73, 82]. For instance, explicit diversity objectives and semantic guidance in conditional synthesis have enabled user-driven manipulation and improved representation of plausible outputs [69]. Concurrently, adaptive data augmentation has evolved from manual techniques to meta-learned, contextually optimized curricula. These methods promote generalization, fairness, and robustness, even in low-resource or distributionally-shifted settings [34, 48, 53, 77, 91]. Collectively, these innovations yield not merely richer synthetic datasets but dynamic frameworks, wherein synthetic data generation, augmentation, and feedback from subsequent tasks inform one another iteratively.

## 10.2 Principles for Future AI

A critical evaluation of generative modeling's current ecosystem surfaces several guiding principles shaping the discipline's trajectory. Foremost among these is generalization: models must transcend idiosyncratic dataset artifacts to demonstrate robustness across divergent tasks, populations, and environments. Fairness and alignment require ongoing diligence, addressing historical biases, demographic underrepresentation, and minimizing social or regulatory harms at the data, model, and system levels [48, 53, 77, 88, 91]. Scalability remains essential, especially as generative AI expands from narrowly defined applications to open-world or multi-modal environments; here, hybrid architectures and foundation models must uphold both efficiency and adaptability [21, 56, 71, 79, 82]. Above all, ethical design—grounded in transparency, accountability, and human-centered values—must underlie every generative modeling system, safeguarding against unintended consequences and sustaining public trust.

Generalization is essential for ensuring that AI models exhibit robustness across diverse tasks, datasets, and operational conditions. Models achieving high fidelity and diversity, as demonstrated in image synthesis tasks [48, 71, 79, 91], illustrate the importance of architectures and loss functions that improve visual realism and maintain semantic consistency.

Fairness and alignment remain foundational for trustworthy AI. This involves continuous effort to mitigate historical biases and ethical risks, integrating mechanisms both at the data curation stage and within model design [53, 77]. Notably, quality assessment methods [77] that better align with human subjective judgment are critical for identifying and addressing potential artifacts or sources of discrimination.

Scalability supports efficient transfer of generative models to new domains while maintaining flexibility. The evolution from convolutional approaches to transformer-based and hybrid models demonstrates quantitative gains in fidelity, diversity, and adaptability [21, 71, 82, 91]. Efficient latent representation [71, 91] and architectural innovations are central to ensuring models remain adaptable as application contexts broaden.

Ethical design must be embedded at every level of generative modeling. Transparency, explainability, and accountability—in both the system's underlying mechanisms and its outputs—are necessary to guard against unintended consequences and maintain public trust.

## 10.3 Bridging Technical and Responsible Innovation

Technical advances in generative modeling demand parallel progress in interpretability, security, policy compliance, and responsible deployment. Among these, interpretability is especially pressing. As models grow in complexity and opacity, mechanisms for human-understandable explanations, provenance tracking, watermarking, and information disclosure become critical, especially within high-stakes environments such as healthcare, law, and critical infrastructure [27, 74, 85]. Security and privacy also require multifaceted solutions, comprising both technical guarantees (differential privacy,  $k$ -anonymity, adversarial robustness) and best operational practices aimed at preventing reidentification, memorization, and

membership inference attacks [27, 76, 77, 85]. The regulatory landscape is rapidly adapting to generative AI's expanding influence, with ongoing efforts to establish standardized evaluation protocols, maintain auditable usage records, and introduce oversight at the dataset and model levels [72, 74, 76, 85, 88].

Responsible deployment is therefore best understood as a sociotechnical enterprise, one balancing utility against risk to ensure generative models are not only accurate but also equitable, explainable, and responsive to shifting legal and societal imperatives. This necessitates interdisciplinary research, continual stakeholder input, and sustained investment at the intersection of machine learning, ethics, and governance [27, 72, 74, 85].

Integration of interpretability tools and protocols is vital for ensuring transparent, traceable outputs throughout the model development lifecycle. Robust privacy-guarding mechanisms must be implemented to address risks such as bias, synthetic data contamination, and membership inference. The adoption of standardized evaluation and audit processes supports ongoing accountability, particularly as the absence of domain-tailored protocols hinders deployment in sensitive fields like healthcare [76, 85]. Finally, effective progress depends on sustained and open dialogue among technical, ethical, and policy stakeholders, and the establishment of responsible usage standards and collaborative practices to guide the evolution of generative AI [27, 74, 85].

## 10.4 Outlook

Looking ahead, the sustainable and responsible evolution of synthetic data research will depend on interdisciplinarity, transparency, and community stewardship. Cross-disciplinary collaboration—engaging computer science, statistics, social science, medicine, and law—is required to address persistent challenges in evaluation, privacy, and bias. Transparent practices, open-source tools, and common benchmarks will be critical for reproducibility and trustworthy comparative evaluation [21, 56, 71, 79]. Engagement from a diverse, global community—comprising both technical innovators and affected stakeholders—will be necessary to align generative AI development with societal interests and mitigate the risk of amplifying inequities or eroding trust.

In conclusion, the convergence of GANs, diffusion models, transformer and hybrid architectures, context-driven generative frameworks, and adaptive augmentation signals a new epoch in artificial intelligence. These technologies are simultaneously fueling innovation and surfacing critical questions regarding fairness, accountability, and the overall value to society. The path forward rests on a deliberate synthesis of technical advancement and responsible innovation—anchored by foundational principles and a steadfast commitment to the public good [21, 22, 32, 34, 48, 53, 56, 69, 71–73, 76, 77, 79, 82, 88, 91].

## References

- [1] D. Adam. 2024. Synthetic data can aid the analysis of clinical outcomes: How much can it be trusted? *Proceedings of the National Academy of Sciences* 121, 32 (2024), e2414310121. doi:10.1073/pnas.2414310121
- [2] I. Ahmed, T. Xie, A. K. Bashir, and A. D. Jurcut. 2024. Computer Vision Based Transfer Learning-Aided Transformer Model for Plant Disease Recognition. *IEEE Access* 12 (2024), 28798–28809. doi:10.1109/ACCESS.2024.3363345
- [3] M. U. Akbar, W. Wang, and A. Eklund. 2025. Beware of diffusion models for synthesizing medical images—a comparison with GANs in terms of memorizing

- brain MRI and chest x-ray images. *Machine Learning: Science and Technology* 6, 1 (2025), 015022. doi:10.1088/2632-2153/ad9a3a
- [4] Duanhua Cao, Geng Chen, Jiaxin Jiang, Jie Yu, Runze Zhang, Mingan Chen, Wei Zhang, Lifan Chen, Feisheng Zhong, Yingying Zhang, Chenghao Lu, Xutong Li, Xiaomin Luo, Sulin Zhang, and Mingyue Zheng. 2024. Generic protein-ligand interaction scoring by integrating physical prior knowledge and data augmentation modelling. *Nature Machine Intelligence* 6, 6 (2024), 688–700. doi:10.1038/s42256-024-00849-z
  - [5] N. Carion, F. Massa, G. Synnaeve, N. Usunier, A. Kirillov, and S. Zagoruyko. 2020. End-to-End Object Detection with Transformers. *arXiv preprint arXiv:2005.12872* (2020). <https://arxiv.org/abs/2005.12872>
  - [6] C. Chadebec, E. Thibaut-Sutre, N. Burgos, and S. Allasoinniere. 2023. Data Augmentation in High Dimensional Low Sample Size Setting Using a Geometry-Based Variational Autoencoder. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 45, 3 (March 2023), 2879–2896. doi:10.1109/TPAMI.2023.3185773 <https://ieeexplore.ieee.org/document/9806307>.
  - [7] T. Chakraborty, U. Reddy K S, S. M. Naik, M. Panja, and B. Manvitha. 2024. Ten years of generative adversarial nets (GANs): a survey of the state-of-the-art. *Machine Learning: Science and Technology* 5, 1 (2024), 011001. doi:10.1088/2632-2153/ad1f77
  - [8] J. Chen, H. Chen, K. Chen, Y. Zhang, Z. Zou, and Z. Shi. 2025. Diffusion Models for Imperceptible and Transferable Adversarial Attack. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 47, 2 (2025), 961–977. doi:10.1109/TPAMI.2024.3480519
  - [9] M. Chen, S. Mei, J. Fan, and M. Wang. 2024. An Overview of Diffusion Models: Applications, Guided Generation, Statistical Rates and Optimization. *arXiv preprint arXiv:2404.07771* (2024). <https://arxiv.org/abs/2404.07771>
  - [10] P. Chen, X. Yu, X. Han, K. Wang, G. Li, L. Xie, Z. Han, and J. Jiao. 2025. P2Object: Single Point Supervised Object Detection and Instance Segmentation. *International Journal of Computer Vision* (2025). doi:10.1007/s11263-025-02441-3
  - [11] S. Chen, E. Dobriban, and J. H. Lee. 2020. A Group-Theoretic Framework for Data Augmentation. *Journal of Machine Learning Research* 21, 245 (2020), 1–71. <https://www.jmlr.org/papers/volume21/20-163/20-163.pdf>
  - [12] Y. Chen, X. Yuan, R. Wu, J. Wang, Q. Wang, L. Zhang, and M.-M. Cheng. 2025. YOLO-MS: Rethinking Multi-Scale Representation Learning for Real-Time Object Detection. *IEEE Transactions on Pattern Analysis and Machine Intelligence* (2025). doi:10.1109/TPAMI.2025.10872821
  - [13] Ziqi Chen, Bo Peng, Xia Ning, Mengxuan Zhou, Jiaqi Zhu, Pengyong Li, Rongrong Jin, and Long Zhang. 2025. Generating 3D small binding molecules using shape-conditioned diffusion models with guidance. *Nature Machine Intelligence* 7 (2025), 588–591. <https://www.nature.com/articles/s42256-025-01009-4>
  - [14] V. Claveau, A. Chaffin, and E. Kijak. 2021. Generating artificial texts as substitution or complement of training data. *Artificial Intelligence* 298 (2021), 103528. <https://arxiv.org/pdf/2110.13016>
  - [15] E. De Cristofaro. 2024. Synthetic Data: Methods, Use Cases, and Risks. *arXiv preprint arXiv:2303.01230* (2024). <https://arxiv.org/abs/2303.01230>
  - [16] F. A. Croitoru, V. Hondru, R. T. Ionescu, and M. Shah. 2023. Diffusion Models in Vision: A Survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 45, 9 (2023), 10850–10869. doi:10.1109/TPAMI.2023.3261988
  - [17] M. Dedeoglu, S. Lin, Z. Zhang, and J. Zhang. 2024. Continual Learning of Generative Models With Limited Data: From Wasserstein-1 Barycenter to Adaptive Coalescence. *IEEE Transactions on Neural Networks and Learning Systems* 35, 9 (2024), 12042–12055. doi:10.1109/TNNLS.2024.10070745
  - [18] Yashar Deldjoo, Zhi He, Julian McAuley, Andrey Korikov, Scott Sanner, Arnau Ramisa, Ramon Vidal, Mahinthan Sathiamoorthy, Amirhossein Kasrizadeh, Simone Milano, and Francesco Ricci. 2024. Recommendation with Generative Models. *arXiv preprint arXiv:2409.15173* (sep 2024). <https://arxiv.org/abs/2409.15173>
  - [19] A. P. Dempster, N. M. Laird, and D. B. Rubin. 2023. Artificial Data in Bootstrapping: Perspectives on the Origins and Implications. *Artificial Intelligence* 323 (2023), 103194. <https://www.sciencedirect.com/science/article/pii/S004370222001456>.
  - [20] J. Ding, N. Xue, Y. Long, G.-S. Xia, Q. Lu, A. Bai, W. Yang, and X. Yao. 2022. Object Detection in Aerial Images: A Large-Scale Benchmark and Challenges. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 44, 11 (2022), 7778–7796. doi:10.1109/TPAMI.2021.3117988
  - [21] C. Du, J. Zhu, and B. Zhang. 2017. Learning Deep Generative Models With Doubly Stochastic MCMC. *IEEE Transactions on Neural Networks and Learning Systems* 28, 6 (2017), 1371–1383. doi:10.1109/TNNLS.2017.2688499
  - [22] X. Du, N. Kolkin, G. Shakhnarovich, and A. Bhattad. 2024. Generative Models: What Do They Know? Do They Know Things? Let's Find Out! *arXiv preprint arXiv:2311.17137* (Oct. 2024). <https://arxiv.org/abs/2311.17137>
  - [23] L. L. Duan, J. E. Johndrow, and D. B. Dunson. 2018. Scaling up Data Augmentation MCMC via Calibration. *Journal of Machine Learning Research* 19, 64 (2018), 1–34. <https://jmlr.org/papers/volume19/17-573/17-573.pdf>
  - [24] Chukwuebuka Joseph Ejayi, Dongsheng Cai, Francis Ofoma Eze, Makuchukwu Bennedith Ejayi, Jennifer Ene Idoko, Sarpong Kwadwo Asere, and Thomas Ugochukwu Ejayi. 2025. Polynomial-SHAP as a SMOTE alternative in conglomerate neural networks for realistic data augmentation in cardiovascular and breast cancer diagnosis. *Journal of Big Data* 12 (2025), Article 97. doi:10.1186/s40537-025-01152-3
  - [25] L. Fan, Y. Yang, F. Wang, N. Wang, Z. Zhang, Z. Cao, and D. Lin. 2023. Super Sparse 3D Object Detection. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 45, 10 (2023), 12490–12505. doi:10.1109/TPAMI.2023.3286409
  - [26] F. Fürutter, G. Muñoz-Gil, and H. J. Briegel. 2024. Quantum circuit synthesis with diffusion models. *Nature Machine Intelligence* 6, 5 (2024), 515–524. doi:10.1038/s42256-024-00831-9
  - [27] Cong Gao, Benjamin D. Killeen, Yicheng Hu, Robert B. Grupp, Russell H. Taylor, Mehran Armand, and Mathias Unberath. 2023. Synthetic data accelerates the development of generalizable learning-based algorithms for X-ray image analysis. *Nature Machine Intelligence* 5 (2023), 294–308. doi:10.1038/s42256-023-00629-1
  - [28] M. Goyal and Q. H. Mahmoud. 2024. A Systematic Review of Synthetic Data Generation Techniques Using Generative AI. *Electronics* 13, 17 (2024), 3509. <https://www.mdpi.com/2079-9292/13/17/3509>
  - [29] Kaiming He, Georgia Gkioxari, Piotr Dollár, and Ross Girshick. 2017. Mask R-CNN. *arXiv preprint arXiv:1703.06870* (2017). <https://arxiv.org/abs/1703.06870>
  - [30] Philipp Hess, Michael Aich, Baoxiang Pan, and Niklas Boers. 2025. Fast, scale-adaptive and uncertainty-aware downscaling of Earth system model fields with generative machine learning. *Nature Machine Intelligence* 7 (2025), 363–373. doi:10.1038/s42256-025-00980-5
  - [31] C. Hou, J. Zhang, and T. Zhou. 2023. When to Learn What: Model-Adaptive Data Augmentation Curriculum. *arXiv preprint arXiv:2309.04747* (sep 2023). <https://arxiv.org/abs/2309.04747>
  - [32] M. Ibrahim, Y. Al Khalil, S. Amirrajab, C. Sun, M. Breeuwer, J. Pluim, B. Elen, G. Ertaylan, and M. Dumontier. 2024. Generative AI for Synthetic Data Across Multiple Medical Modalities: A Systematic Review of Recent Developments and Challenges. *arXiv preprint arXiv:2407.00116* (2024). <https://arxiv.org/abs/2407.00116>
  - [33] I. Igashov, H. Stärk, C. Vignac, A. Schneuing, V. Garcia Satorras, P. Frossard, M. Welling, M. Bronstein, and B. Correia. 2024. Equivariant 3D-conditional diffusion model for molecular linker design. *Nature Machine Intelligence* 6, 4 (2024), 417–427. doi:10.1038/s42256-024-00815-9
  - [34] Junjun Jiang, Yi Yu, Zheng Wang, Xianming Liu, and Jiayi Ma. 2019. Graph-Regularized Locality-Constrained Joint Dictionary and Residual Learning for Face Sketch Synthesis. *IEEE Transactions on Image Processing* 28, 2 (February 2019), 628–641. doi:10.1109/TIP.2018.2870936 <https://doi.org/10.1109/TIP.2018.2870936>
  - [35] X. Jiang and Y. Wu. 2023. Remote Sensing Object Detection Based on Convolution and Swin Transformer. *IEEE Access* 11 (2023), 38643–38656. doi:10.1109/ACCESS.2023.3267435
  - [36] I. Joshi, M. Grimmer, J. Fierrez, A. Dantcheva, J. L. Crowley, and J. Busch. 2024. Synthetic Data in Human Analysis: A Survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 46, 7 (July 2024), 4957–4976. <https://ieeexplore.ieee.org/document/10423161>.
  - [37] A. Karev and D. Xu. 2025. ConSCompF: Consistency-focused Similarity Comparison Framework for Generative Large Language Models. *Journal of Artificial Intelligence Research* 82 (2025), 1–32. <https://www.jair.org/index.php/jair/article/view/17028>
  - [38] T. Karras, M. Aittala, T. Aila, and S. Laine. 2022. Elucidating the Design Space of Diffusion-Based Generative Models. *arXiv preprint arXiv:2206.00364* (2022). <https://arxiv.org/abs/2206.00364>
  - [39] S. Kong and D. Ramanan. 2025. OpenGAN: Open-Set Recognition via Open Data Generation. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 47, 5 (May 2025), 3233–3243. <https://ieeexplore.ieee.org/document/9799769>.
  - [40] Lisa Kühnel, Julian Schneider, Ines Perrar, Tim Adams, Sobhan Moazemi, Fabian Prasser, Ute Nöthlings, Holger Fröhlich, and Juliane Fluck. 2024. Synthetic data generation for a longitudinal cohort study – evaluation, method extension and reproduction of published data analysis results. *Scientific Reports* 14 (2024), 14412. doi:10.1038/s41598-024-62102-2
  - [41] V. Kumar, P. Singh, S. Sharma, and R. Prasad. 2024. Computer Vision-Based Framework for Data Extraction From Document Images Using Deep Learning Techniques. *IEEE Access* 13 (2024), 17706–17723. doi:10.1109/ACCESS.2024.3351125 <https://ieeexplore.ieee.org/document/10813363/>.
  - [42] B. Li, Y. Hou, and W. Che. 2022. Data Augmentation Approaches in Natural Language Processing: A Survey. *arXiv preprint arXiv:2110.01852 [cs.CL]* (June 2022). <https://arxiv.org/abs/2110.01852>
  - [43] T. Li, L. Biferale, F. Bonaccorso, M. A. Scarpolini, and M. Buzzicotti. 2024. Synthetic Lagrangian turbulence by generative diffusion models. *Nature Machine Intelligence* 6, 4 (2024), 393–403. doi:10.1038/s42256-024-00810-0
  - [44] X. Li, H. Ding, H. Yuan, W. Zhang, J. Pang, G. Cheng, K. Chen, Z. Liu, and C. C. Loy. 2024. Transformer-Based Visual Segmentation: A Survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 46, 6 (2024), 10138–10163. doi:10.1109/TPAMI.2024.10613466
  - [45] C.-H. Lin, C. Kaushik, E. L. Dyer, and V. Muthukumar. 2024. The good, the bad and the ugly sides of data augmentation: An implicit spectral regularization perspective. *Journal of Machine Learning Research* 25, 91 (2024), 1–85. <https://jmlr.org/papers/volume25/22-1312/22-1312.pdf>
  - [46] L. Lin, H. Mu, Z. Zhai, M. Wang, Y. Wang, R. Wang, J. Gao, Y. Zhang, W. Che, T. Baldwin, X. Han, and H. Li. 2025. Against The Achilles' Heel: A Survey on



- Red Teaming for Generative Models. *Journal of Artificial Intelligence Research* 82 (2025), 1–84. <https://www.jair.org/index.php/jair/article/view/17654>
- [47] L. Liu, W. Ouyang, X. Wang, P. Fieguth, J. Chen, X. Liu, and M. Pietikäinen. 2020. Deep learning for generic object detection: A survey. *Pattern Recognition* 102 (2020), 107198. <https://www.sciencedirect.com/science/article/pii/S0031320319303851>
- [48] R. Liu, J. Wei, F. Liu, C. Si, Y. Zhang, J. Rao, S. Zheng, D. Peng, D. Yang, D. Zhou, and A. M. Dai. 2024. Best Practices and Lessons Learned on Synthetic Data. *arXiv preprint arXiv:2404.07503* (Aug. 2024). <https://arxiv.org/abs/2404.07503>
- [49] Y. Liu, R. Shen, and X. Shen. 2024. Novel Uncertainty Quantification Through Perturbation-Assisted Sample Synthesis. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 46, 12 (2024), 7813–7824. <https://ieeexplore.ieee.org/document/10508110>
- [50] Ze Liu, Yutong Lin, Yue Cao, Han Hu, Yixuan Wei, Zheng Zhang, Stephen Lin, and Baining Guo. 2021. Swin Transformer: Hierarchical Vision Transformer Using Shifted Windows. *arXiv preprint arXiv:2103.14030* (2021). <https://arxiv.org/abs/2103.14030>
- [51] L. Long, R. Wang, R. Xiao, J. Zhao, X. Ding, G. Chen, and H. Wang. 2024. On LLMs-Driven Synthetic Data Generation, Curation, and Evaluation: A Survey. *arXiv preprint arXiv:2406.15126* (2024). <https://arxiv.org/abs/2406.15126>
- [52] Y. Lu, L. Chen, Y. Zhang, M. Shen, H. Wang, X. Wang, C. van Rechem, T. Fu, and W. Wei. 2023. Machine Learning for Synthetic Data Generation: A Review. *arXiv preprint arXiv:2302.04062* (2023). <https://arxiv.org/abs/2302.04062>
- [53] Y. Luo, Q. Yang, Y. Fan, H. Qi, and M. Xia. 2024. Measurement Guidance in Diffusion Models: Insight from Medical Image Synthesis. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 46, 12 (2024), 7983–7997. doi:10.1109/TPAMI.2024.3363309
- [54] S. Nie, F. Zhu, Z. You, X. Zhang, J. Ou, J. Hu, J. Zhou, Y. Lin, J.-R. Wen, and C. Li. 2025. Large Language Diffusion Models. *arXiv preprint arXiv:2502.09992* (2025). <https://arxiv.org/abs/2502.09992>
- [55] E. Papadaki, A. G. Vrahatis, and S. Kotsiantis. 2024. Exploring Innovative Approaches to Synthetic Tabular Data Generation. *Electronics* 13, 10 (2024), 1965. doi:10.3390/electronics13101965
- [56] S. Perdikis, R. Leeb, R. Chavarriaga, and J. d. R. Millán. 2021. Context-Aware Learning for Generative Models. *IEEE Transactions on Neural Networks and Learning Systems* 32, 8 (2021), 3471–3483. doi:10.1109/TNNLS.2020.3014897
- [57] Zhaozhi Qian, Thomas Callender, Baiba Cebere, Sarah M. Janes, Nishant Navani, and Mihaela van der Schaar. 2024. Synthetic data for privacy-preserving clinical risk prediction. *Scientific Reports* 14 (2024), Article number: 25676. doi:10.1038/s41598-024-72894-y
- [58] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. 2021. Learning Transferable Visual Models From Natural Language Supervision. *arXiv preprint arXiv:2103.00020* (2021). <https://arxiv.org/abs/2103.00020>
- [59] H. H. Rashidi, S. Albahra, B. P. Rubin, and B. Hu. 2024. A novel and fully automated platform for synthetic tabular data generation and validation. *Scientific Reports* 14 (2024). <https://www.nature.com/articles/s41598-024-73608-0>
- [60] D. B. Resnik, M. Hosseini, J. J. H. Kim, G. Epiphaniou, and C. Maple. 2025. GenAI synthetic data create ethical challenges for scientists. Here’s how to address them. *Proceedings of the National Academy of Sciences* 122, 9 (2025), e2409182122. doi:10.1073/pnas.2409182122
- [61] Y. Rong, T. Leemann, T. T. Nguyen, L. Fiedler, P. Qian, V. V. Unhelkar, T. Seidel, G. Kasneci, and E. Kasneci. 2024. Towards Human-Centered Explainable AI: A Survey of User Studies for Model Explanations. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 46, 4 (2024), 2104–2122. <https://dblp.org/rec/journals/pami/RongLNFQUSKK24>
- [62] S. Ruggles. 2025. The shortcomings of synthetic census microdata. *Proceedings of the National Academy of Sciences* 122, 11 (2025), e2424655122. doi:10.1073/pnas.2424655122
- [63] J. Rydzewski, J. M. A. Grimme, and M. W. F. Fischer. 2023. Manifold learning in atomistic simulations: a conceptual review. *Machine Learning: Science and Technology* 4, 3 (2023), 031001. doi:10.1088/2632-2153/ace81a
- [64] V. M. Sánchez-Cartagena, M. Esplà-Gomis, J. A. Pérez-Ortiz, and F. Sánchez-Martínez. 2024. Non-Fluent Synthetic Target-Language Data Improve Neural Machine Translation. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 46, 2 (2024), 837–850. doi:10.1109/TPAMI.2023.3333949
- [65] Hualian Sheng, Sijia Cai, Na Zhao, Bing Deng, Qiao Liang, Min-Jian Zhao, and Jieping Ye. 2024. CT3D++: Improving 3D Object Detection with Keypoint-induced Channel-wise Transformer. *CoRR abs/2406.08152* (2024). doi:10.48550/ARXIV.2406.08152
- [66] J. Shi, L. Ma, X. Wei, C. Fang, and Y. Zhang. 2024. Differentiable Image Data Augmentation and Its Applications: A Survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 46, 2 (February 2024), 1148–1164. doi:10.1109/TPAMI.2023.3330862 <https://ieeexplore.ieee.org/document/10310260>
- [67] Connor Shorten and Taghi M. Khoshgoftaar. 2019. A survey on Image Data Augmentation for Deep Learning. *Journal of Big Data* 6, 1 (2019), 60. doi:10.1186/s40537-019-0197-0
- [68] C. Shorten, T. M. Khoshgoftaar, and B. Furht. 2021. Text Data Augmentation for Deep Learning. *Journal of Big Data* 8 (2021), 101. <https://journalofbigdata.springeropen.com/articles/10.1186/s40537-021-00492-0>
- [69] Serban Stan and Mohammad Rostami. 2024. Preserving Fairness in AI under Domain Shift. *Journal of Artificial Intelligence Research* 81 (2024), 1966–2003. doi:10.1613/jair.1.16694
- [70] Y. Sun, J. Zhang, Y. Liu, and X. Li. 2024. The evolution of object detection methods. *Pattern Recognition* 148 (2024), 110438. <https://www.sciencedirect.com/science/article/pii/S003132032400136X>
- [71] Hongchen Tan, Xiuping Liu, Meng Liu, Baocai Yin, and Xin Li. 2021. KT-GAN: Knowledge-Transfer Generative Adversarial Network for Text-to-Image Synthesis. *IEEE Transactions on Image Processing* 30 (2021), 1275–1290. doi:10.1109/TIP.2021.3049544
- [72] Y. Tang, J. Weng, and P. Zhang. 2023. Neural-network solutions to stochastic reaction networks. *Nature Machine Intelligence* 5 (2023), 376–385. doi:10.1038/s42256-023-00632-6
- [73] R. Timponi and Y. Yang. 2024. Artificial Data, Real Insights: Evaluating Opportunities and Risks of Expanding the Data Ecosystem with Synthetic Data. *arXiv preprint arXiv:2408.15260* (Aug. 2024). <https://arxiv.org/abs/2408.15260>
- [74] M. C. Tsakiris. 2023. Low-Rank Matrix Completion Theory via Plücker Coordinates. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 45, 8 (2023), 10084–10099. <https://ieeexplore.ieee.org/document/10056236>
- [75] P.-D. Tudosiu, W. H. L. Pinaya, P. Ferreira Da Costa, J. Dafflon, A. Patel, P. Borges, V. Fernandez, M. S. Graham, R. J. Gray, P. Nachev, S. Ourselin, and M. J. Cardoso. 2024. Realistic morphology-preserving generative modelling of the brain. *Nature Machine Intelligence* 6, 7 (2024), 811–819. doi:10.1038/s42256-024-00864-0
- [76] Wijnand van Woerkom, Davide Grossi, Henry Prakken, and Bart Verheij. 2024. A Fortiori Case-Based Reasoning: From Theory to Data. *Journal of Artificial Intelligence Research* 81 (2024), 2113–2151. doi:10.1613/jair.1.15178
- [77] Z. Wan, J. Zhang, D. Chen, and J. Liao. 2024. High-Fidelity and Efficient Pluralistic Image Completion With Transformers. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 46, 12 (2024), 9612–9629. doi:10.1109/TPAMI.2024.3382342
- [78] A. Wang, H. Chen, L. Liu, K. Chen, Z. Lin, J. Han, and G. Ding. 2024. YOLOv10: Real-Time End-to-End Object Detection. *arXiv preprint arXiv:2405.14458* (2024). <https://arxiv.org/abs/2405.14458>
- [79] G. Wang, J. Zhang, K. Zhang, R. Huang, and L. Fang. 2024. GiganticNVS: Gigapixel Large-Scale Neural Rendering With Implicit Meta-Deformed Manifold. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 46, 1 (2024), 338–353. doi:10.1109/TPAMI.2023.3323069
- [80] X. Wang, R. Girdhar, S. X. Yu, and I. Misra. 2023. Cut and Learn for Unsupervised Object Detection and Instance Segmentation. *arXiv preprint arXiv:2301.11320* (2023). <https://arxiv.org/abs/2301.11320>
- [81] Y. Wang, G. Huang, S. Song, X. Pan, Y. Xia, and C. Wu. 2022. Regularizing Deep Networks With Semantic Data Augmentation. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 44, 7 (2022), 3733–3748. doi:10.1109/TPAMI.2021.3056495
- [82] Ying Wang, Fuyuan Ma, Zhaoqi Yang, Yaodi Zhu, Bo Yang, Pengfei Shen, and Lei Yun. 2025. Rumor Detection with Adaptive Data Augmentation and Adversarial Training. *Journal of Artificial Intelligence Research* 82 (2025), 1–34. doi:10.1613/jair.1.16963
- [83] Z. Wang, P. Wang, K. Liu, P. Wang, Y. Fu, C.-T. Lu, C. C. Aggarwal, J. Pei, and Y. Zhou. 2024. A Comprehensive Survey on Data Augmentation. *arXiv preprint arXiv:2405.09591 [cs.LG]* (May 2024). <https://arxiv.org/abs/2405.09591>
- [84] Genta Indra Winata, Hanyang Zhao, Anirban Das, Wenpin Tang, David D. Yao, Shi-Xiong Zhang, and Sambit Sahu. 2025. Preference Tuning with Human Feedback on Language, Speech, and Vision Tasks: A Survey. *Journal of Artificial Intelligence Research* 82 (2025), 1–55. doi:10.1613/jair.1.17541
- [85] Xiaodan Xing, Fadong Shi, Jiahao Huang, Yinzhe Wu, Yang Nan, Sheng Zhang, Yingying Fang, Michael Roberts, Carola-Bibiane Schönlieb, Javier Del Ser, and Guang Yang. 2025. On the caveats of AI autophagy. *Nature Machine Intelligence* 7 (2025), 172–180. doi:10.1038/s42256-025-00984-1
- [86] H. Yang, J. Li, K. Z. Lim, C. Pan, T. V. Truong, Q. Wang, K. Li, S. Li, X. Xiao, M. Ding, T. Chen, X. Liu, Q. Xie, P. Valdivia y. Alvarado, A. Wang, and P.-Y. Chen. 2022. Automatic strain sensor design via active learning and data augmentation for soft machines. *Nature Machine Intelligence* 4, 1 (2022), 84–94. doi:10.1038/s42256-021-00434-8
- [87] J. Yang, B. Hu, H. Li, Y. Liu, X. Gao, J. Han, F. Chen, and X. Wu. 2025. Dynamic VAEs via semantic-aligned matching for continual zero-shot learning. *Pattern Recognition* 160 (2025), 111199. doi:10.1016/j.patcog.2023.111199
- [88] L. Yang, S. Wang, J. Yang, T. Nguyen, A. S. Aved, and J. Z. Wang. 2021. Artificial Data Synthesis for Machine Learning: A Review. *Artificial Intelligence* 294 (2021), 103492. <https://www.sciencedirect.com/science/article/pii/S0004370221000644>
- [89] L. Yang, Z. Zhang, Y. Song, S. Hong, R. Xu, Y. Zhao, W. Zhang, B. Cui, and M.-H. Yang. 2024. Diffusion Models: A Comprehensive Survey of Methods and Applications. *Comput. Surveys* (2024). <https://arxiv.org/abs/2209.00796> arXiv preprint arXiv:2209.00796, accepted by ACM Computing Surveys.
- [90] M. Yang, X. Wang, and G. Zhao. 2020. Single-shot object detection with enriched semantics. *Pattern Recognition* 106 (2020), 107404. <https://www.sciencedirect.com>

Conference'17, July 2017, Washington, DC, USA

- com/science/article/pii/S0031320319303711
- [91] Zichen Yang, Haifeng Liu, and Deng Cai. 2019. On the Diversity of Conditional Image Synthesis With Semantic Layouts. *IEEE Transactions on Image Processing* 28, 6 (2019), 2898–2907. doi:10.1109/TIP.2019.2891935
- [92] B. Yin, X. Zhang, L. Liu, M.-M. Cheng, Y. Liu, and Q. Hou. 2025. Camouflaged Object Detection with Adaptive Partition and Background Retrieval. *International Journal of Computer Vision* (2025). doi:10.1007/s11263-025-02406-6
- [93] Chaoyang Zhu and Long Chen. 2024. A Survey on Open-Vocabulary Detection and Segmentation: Past, Present, and Future. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 46, 12 (2024), 8954–8975. doi:10.1109/TPAMI.2024.3413013